

# Deep Learning Bootcamp

A hands-on intensive workshop covering neural fundamentals,  
computer vision, and natural language processing using modern deep learning tools.

by: Lama Ayash

## Day 2

# Course Outlines

**01** Computer Vision

---

**02** Computer Vision Tasks

---

**03** Convolution Layer

---

**04** Padding

---

**05** Stride

---

**06** Output Dimensions in Convolution

---

**07** Output Dimensions Example

**08** Pooling Layer

---

**09** Type of Pooling Layer

---

**10** Dense Layer

---

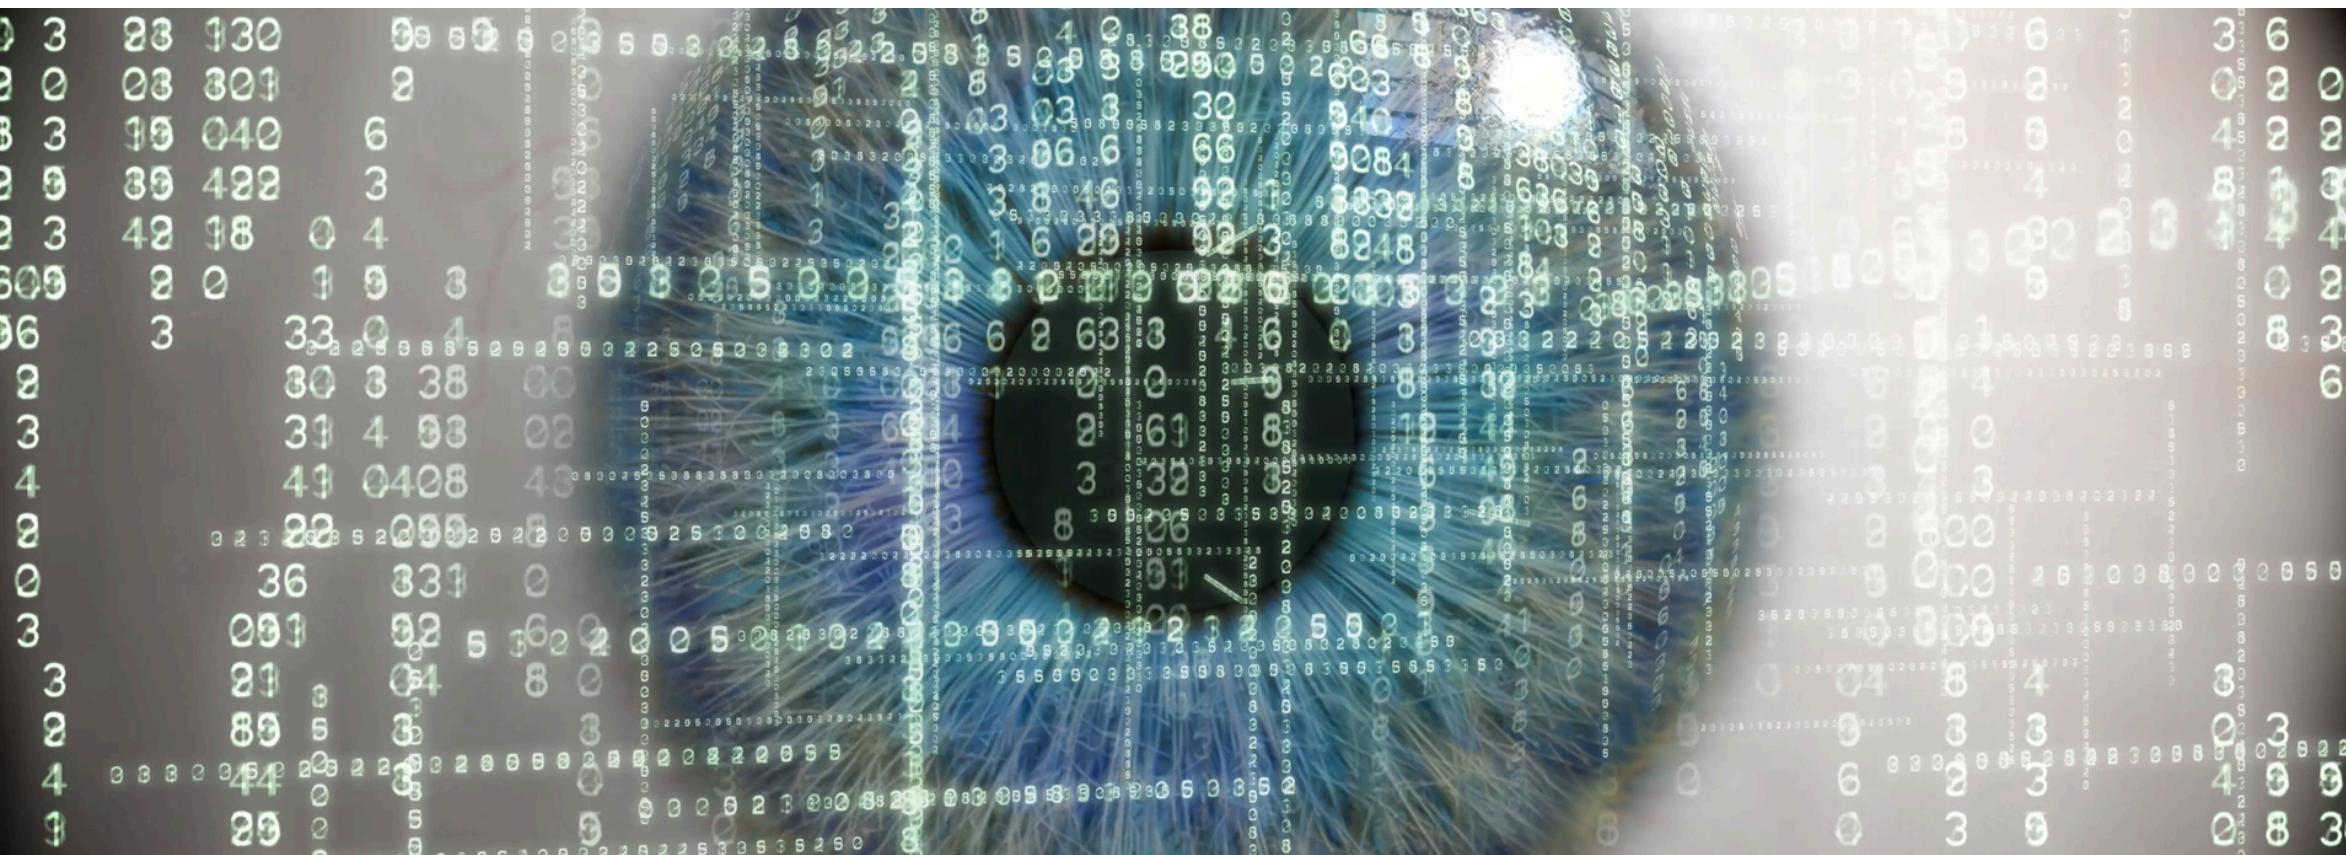
**11** Learnable Parameters in CNN

---

**12** Q&A

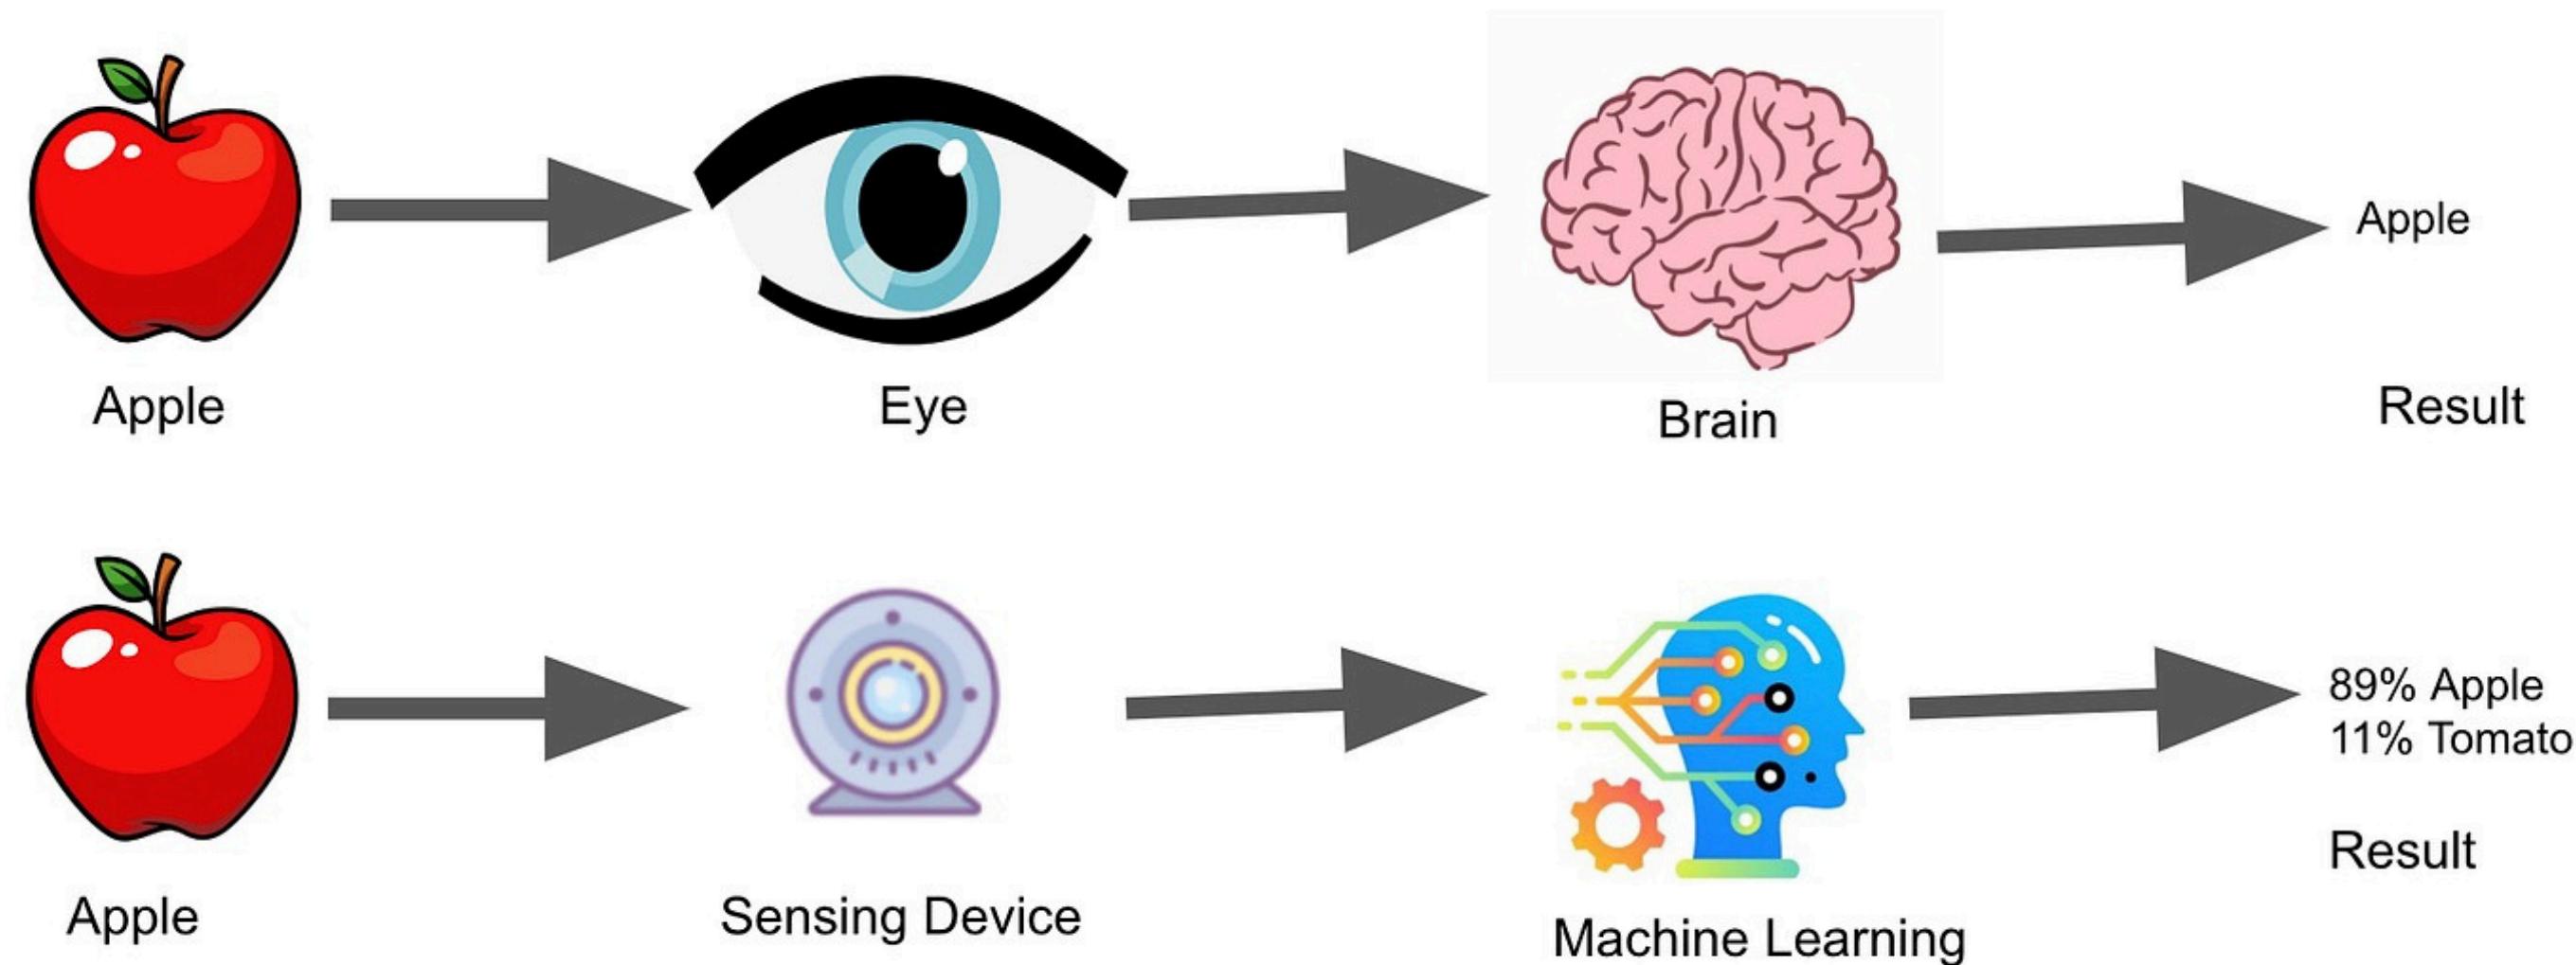
# Computer Vision

**Computer Vision (CV)** enables machines to see, understand, and interpret the visual world. Inspired by how humans perceive and process visual information.



# Computer Vision

**Computer Vision (CV)** enables machines to see, understand, and interpret the visual world. Inspired by how humans perceive and process visual information.



# Computer Vision Tasks

- Image Classification

Classification is the simplest and most common task in computer vision.



*The class in this case is a cat*



*The class in this case is a dog*

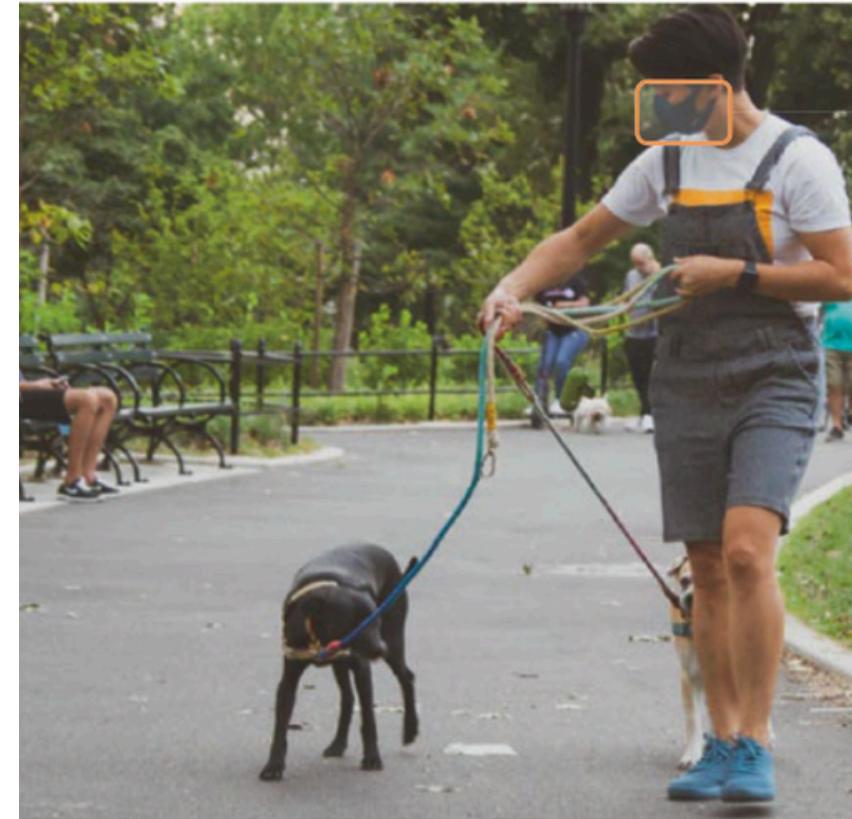
# Computer Vision Tasks

- Object Detection and Localization

The goal is to detect and locate specific objects within a larger image.



*Class: No mask detected*



*Class: A mask is detected in the image*

# Computer Vision Tasks

- **Image Segmentation**

Segmentation divides an image into distinct regions or objects based on color, texture, and boundaries.



# Computer Vision Tasks

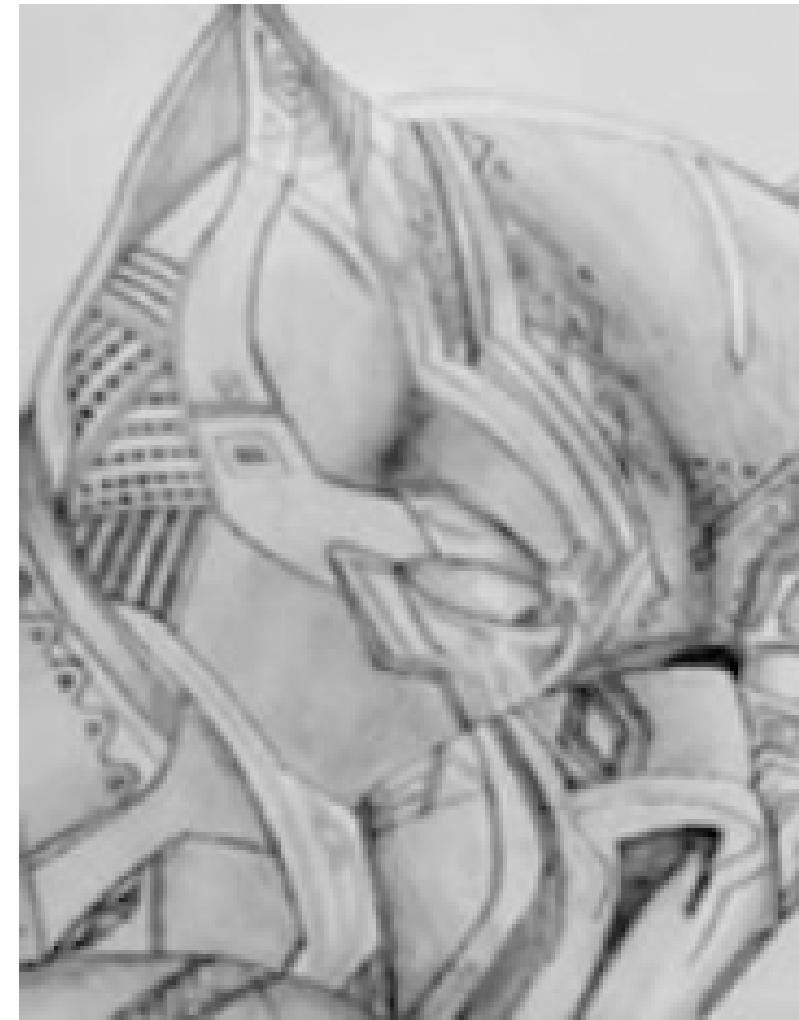
- Anomaly Detection
  - Detects unusual patterns or defects in images by comparing them to normal reference data.



*Perfect examples of steel pipes*

# Channels

- A **channel** represents one dimension of color or intensity information in an image.



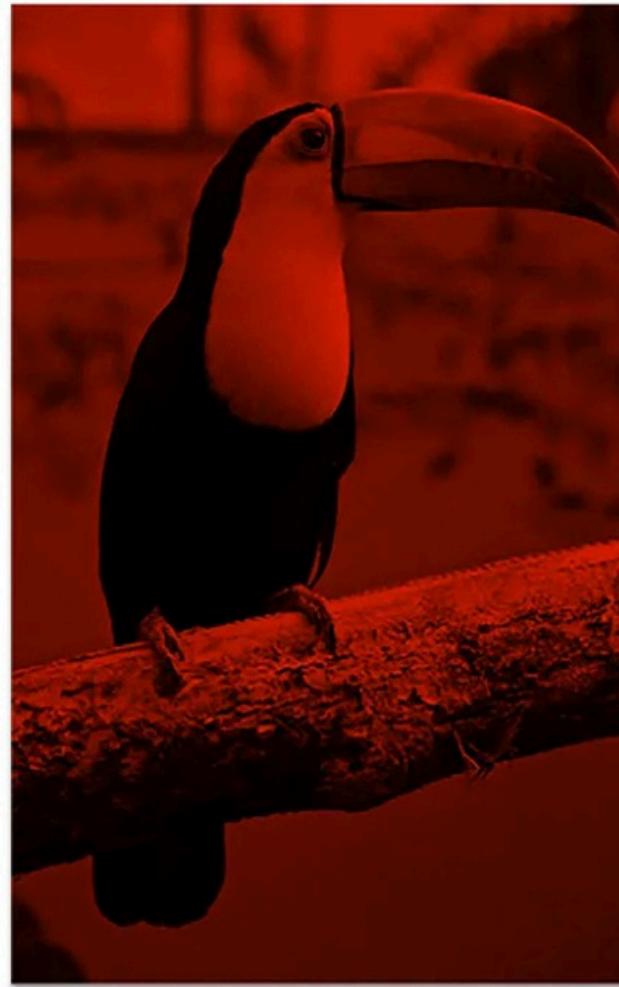
2	36	40	200
195	190	20	180
40	54	200	200
30	40	200	180

# Channels

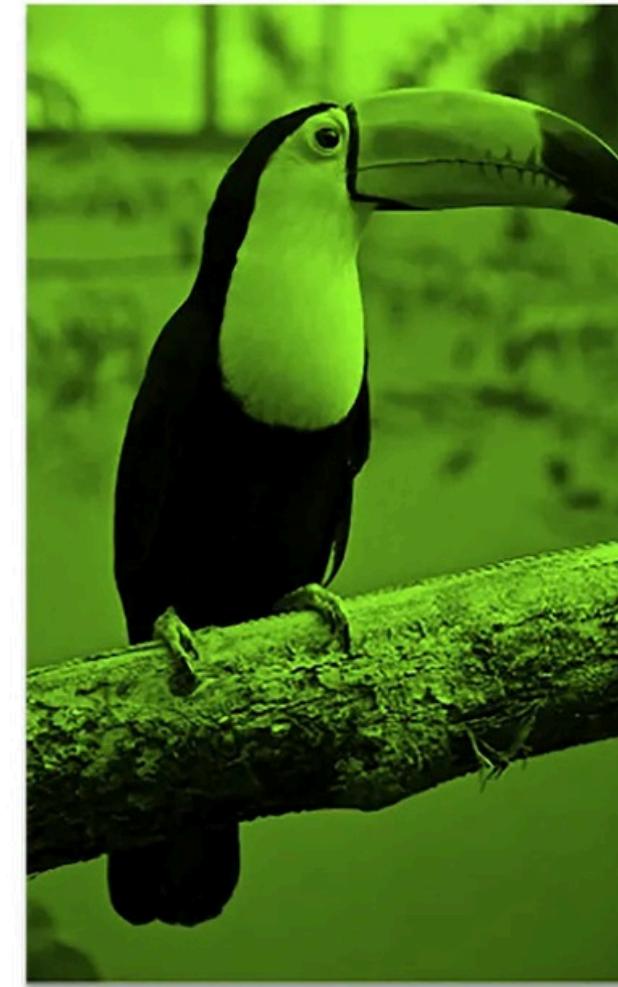
- A color image has three channels: Red, Green, and Blue (RGB).



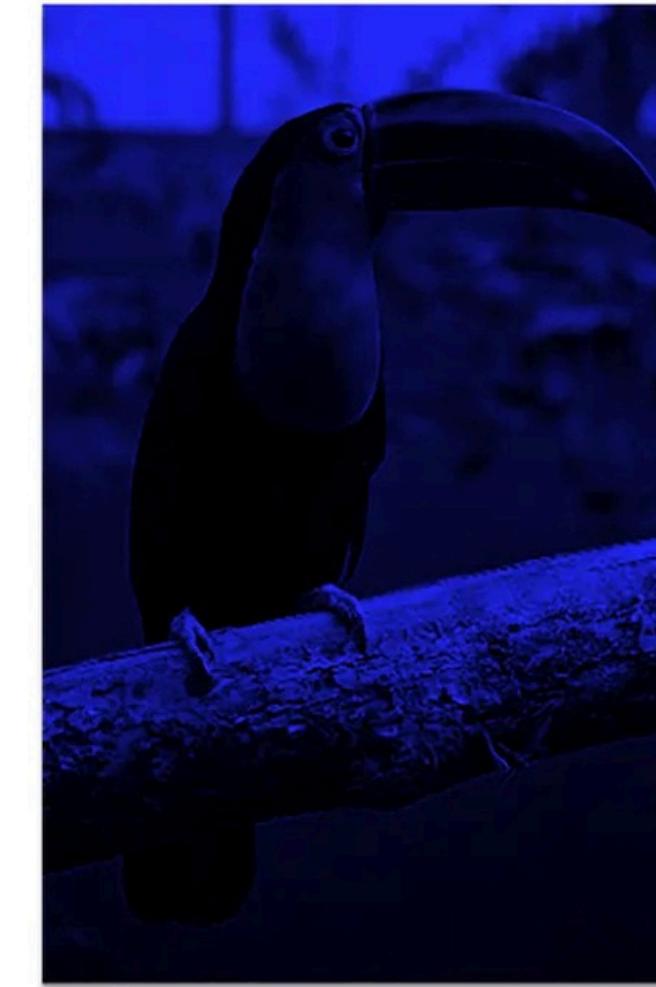
Original Image



Red Channel



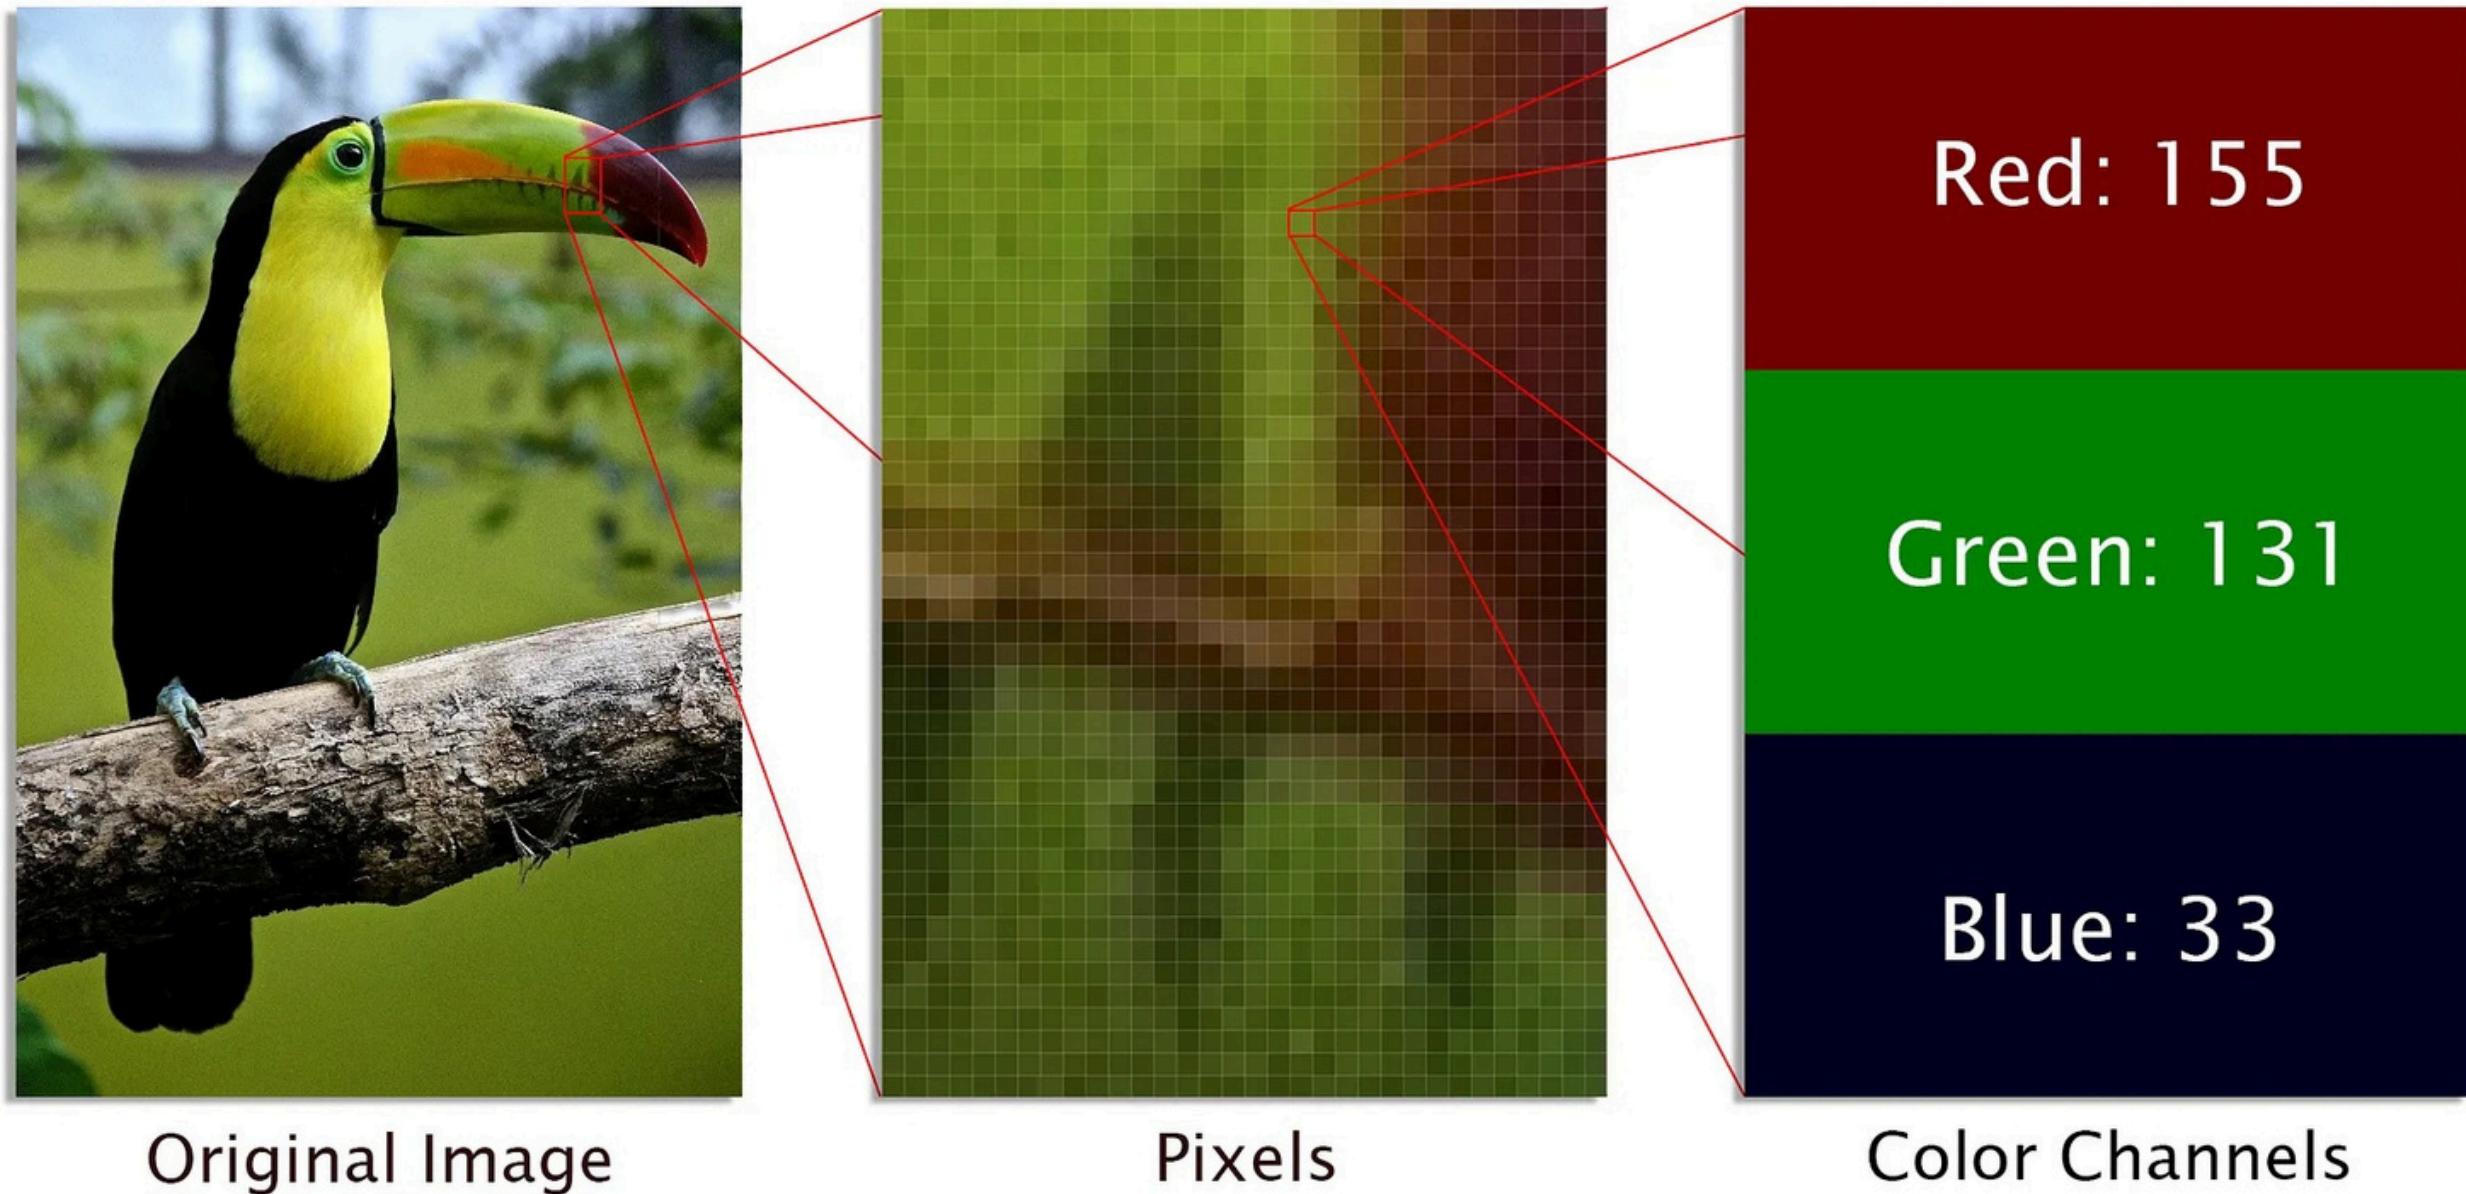
Green Channel



Blue Channel

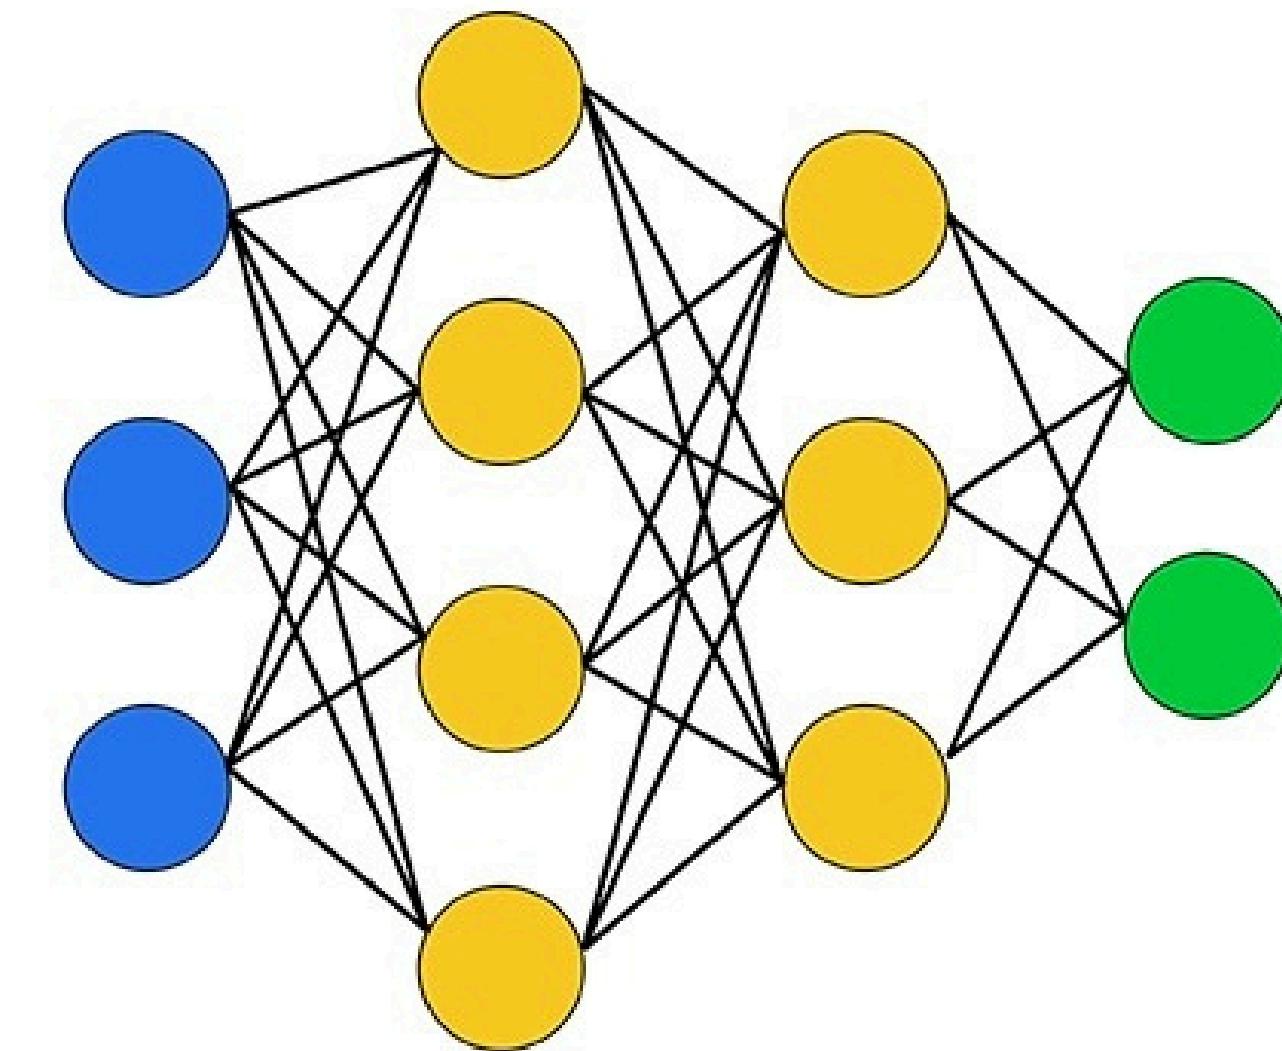
# Channels

- A color image has three channels: Red, Green, and Blue (RGB).



# Vision Models

- A color image has three channels: Red, Green, and Blue (RGB).



Input Layer      Hidden Layer      Output Layer

But can we apply the same approach to images?

# Vision Models

Little or no invariance to shifting, scaling, and other forms of distortion

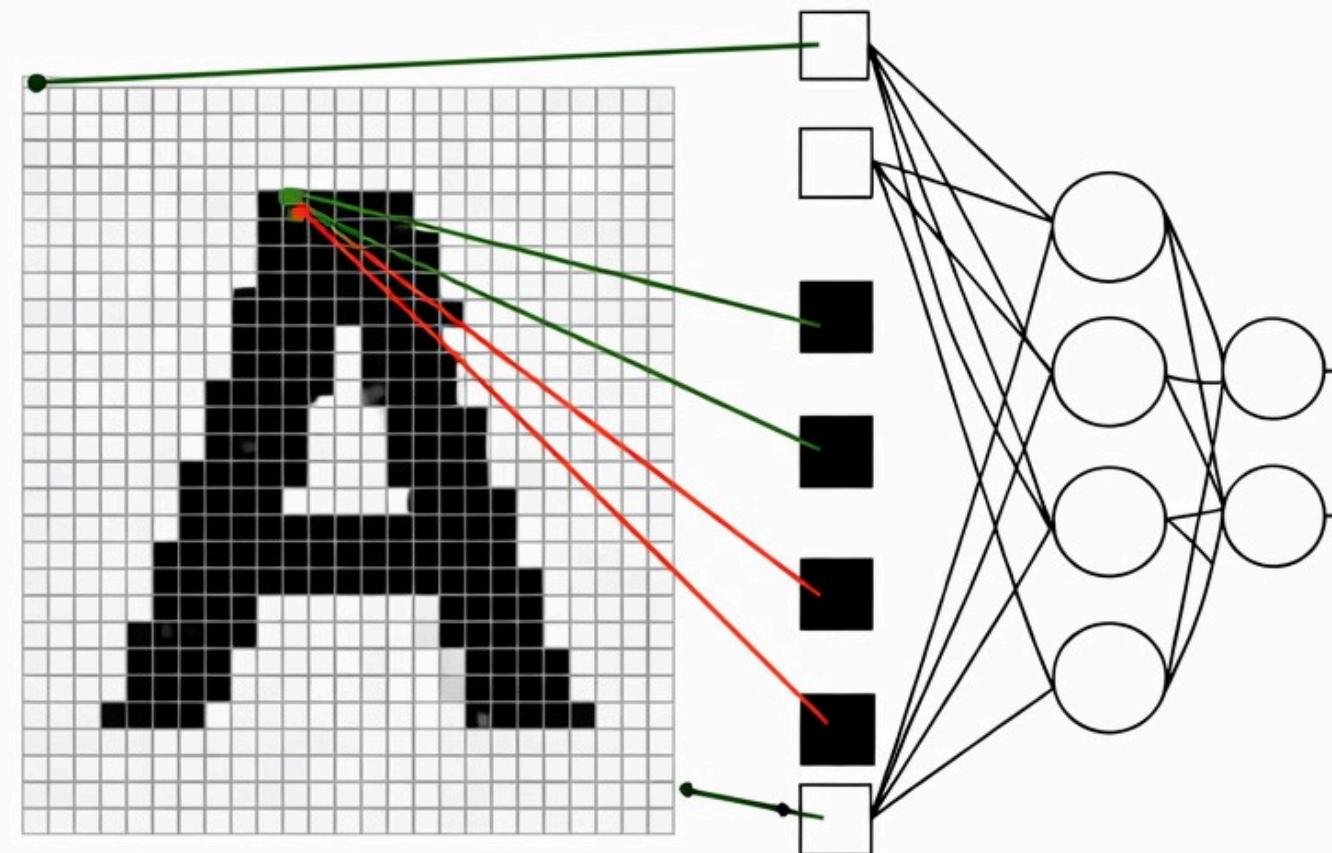


Figure 2: Original "A" character.

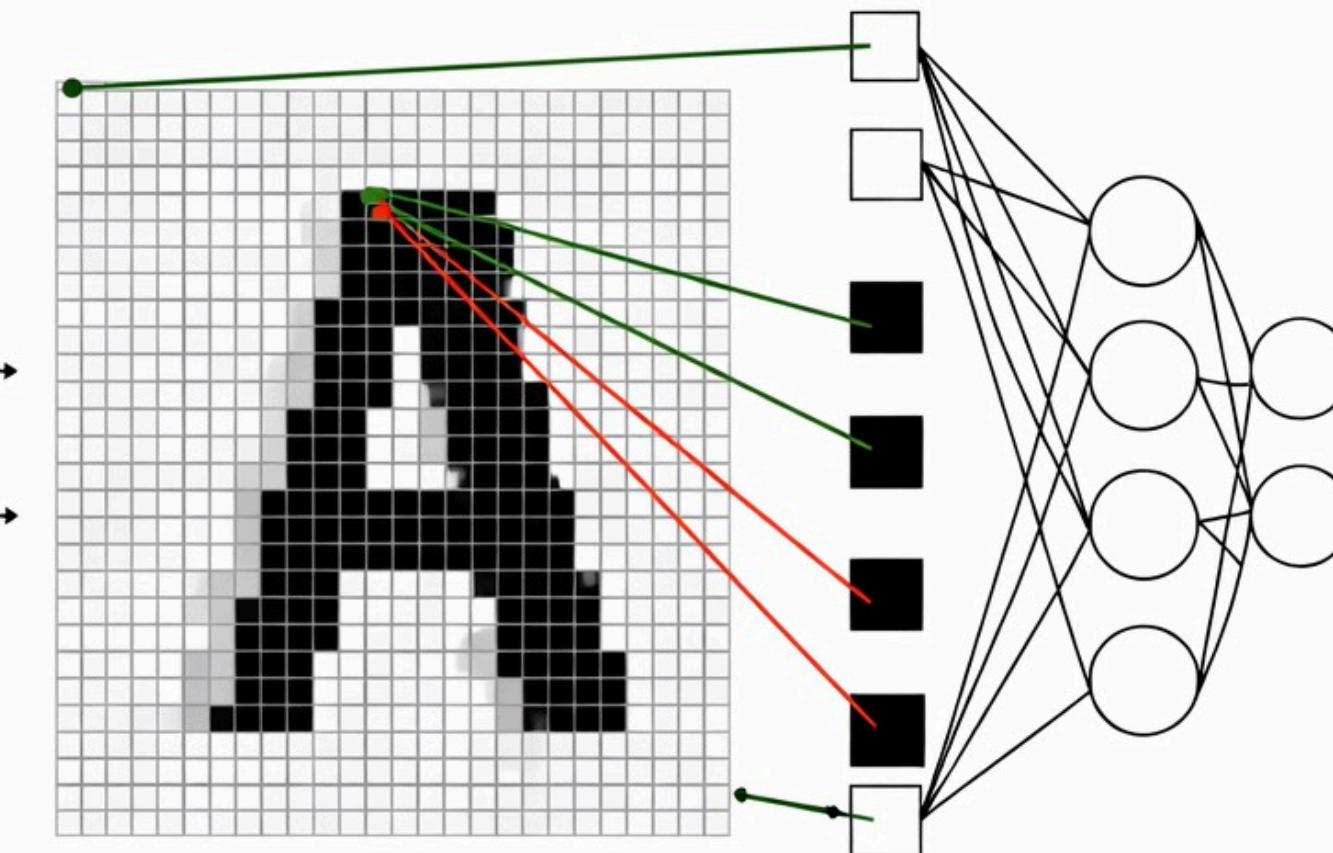
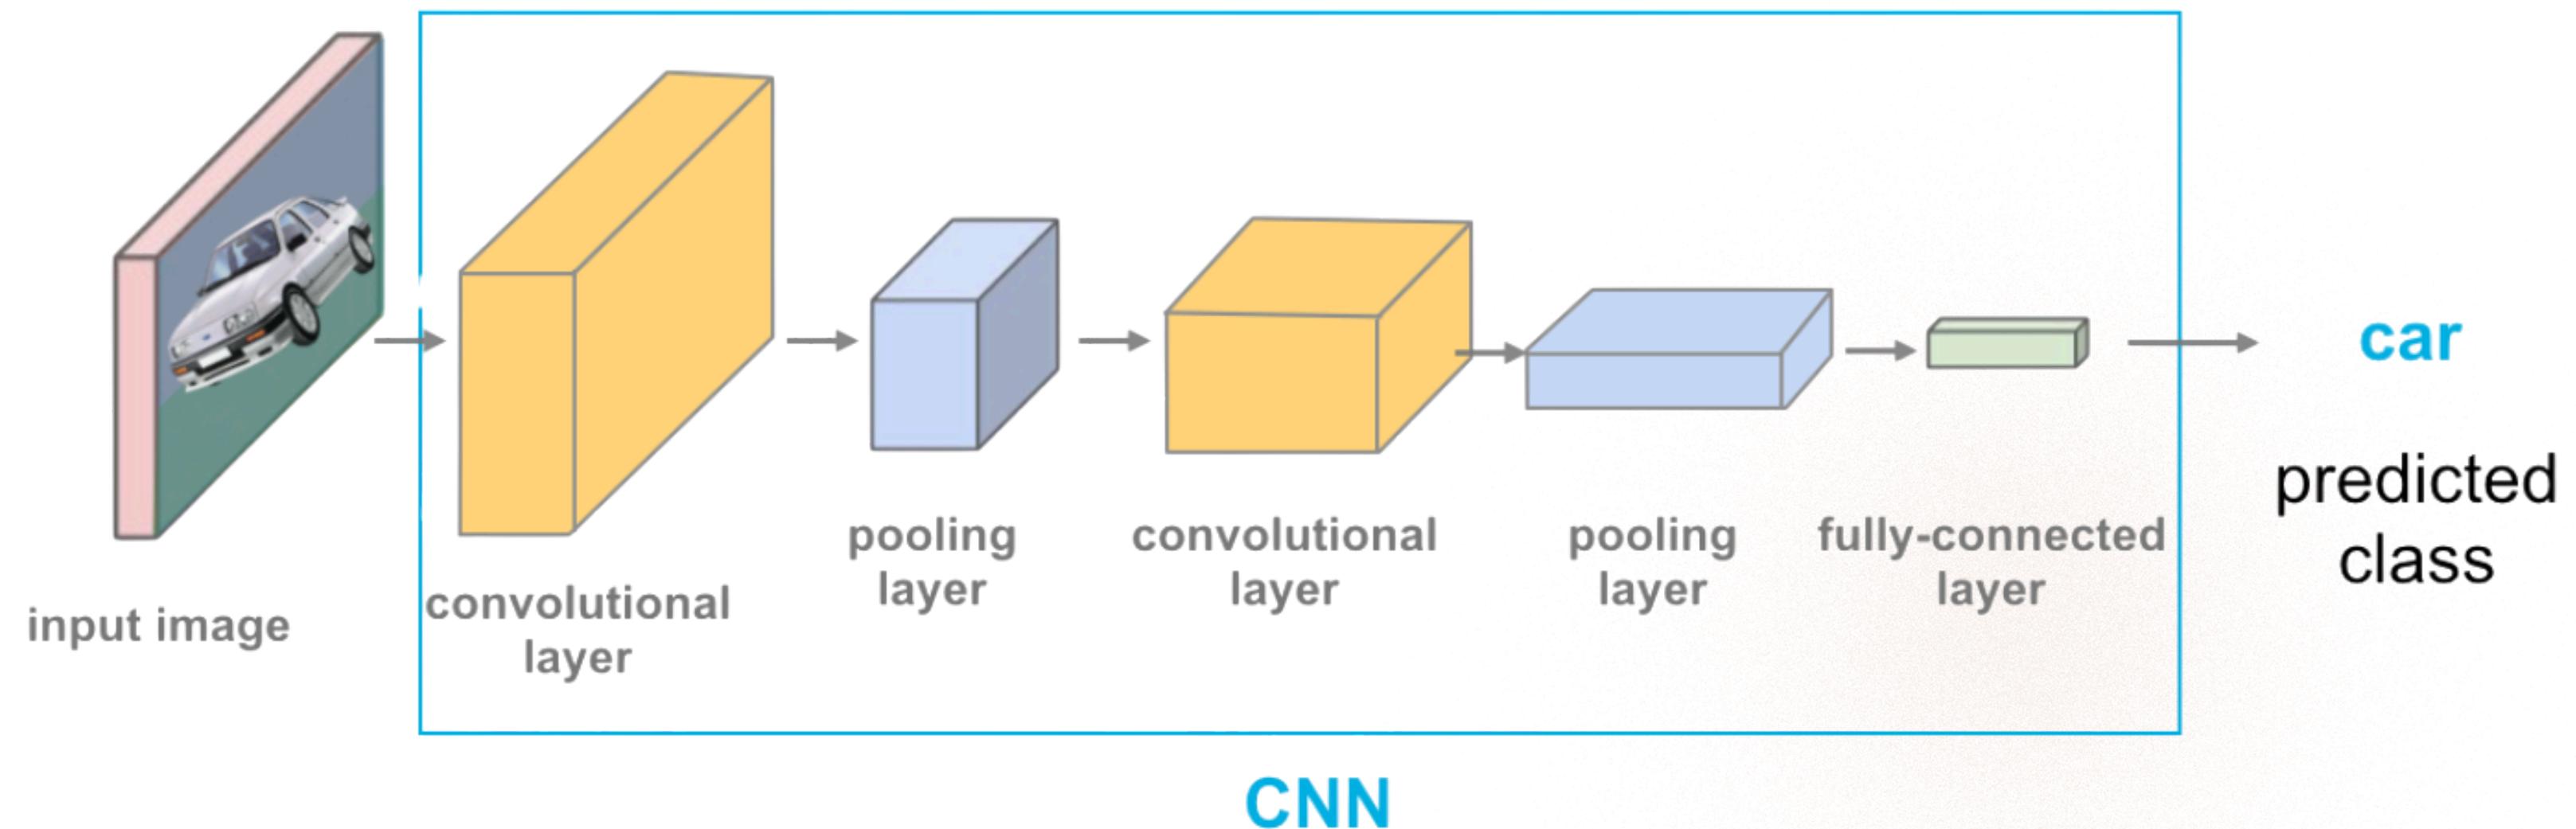


Figure 3: Shifted "A" character.

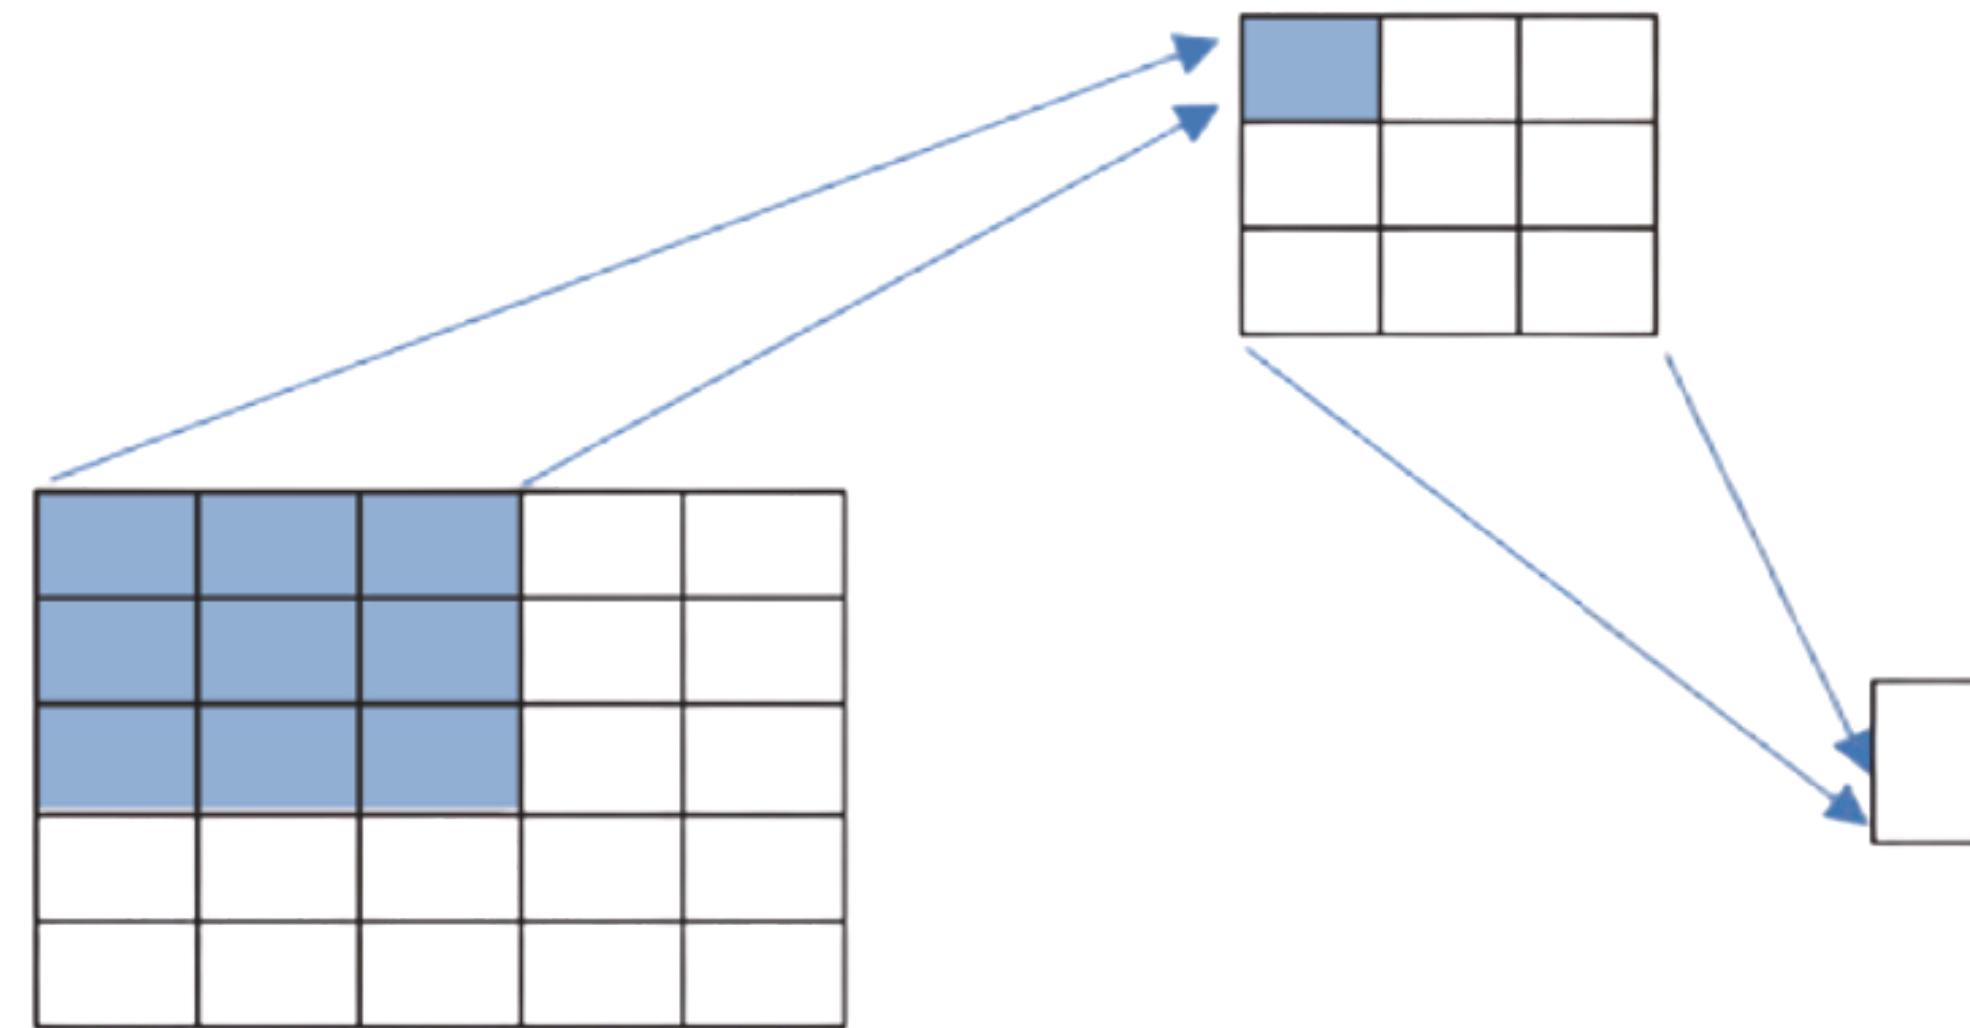
# Convolutional Neural Networks

- A CNN is a deep network of neurons with learnable filters that perform convolution operations on inputs, usually images, to extract hierarchical features.



# Convolutional Neural Networks

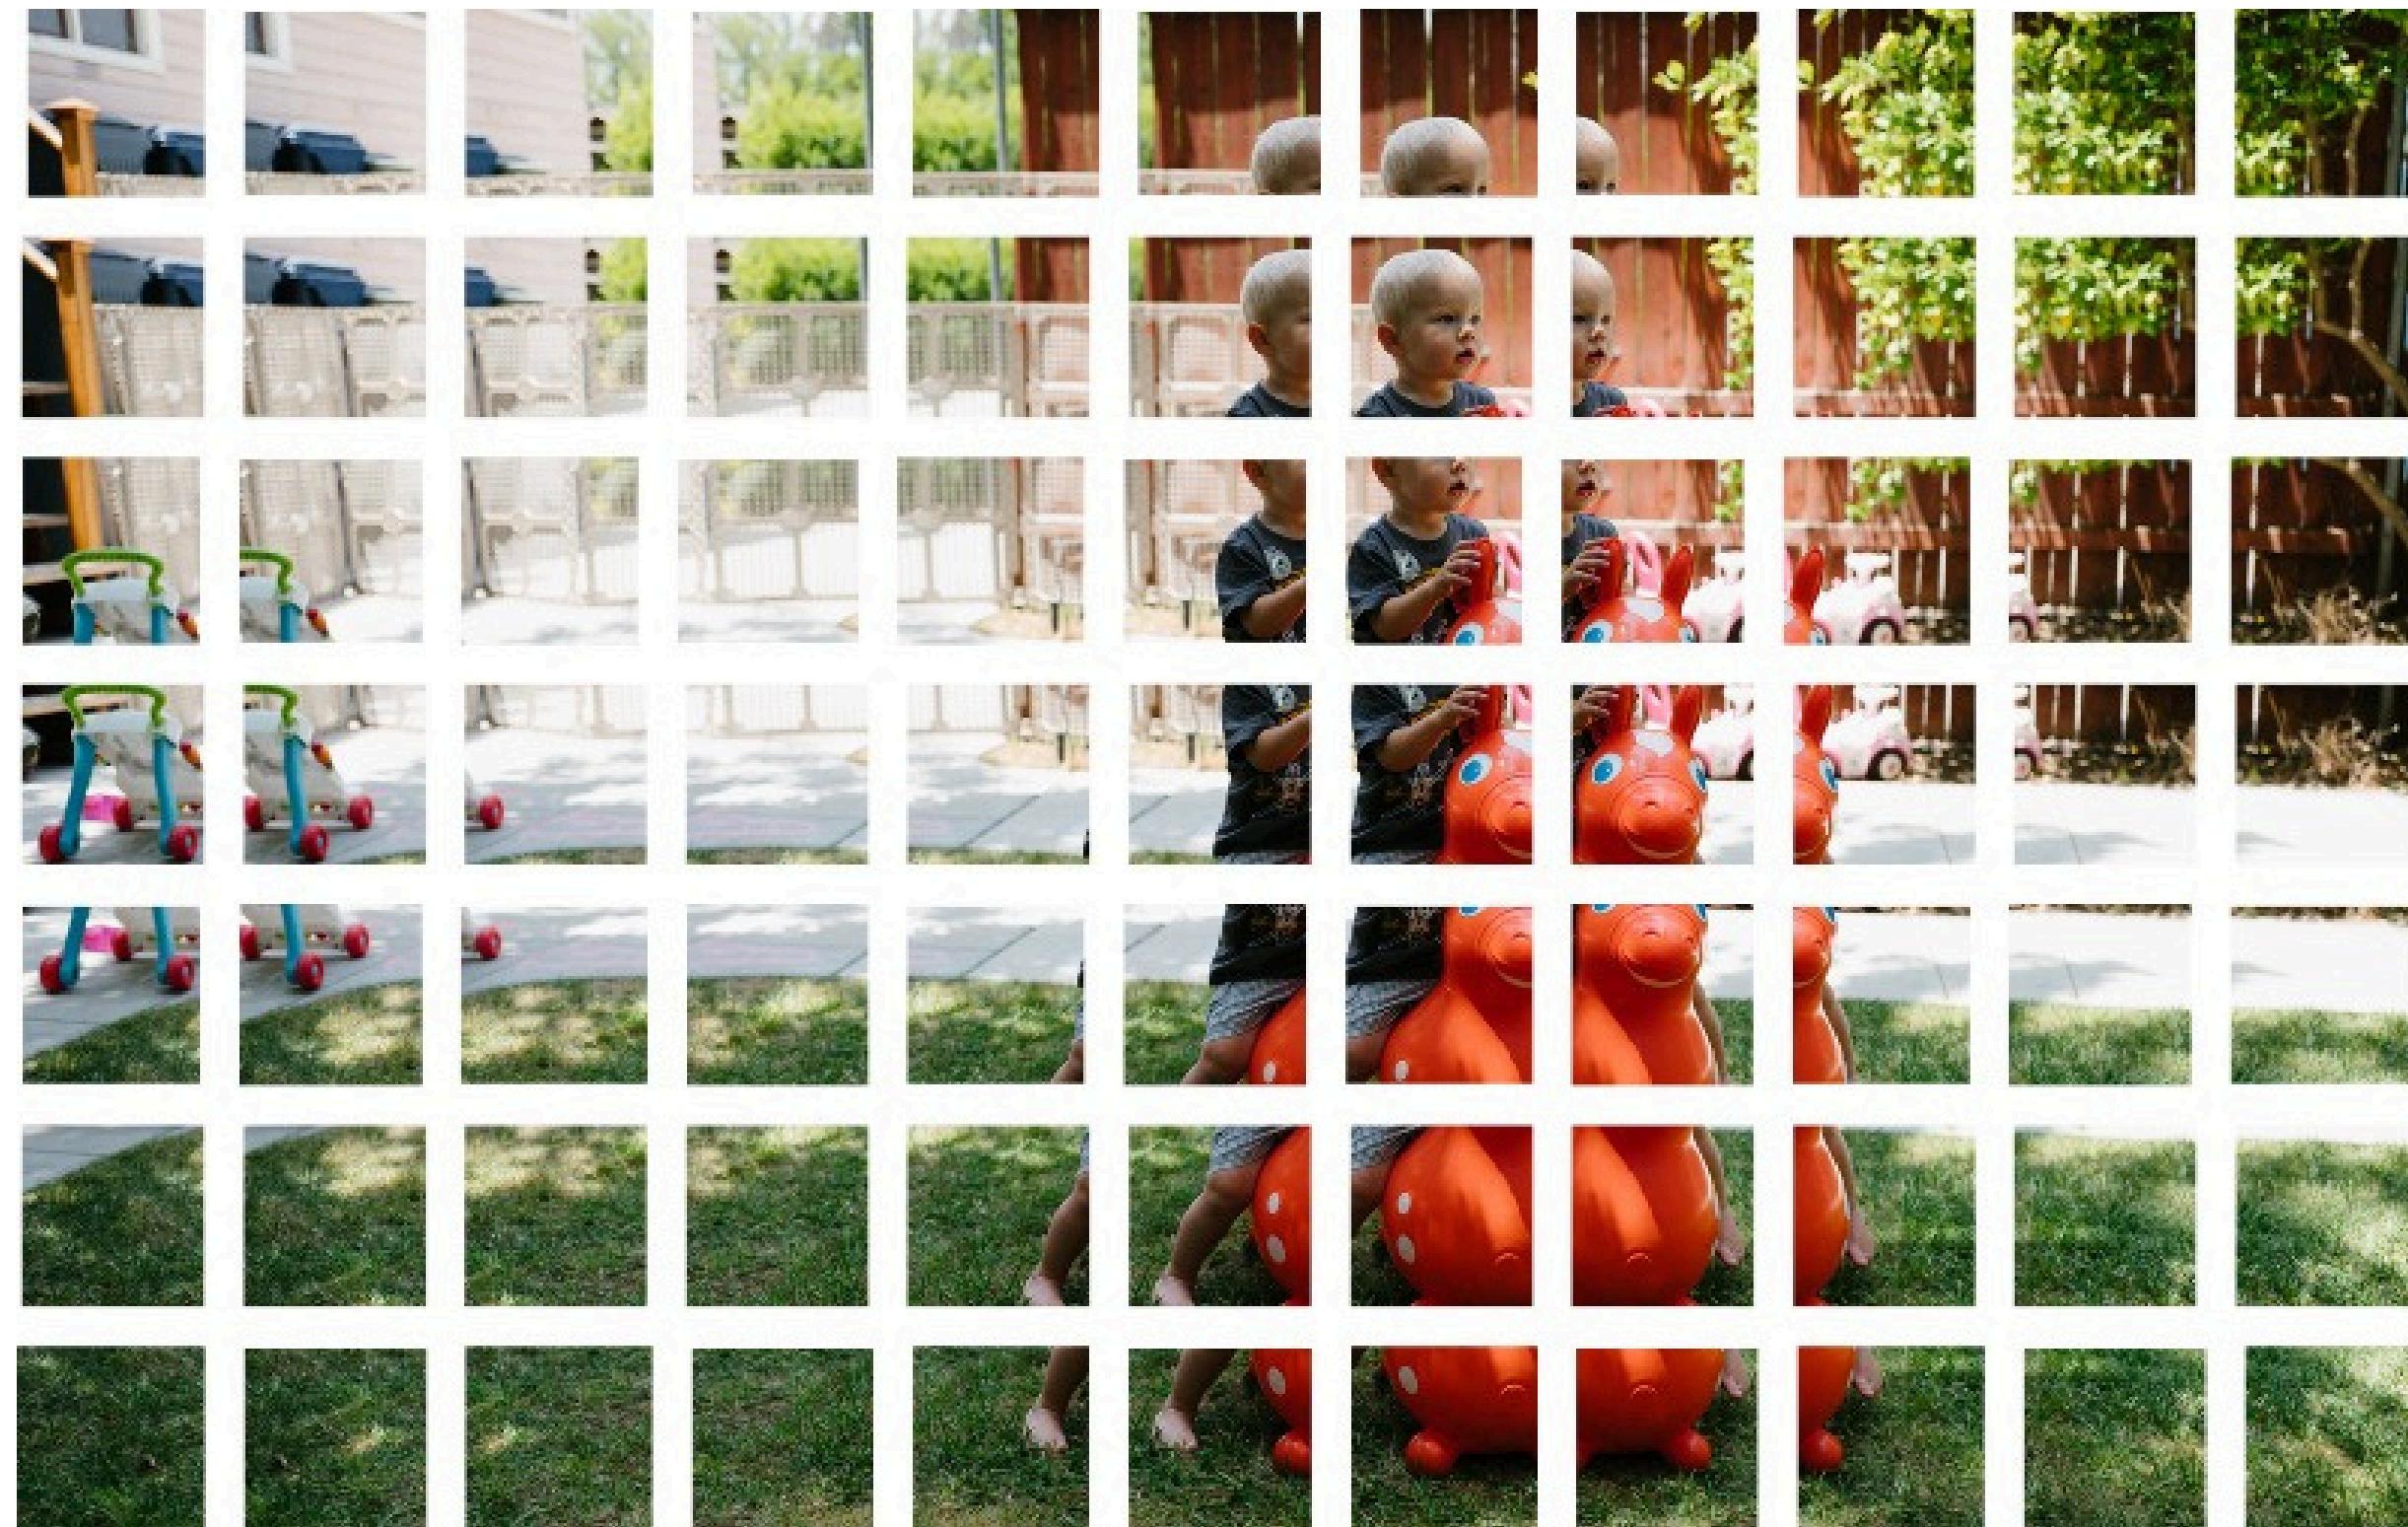
- It applies a set of learnable filters (**kernels**) that slide across the input, performing convolution operations to extract local patterns like edges, textures, and shapes.



# Convolutional Neural Networks

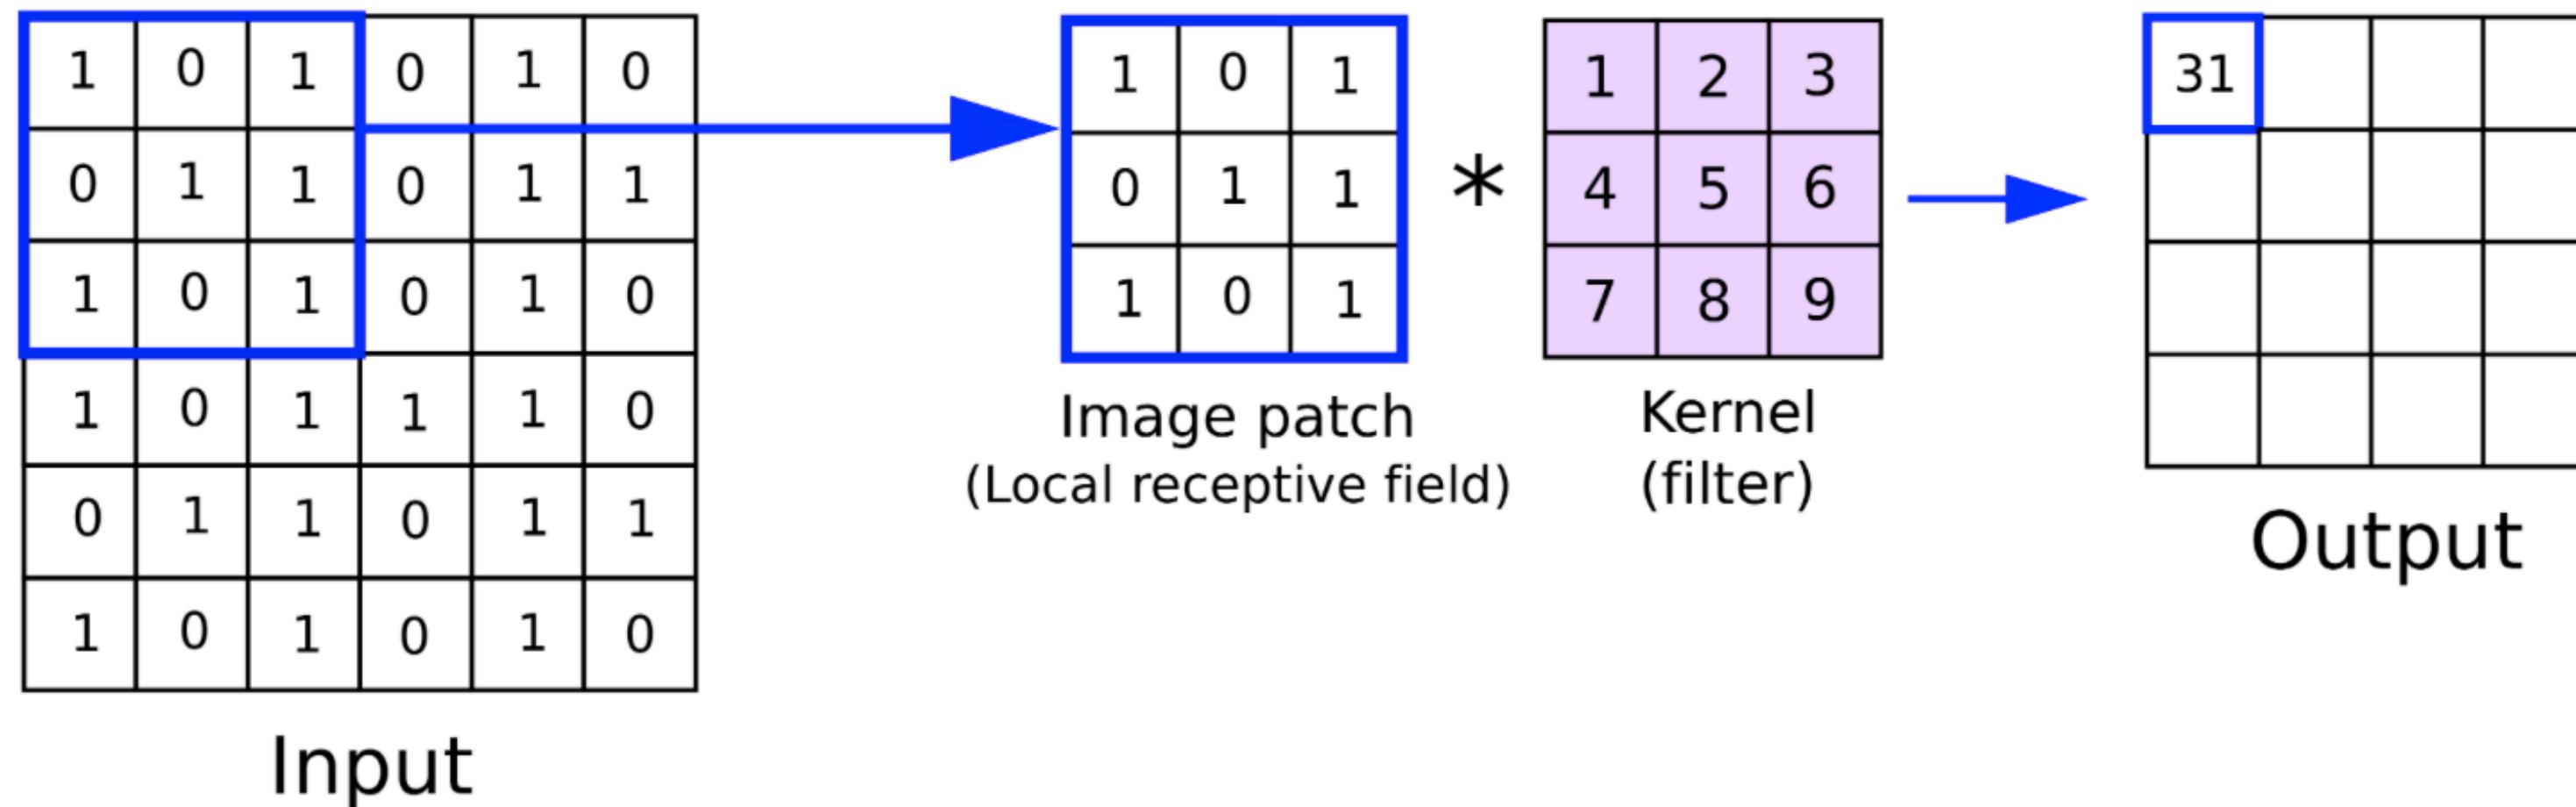


# Convolutional Neural Networks



# Convolution Layer

- Convolution involves applying a kernel (filter) over an image to extract features.
- A kernel is a small matrix (like:  $3\times 3$ ,  $5\times 5$ ).



# Convolution Layer

- Edge detection is an example of the usefulness of convolution in identifying the regions in an image where there is a sharp change in colors or intensity.

$$\begin{array}{|c|c|c|c|c|c|c|c|c|} \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline \end{array} \quad * \quad \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|c|c|c|c|c|c|c|} \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline \end{array}$$

Vertical

The diagram illustrates a convolution operation for vertical edge detection. It shows an input image of size 9x8, a kernel of size 3x3, and an output image of size 7x7. The input image has a blue box around its top-left 3x3 block, a pink box around its middle row, and a green box around its bottom row. The kernel is labeled "Vertical". The resulting output image has a blue box around its top-left 3x3 block, a pink box around its middle row, and a green box around its bottom row, highlighting the value 30 at the bottom-right position.

*An example of vertical edge detection*

# Convolution Layer

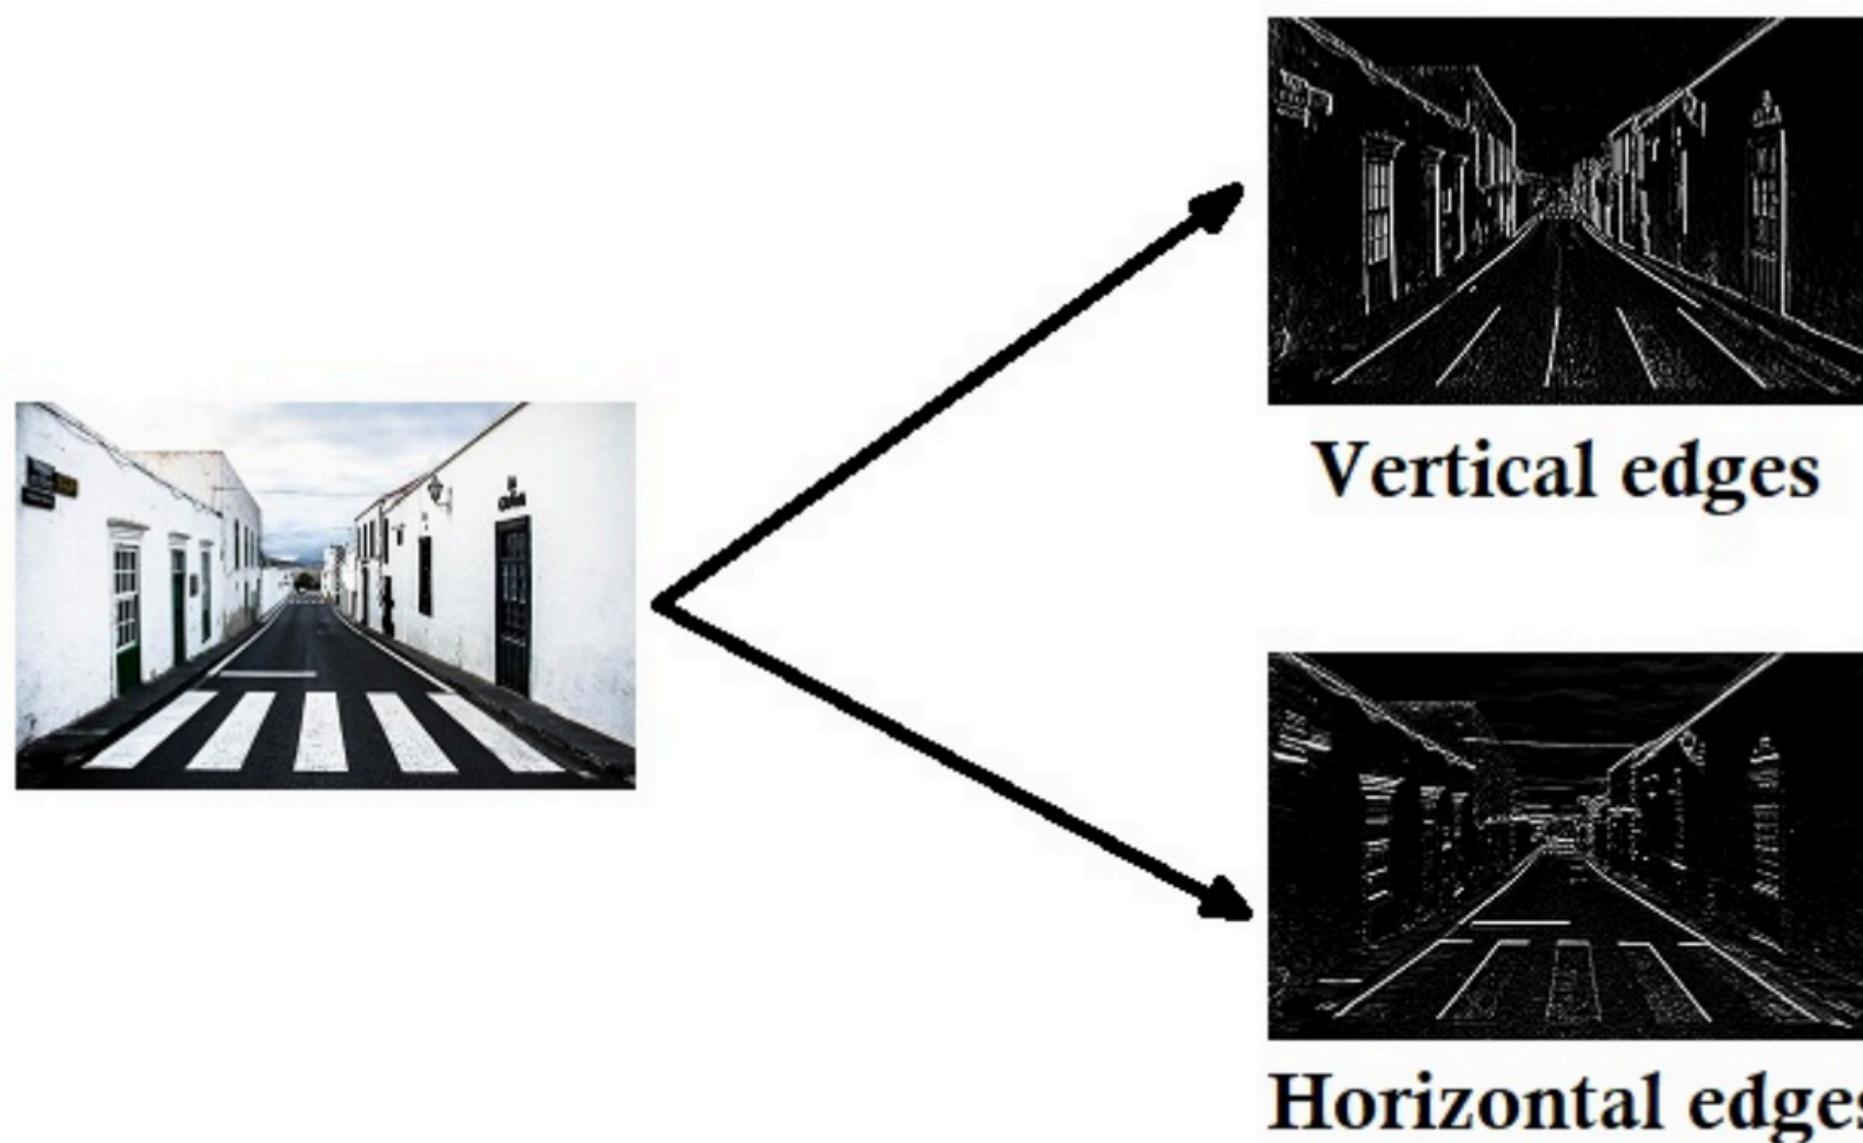
- Edge detection is an example of the usefulness of convolution in identifying the regions in an image where there is a sharp change in colors or intensity.

$$\begin{matrix} 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 10 & 10 & 10 & 10 \\ 0 & 0 & 0 & 0 & 10 & 10 & 10 & 10 \\ 0 & 0 & 0 & 0 & 10 & 10 & 10 & 10 \end{matrix} * \begin{matrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{matrix}_{\text{Horizontal}} = \begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 30 & 30 & 10 & -10 & -30 & -30 \\ 30 & 30 & 10 & -10 & -30 & -30 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{matrix}$$

An example of horizontal edge detection

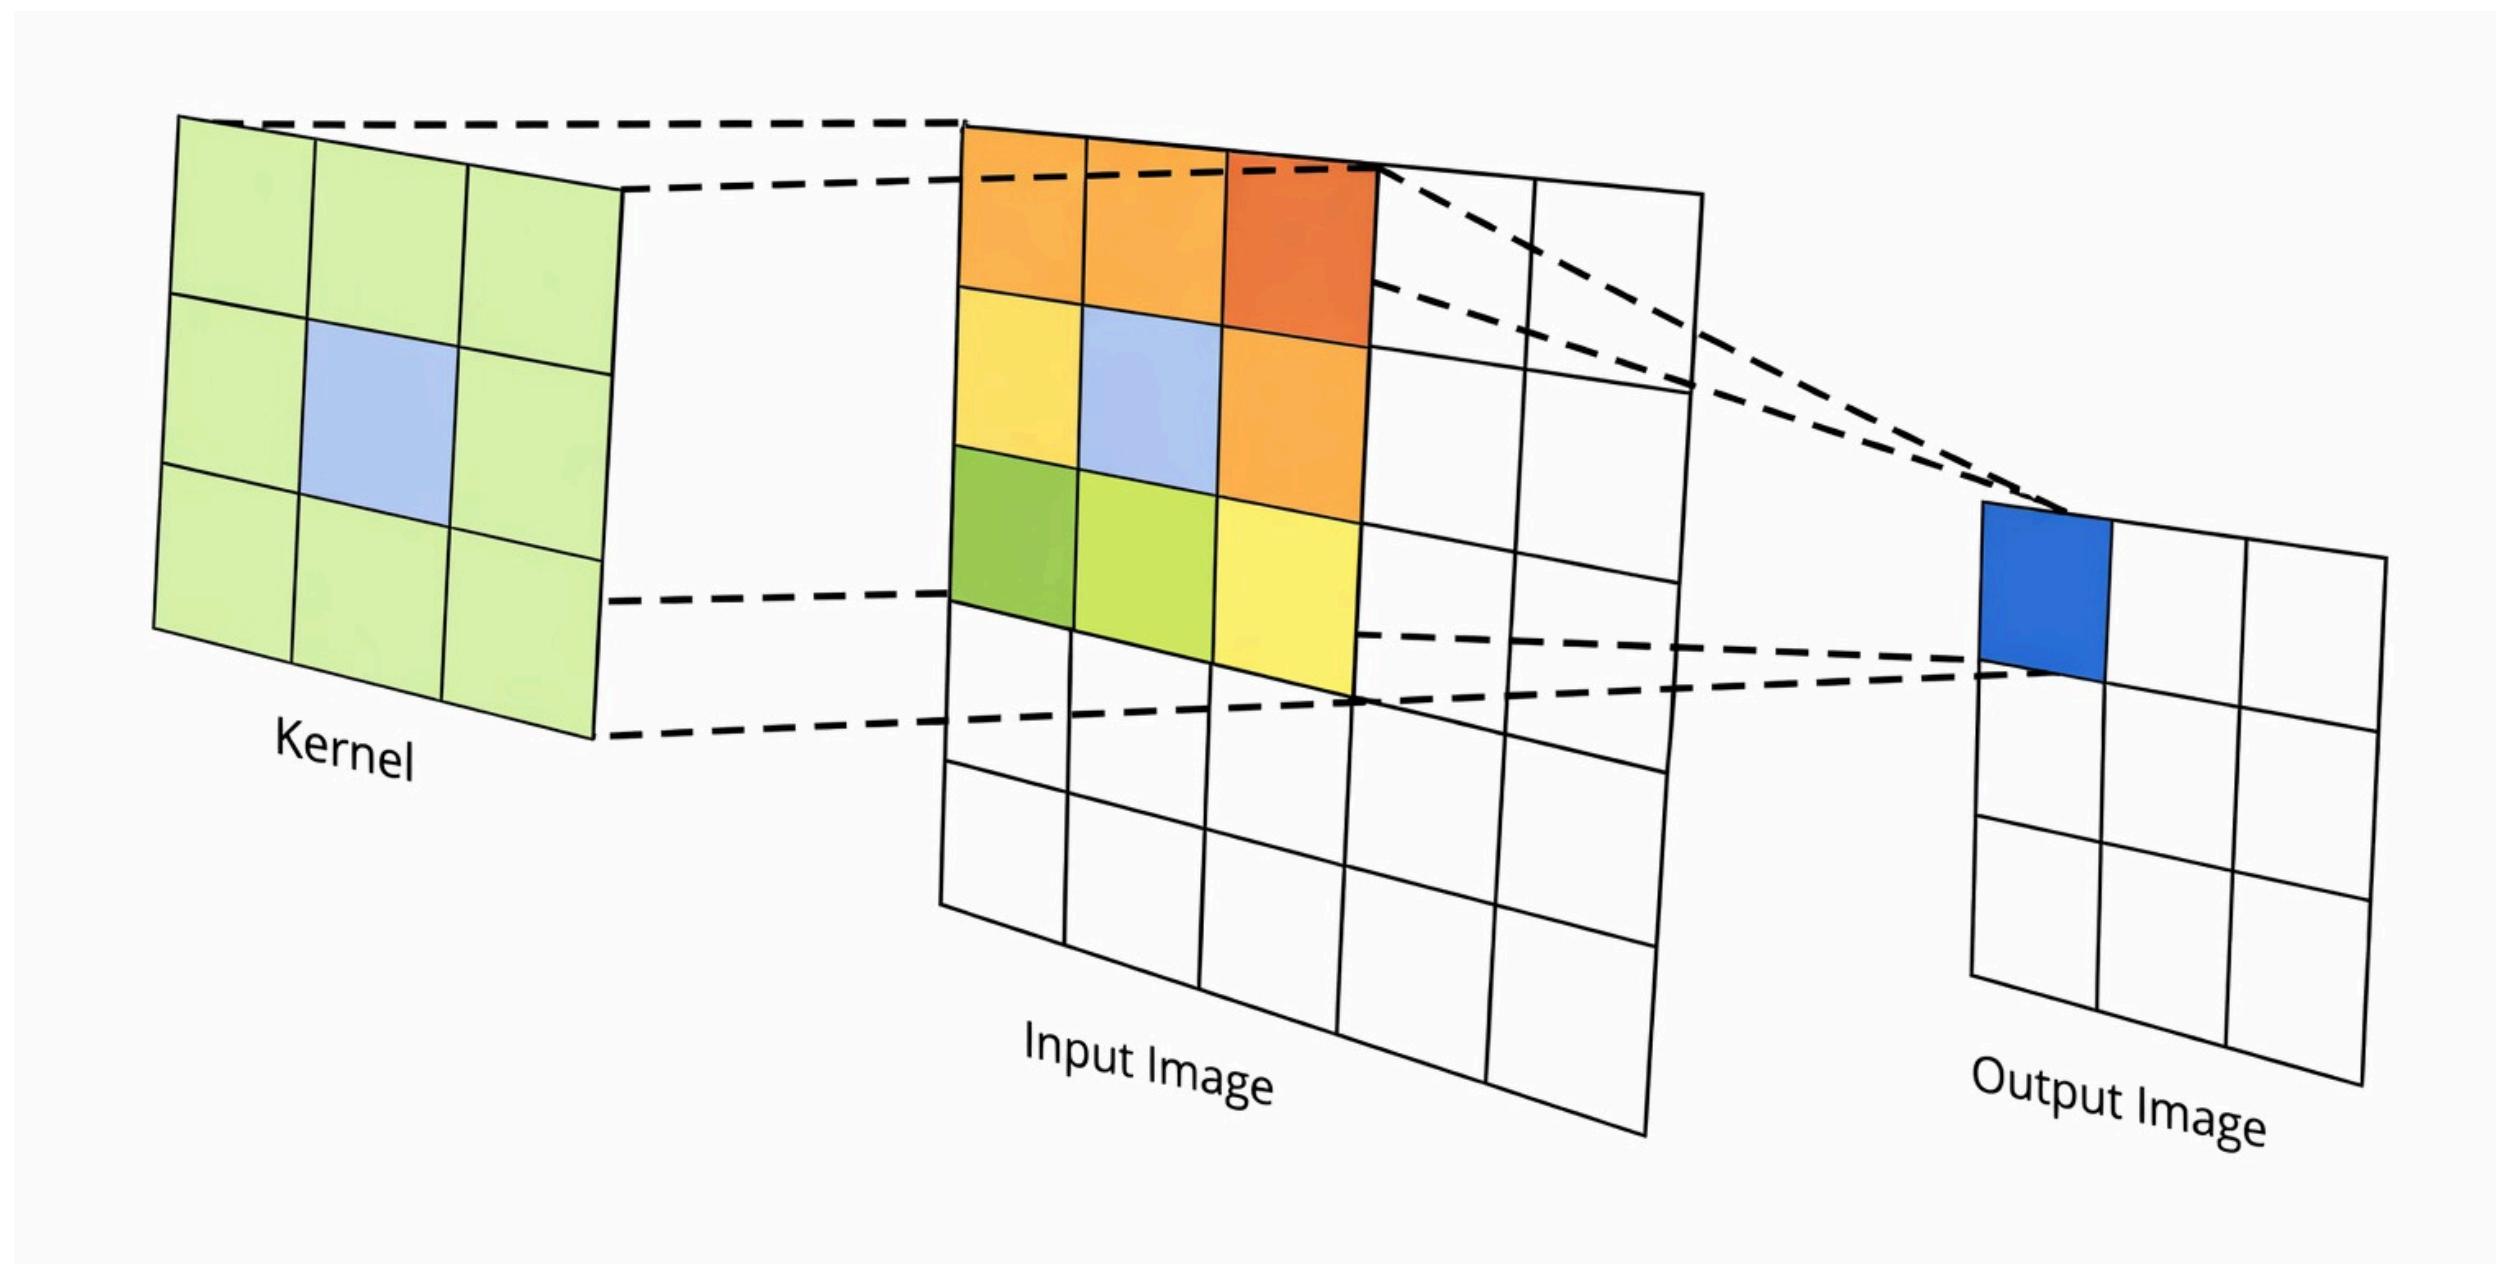
# Convolution Layer

- Edge detection is an example of the usefulness of convolution in identifying the regions in an image where there is a sharp change in colors or intensity.



# Convolution Layer

- The kernel is slid across the image, computing a dot product with local regions at each position.



# Padding

- Padding means adding extra layers of zeros around the image borders.
- If padding size = p, new input dimension becomes:

$$(n + 2p) \times (n + 2p)$$

0	0	0	0	0	0	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	10	10	10	10	10	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

\* Vertical =

1	0	-1
1	0	-1
1	0	-1

-20	0	0	20	20	0	0	0
-30	0	0	30	30	0	0	0
-30	0	0	30	30	0	0	0
-30	0	0	30	30	0	0	0
-30	0	0	30	30	0	0	0
-30	0	0	30	30	0	0	0
-30	0	0	30	30	0	0	0
-20	0	0	20	20	0	0	0

# Stride

- Stride refers to the speed at which a filter translates across an image.

Stride = 1

2	0	1	2	2
0	1	0	1	3
3	2	0	1	1
1	0	0	1	2
2	2	1	0	0

2	0	1	2	2
0	1	0	1	3
3	2	0	1	1
1	0	0	1	2
2	2	1	0	0

...

2	0	1	2	2
0	1	0	1	3
3	2	0	1	1
1	0	0	1	2
2	2	1	0	0

:

Stride = 2

2	0	1	2	2
0	1	0	1	3
3	2	0	1	1
1	0	0	1	2
2	2	1	0	0

2	0	1	2	2
0	1	0	1	3
3	2	0	1	1
1	0	0	1	2
2	2	1	0	0

...

2	0	1	2	2
0	1	0	1	3
3	2	0	1	1
1	0	0	1	2
2	2	1	0	0

:

## Output Dimensions in Convolution

$$n_{\text{out}} = \left\lfloor \frac{n_{\text{in}} + 2p - k}{s} \right\rfloor + 1$$

$n_{\text{in}}$  : number of input features

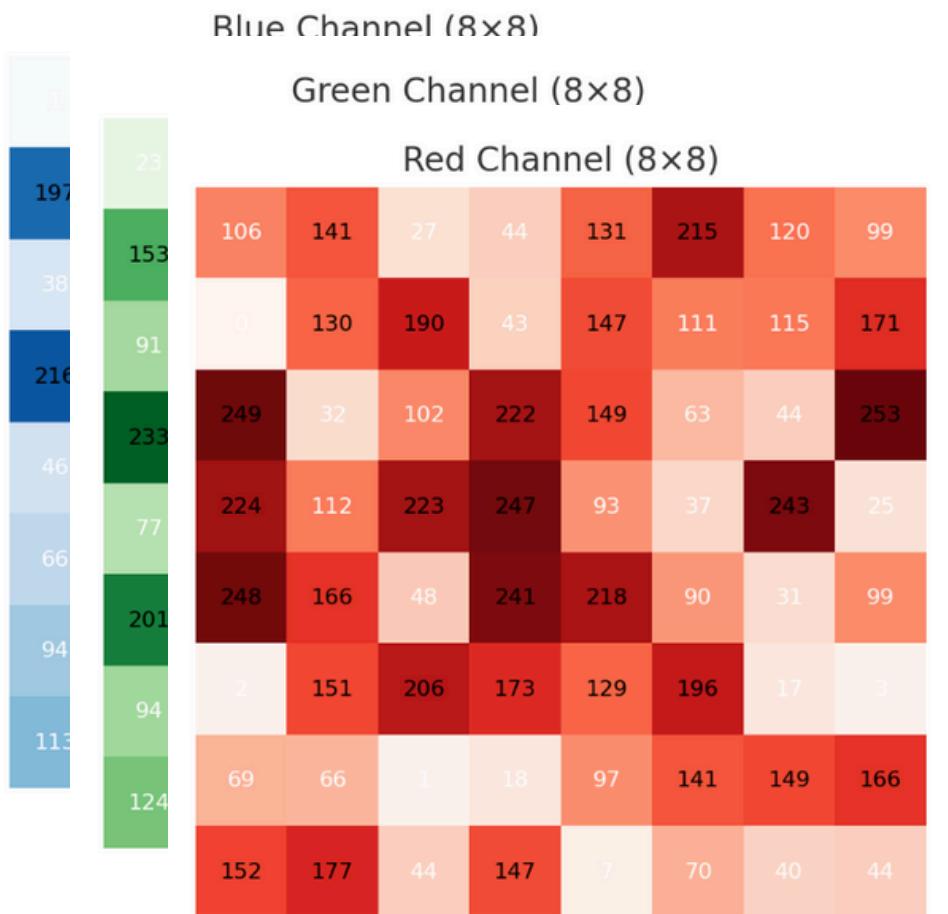
$n_{\text{out}}$  : number of output features

$k$  : convolution kernel size

$p$  : convolution padding size

$s$  : convolution stride

# Output Dimensions in Convolution



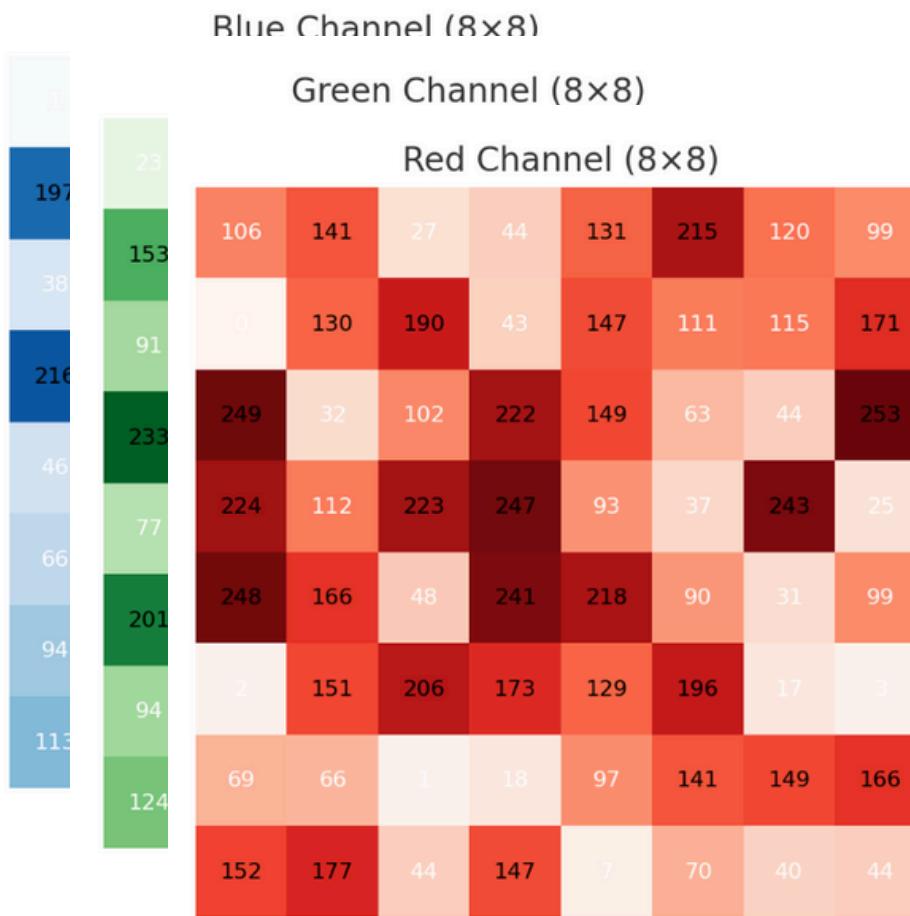
Height=8

Width=8

Depth=3

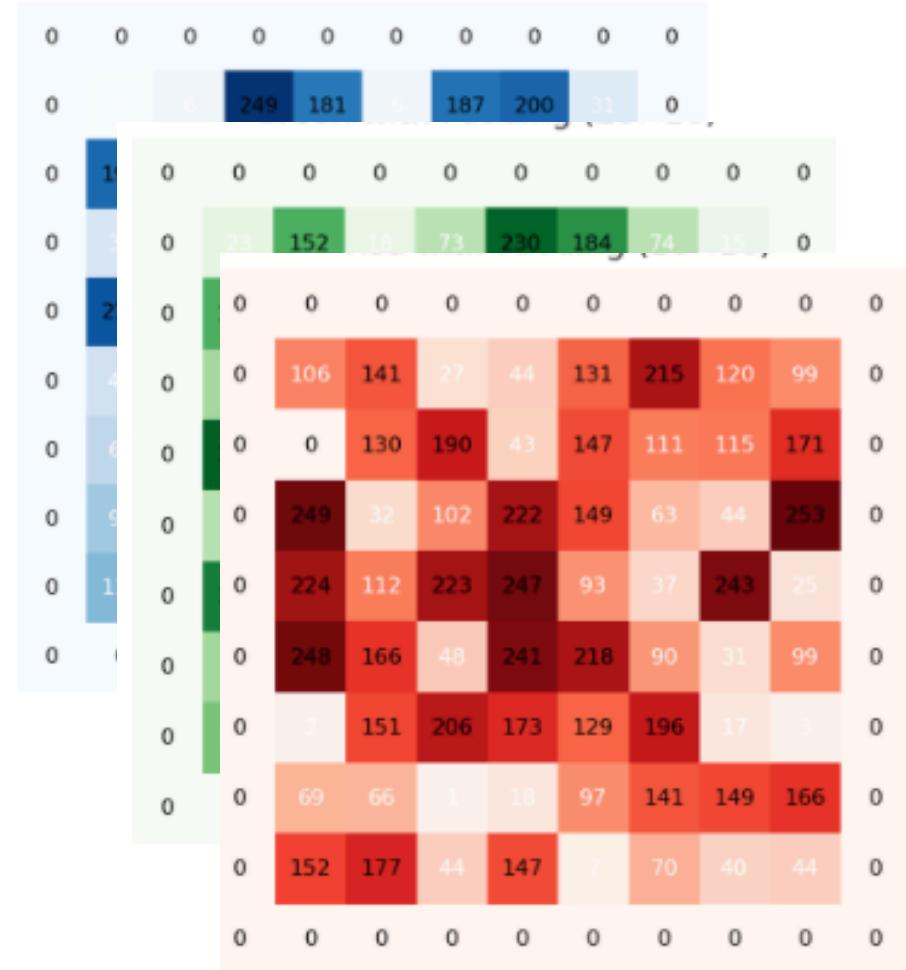
# Output Dimensions in Convolution

Input size:



$(8*8*3)$

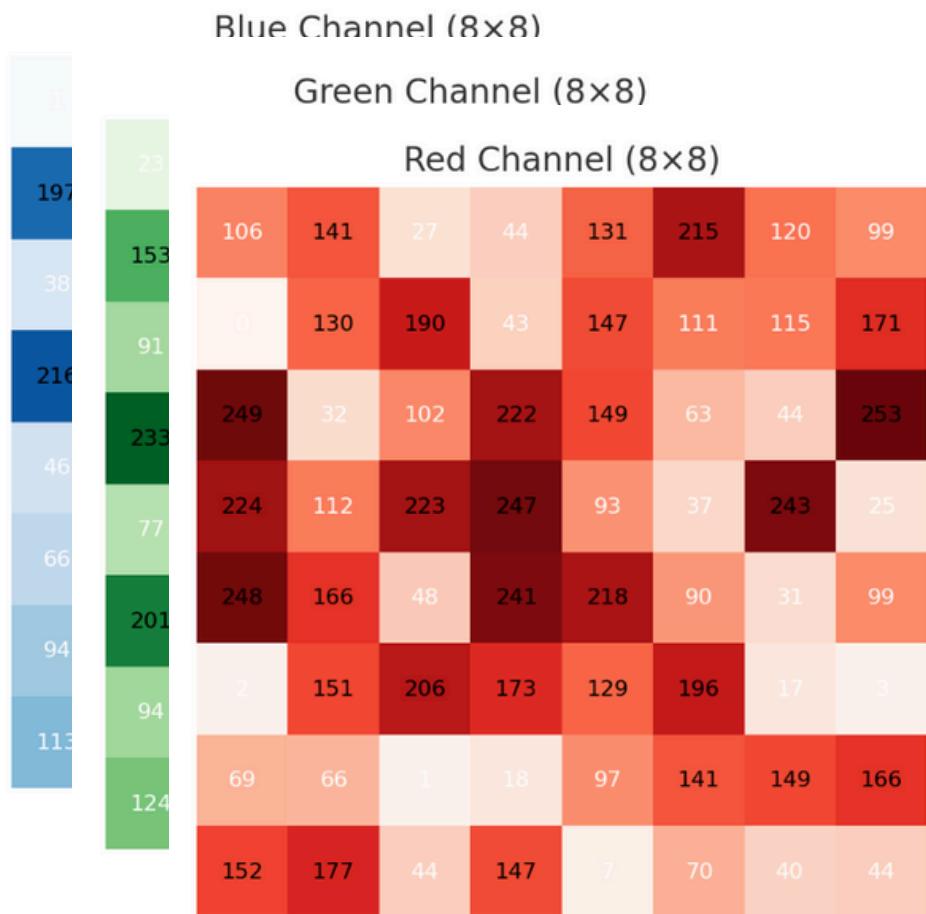
Padding:



(1)

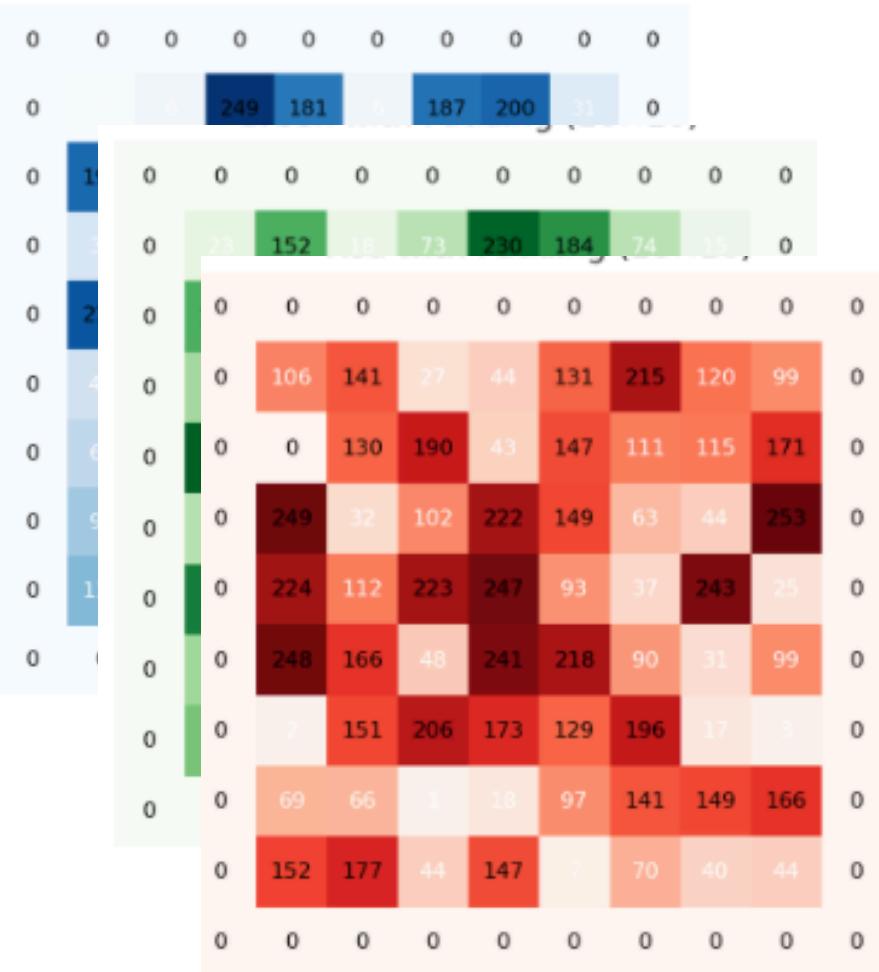
# Output Dimensions in Convolution

Input size:



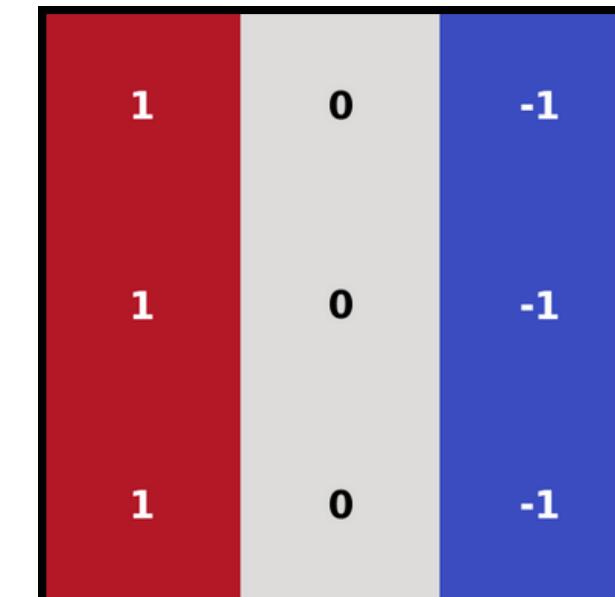
$(8*8*3)$

Padding:



(1)

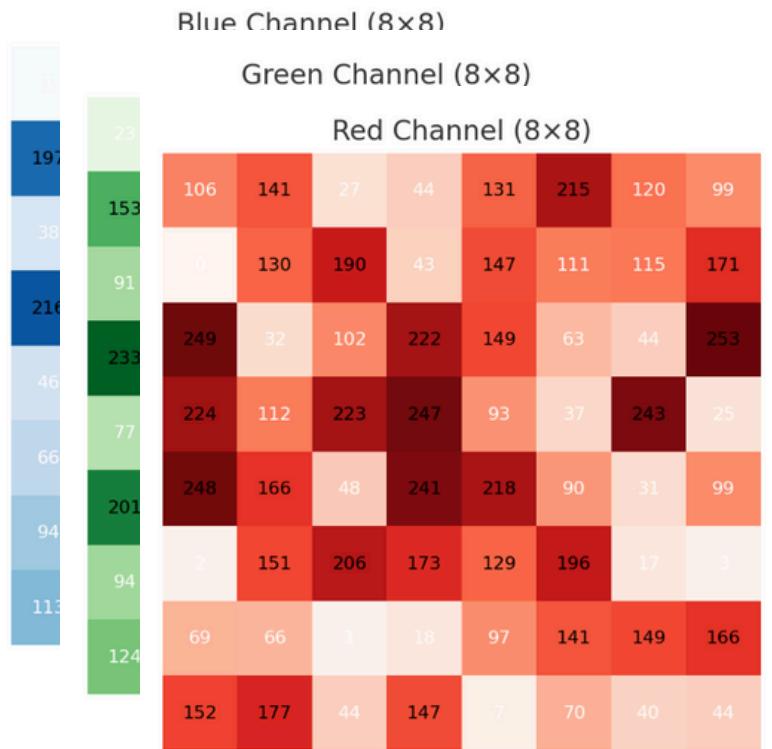
Kernel size:



$(3*3)$

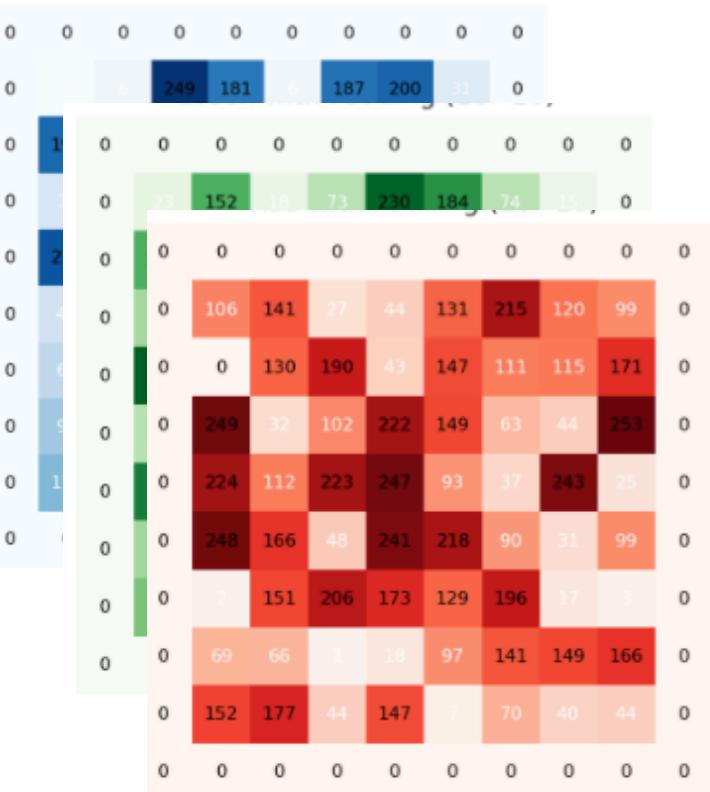
# Output Dimensions in Convolution

Input size:



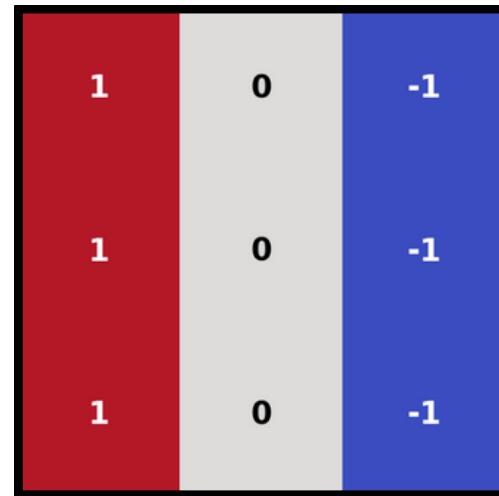
$(8*8*3)$

Padding:



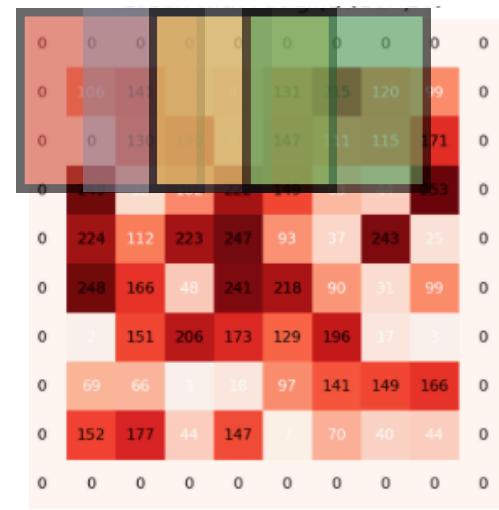
$(1)$

Kernel size:



$(3*3)$

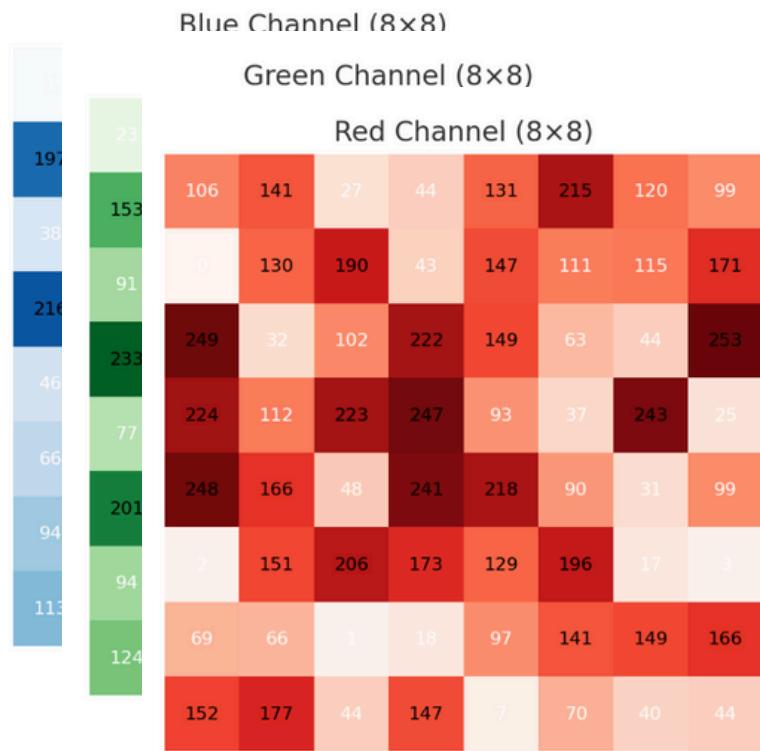
Stride:



$(1)$

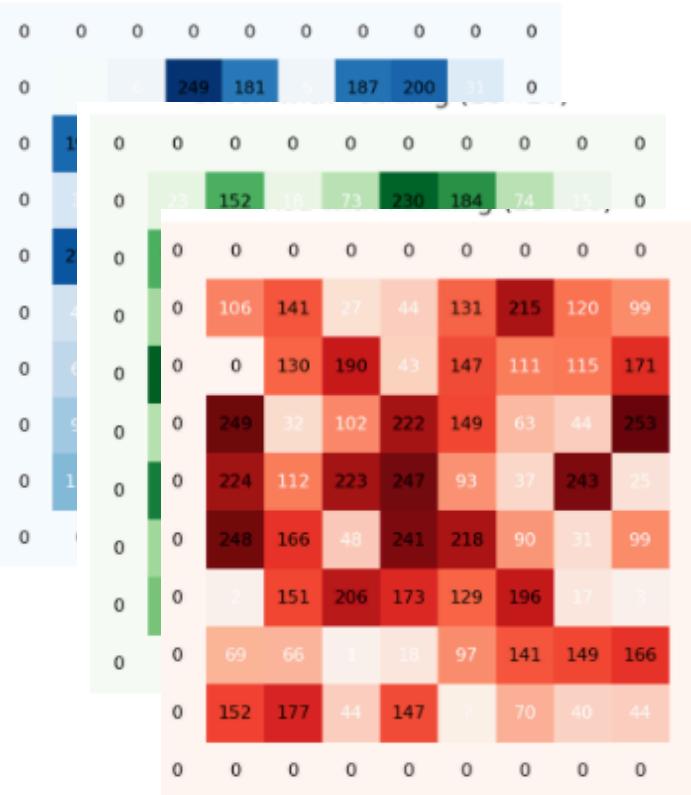
# Output Dimensions in Convolution

Input size:



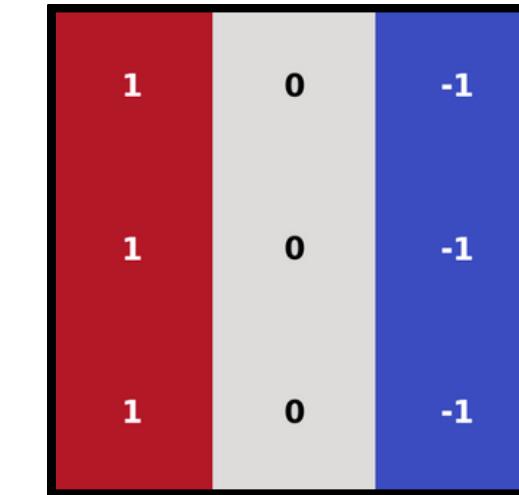
$(8*8*3)$

Padding:



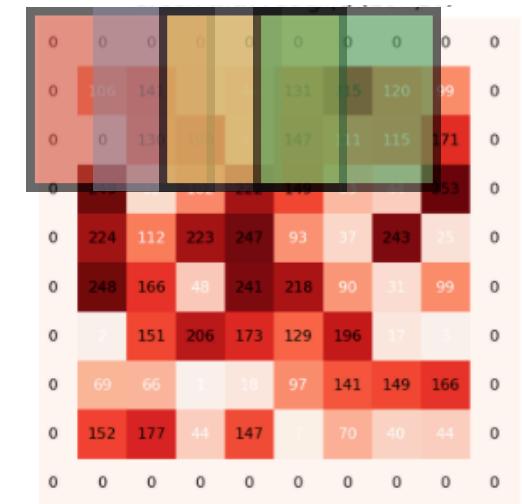
(1)

Kernel size:



$(3*3)$

Stride:



(1)

$$n_H[l] = \frac{n_H[l-1] + 2p[l] - f[l]}{s[l]} + 1$$

$$n_{out} = \frac{8 + 2(1) - 3}{1} + 1 = \frac{7}{1} + 1 = 8$$

## Output Dimensions in Convolution

Input image size:  $10 \times 10 \times 3$

Kernel size:  $5 \times 5$

Padding: 2

Stride: 2

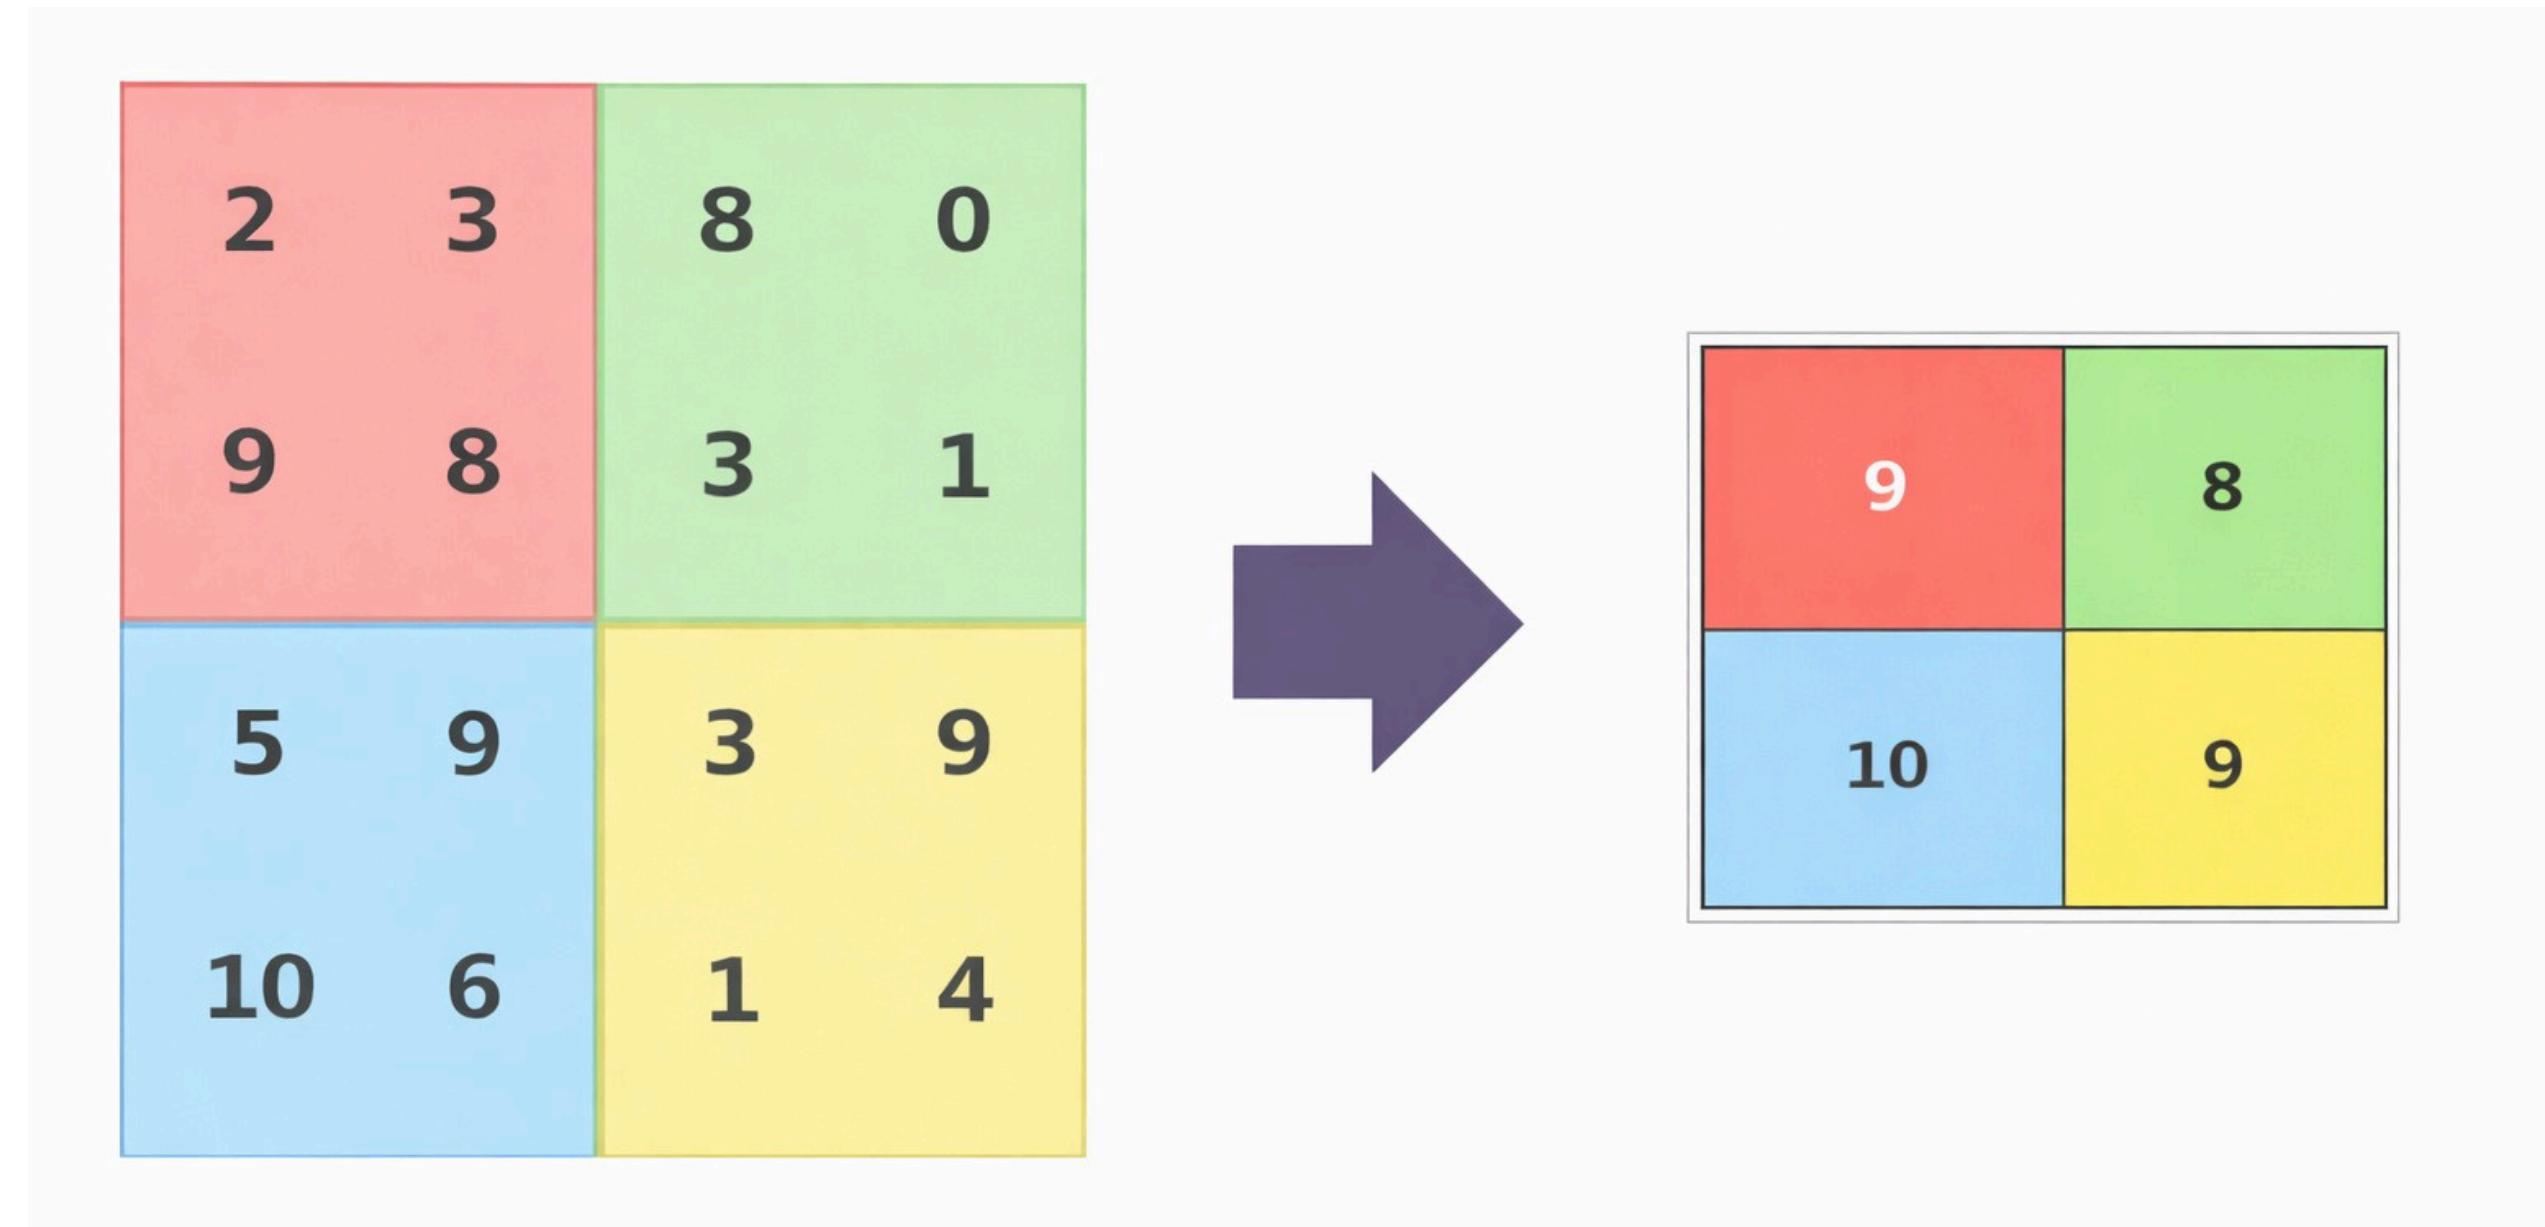
Number of filters: 16

$$n_{\text{out}} = \left\lfloor \frac{n_{\text{in}} + 2p - k}{s} \right\rfloor + 1$$

## Pooling Layer

- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

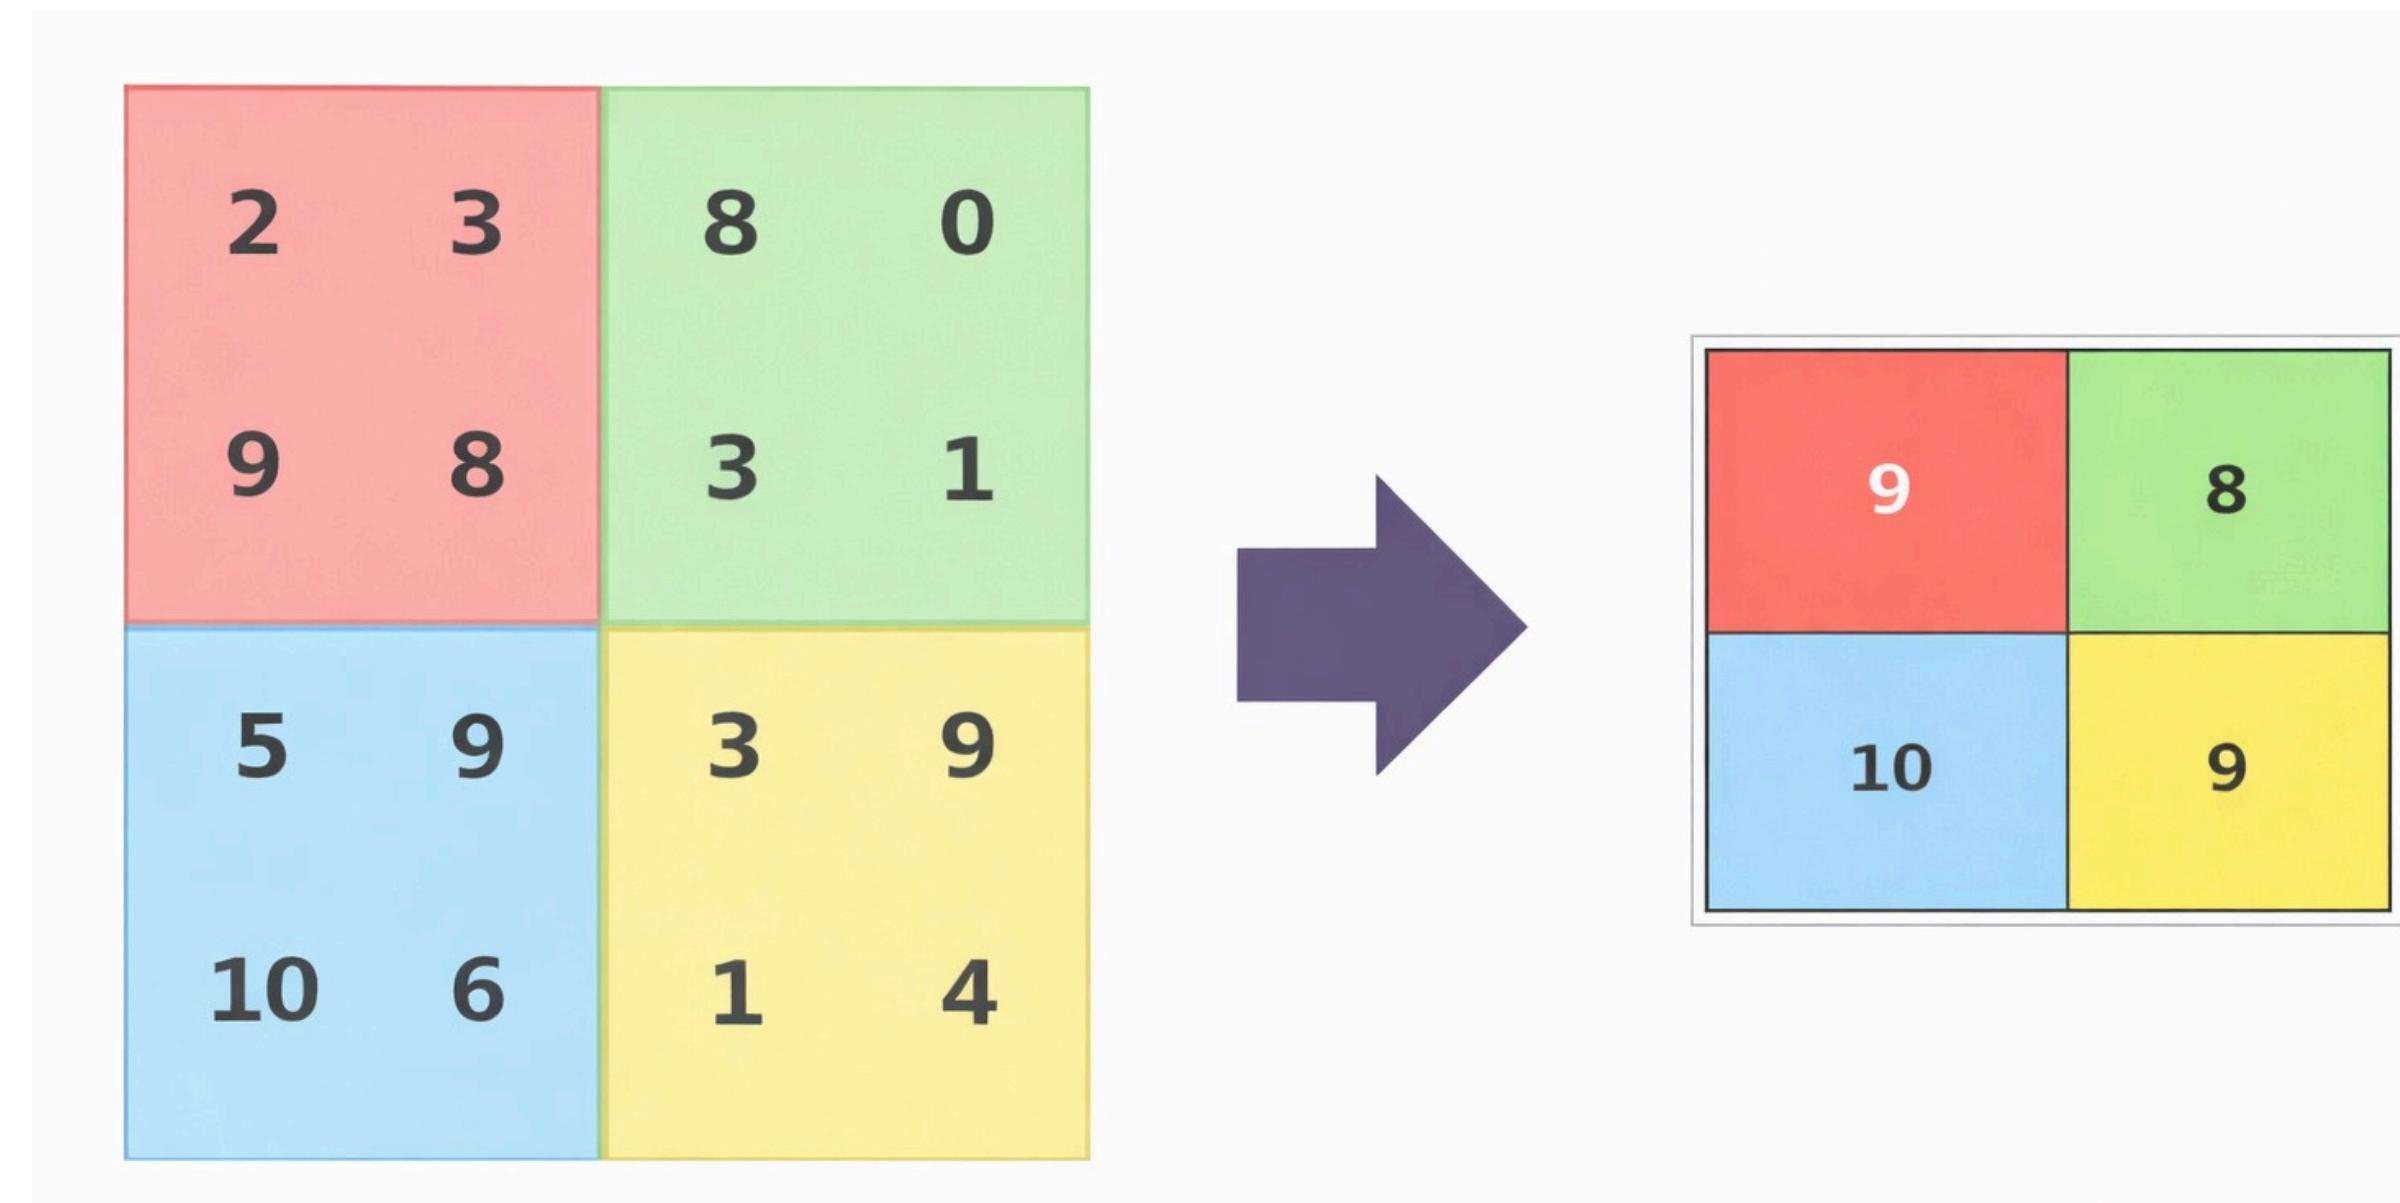
- Max Pooling



## Pooling Layer

- › Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

- › Max Pooling

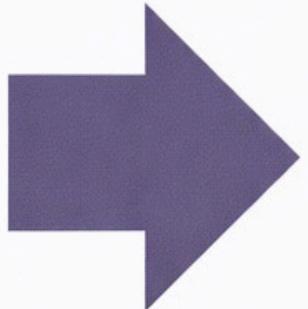


## Pooling Layer

- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

- Average Pooling

2	3	8	0
9	8	3	1
5	9	3	9
10	6	1	4

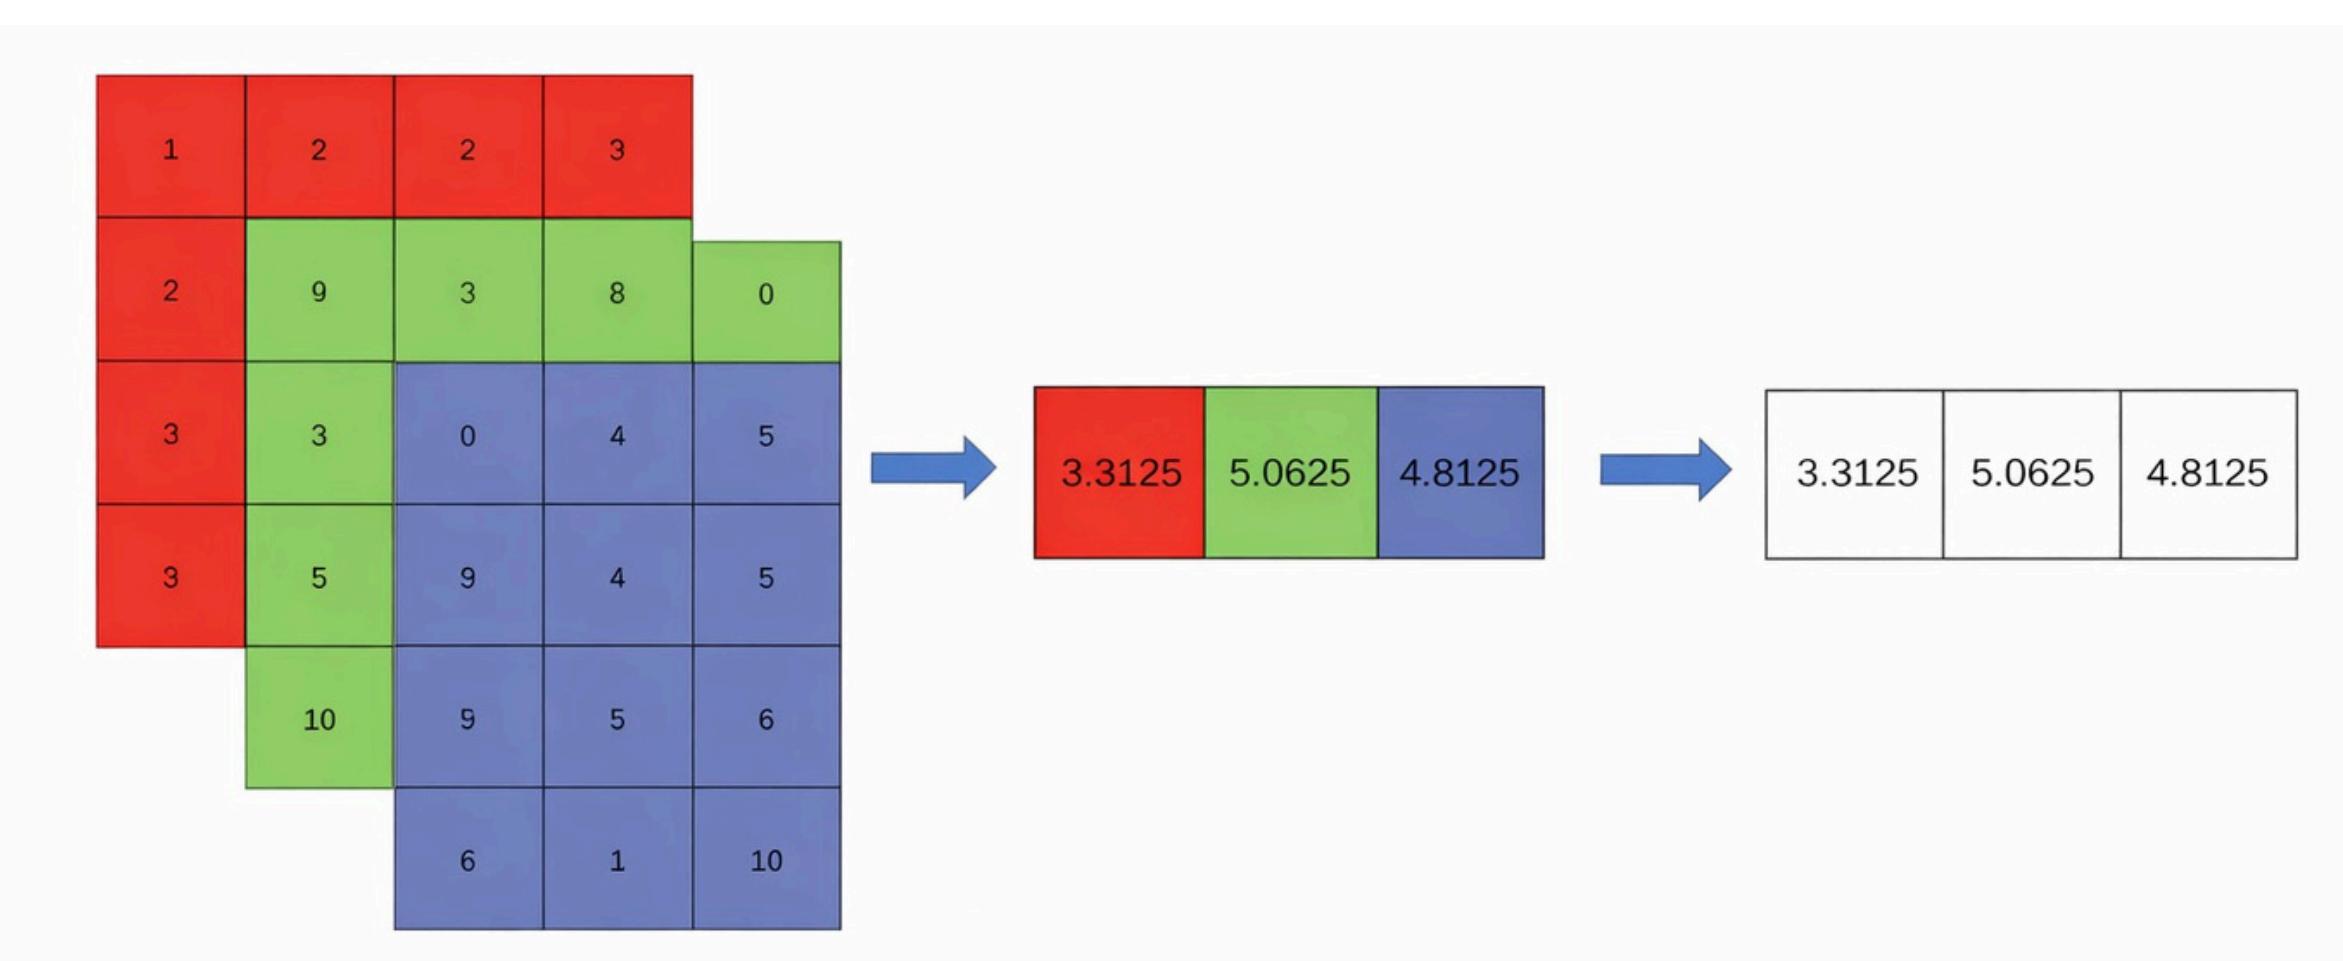


5.5	3
7.5	4.25

# Pooling Layer

- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

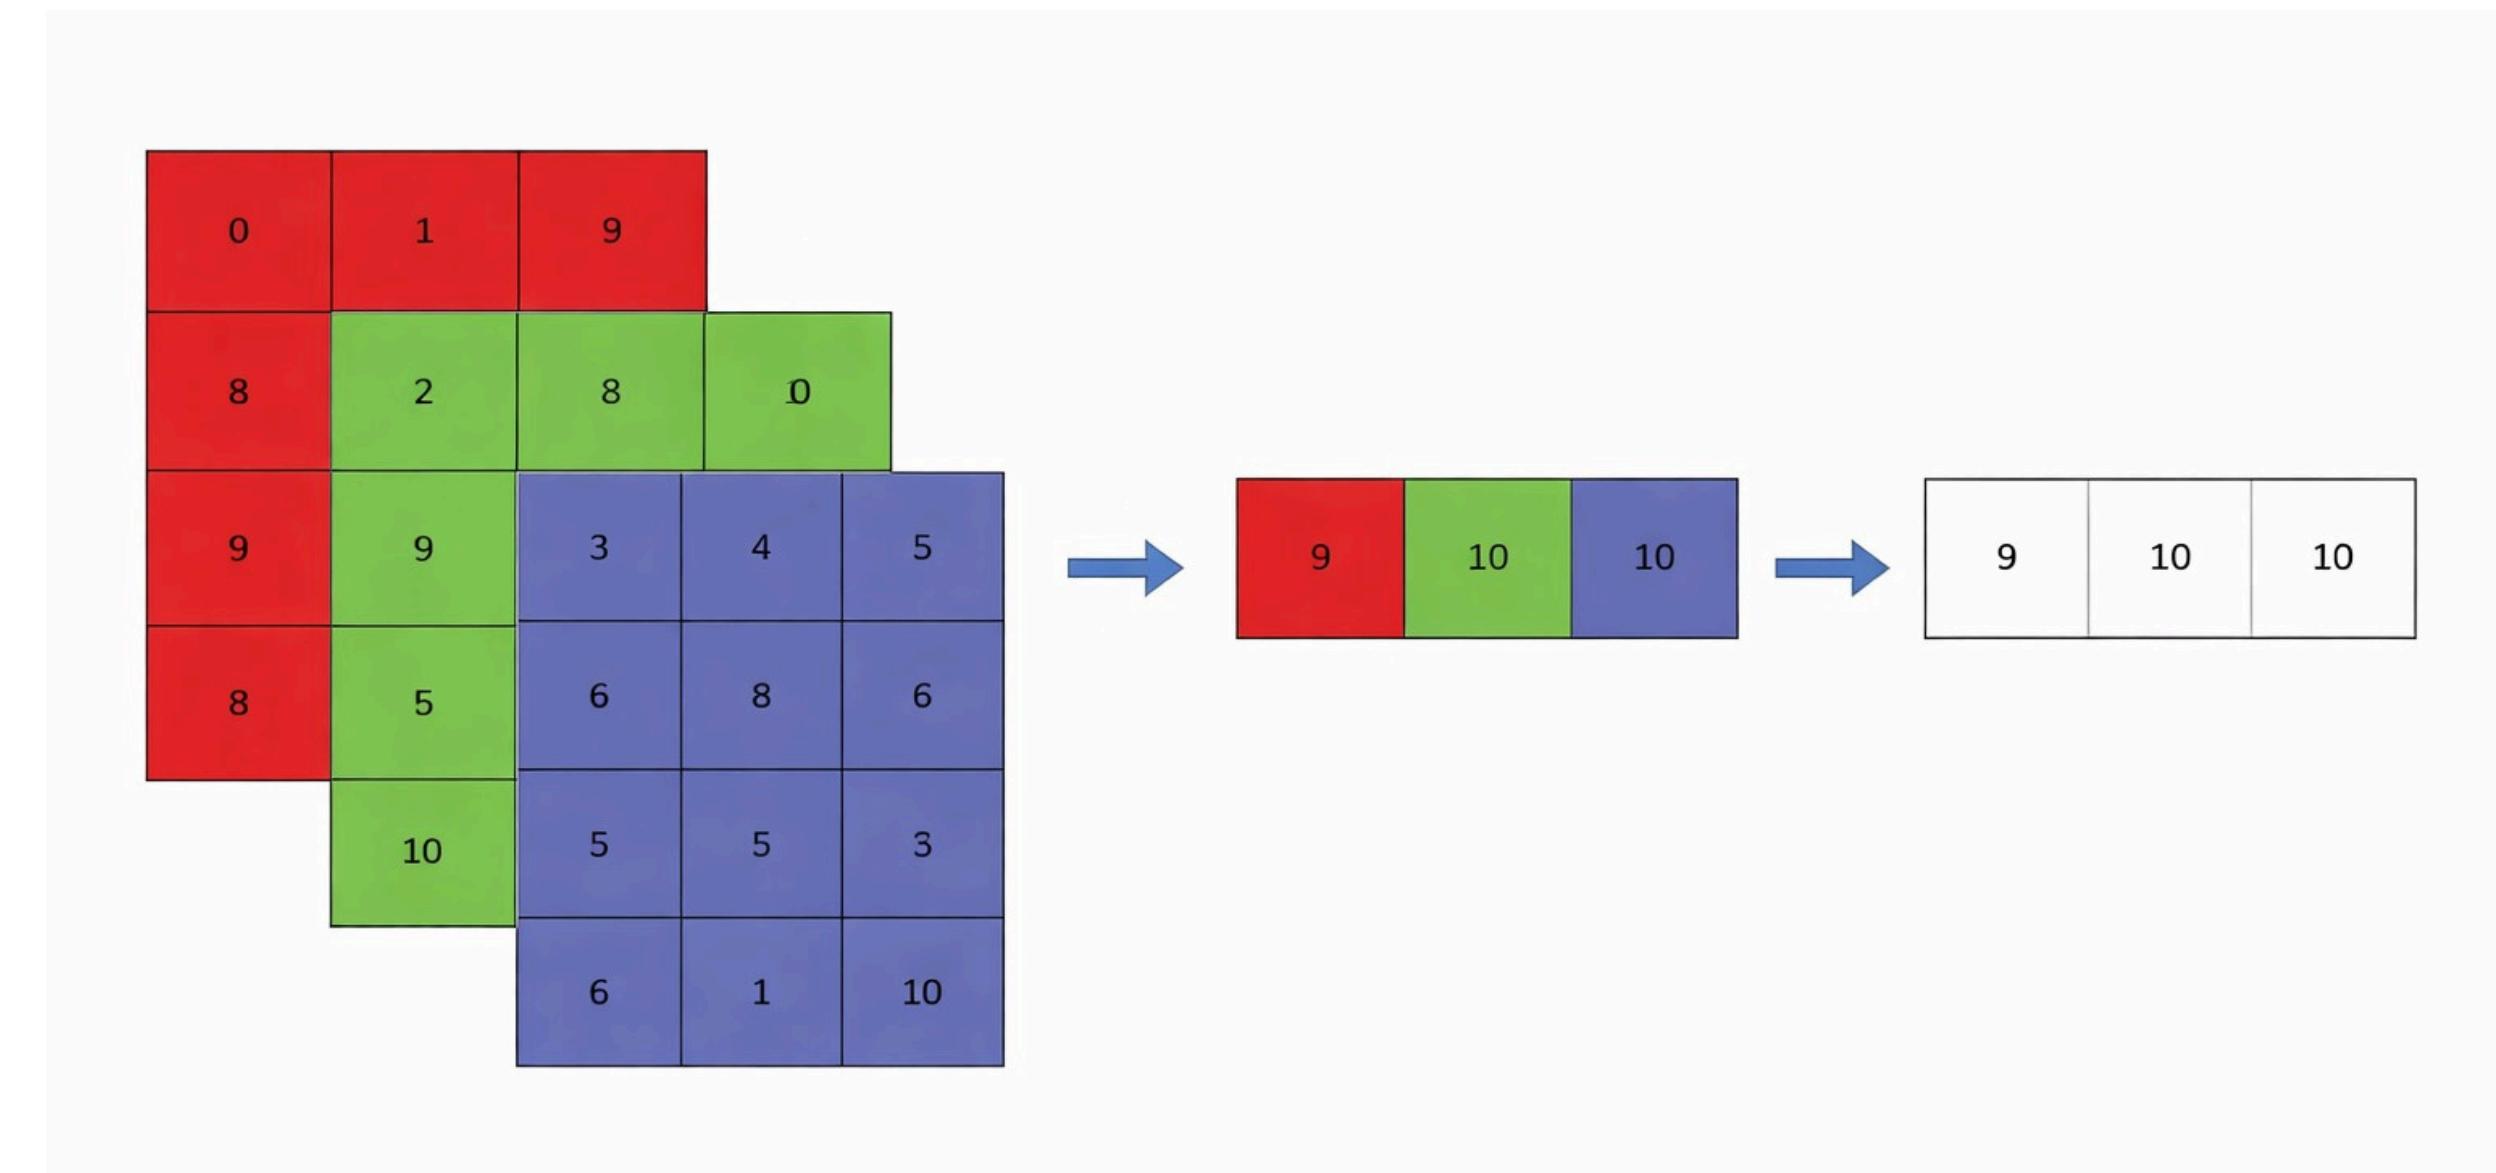
## ➤ Global Average Pooling



# Pooling Layer

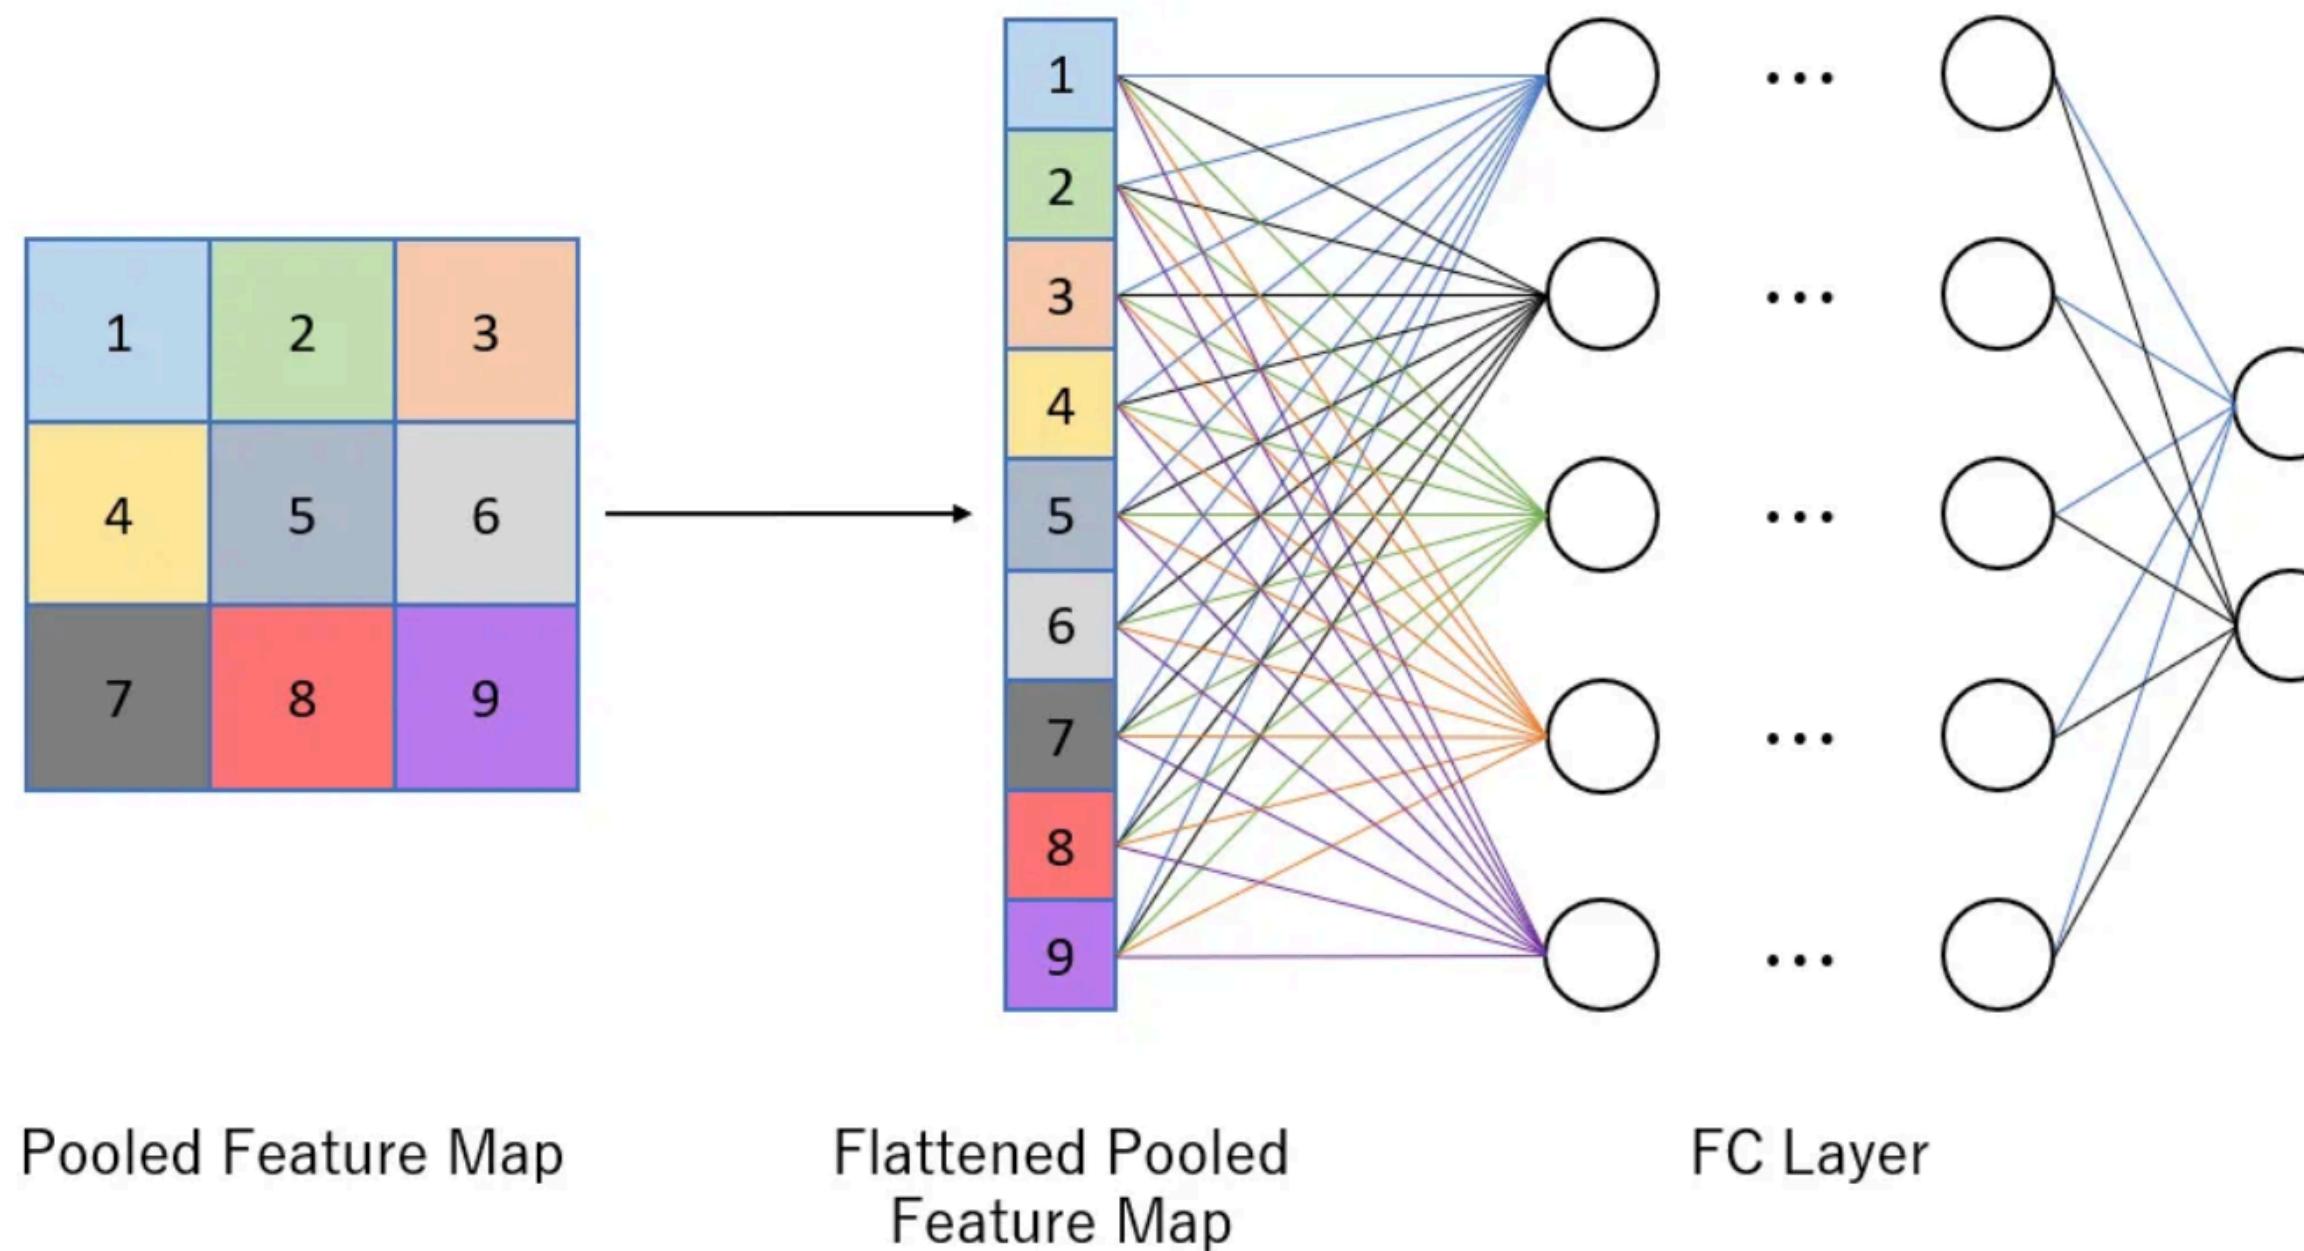
- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

## ➤ Global Max Pooling



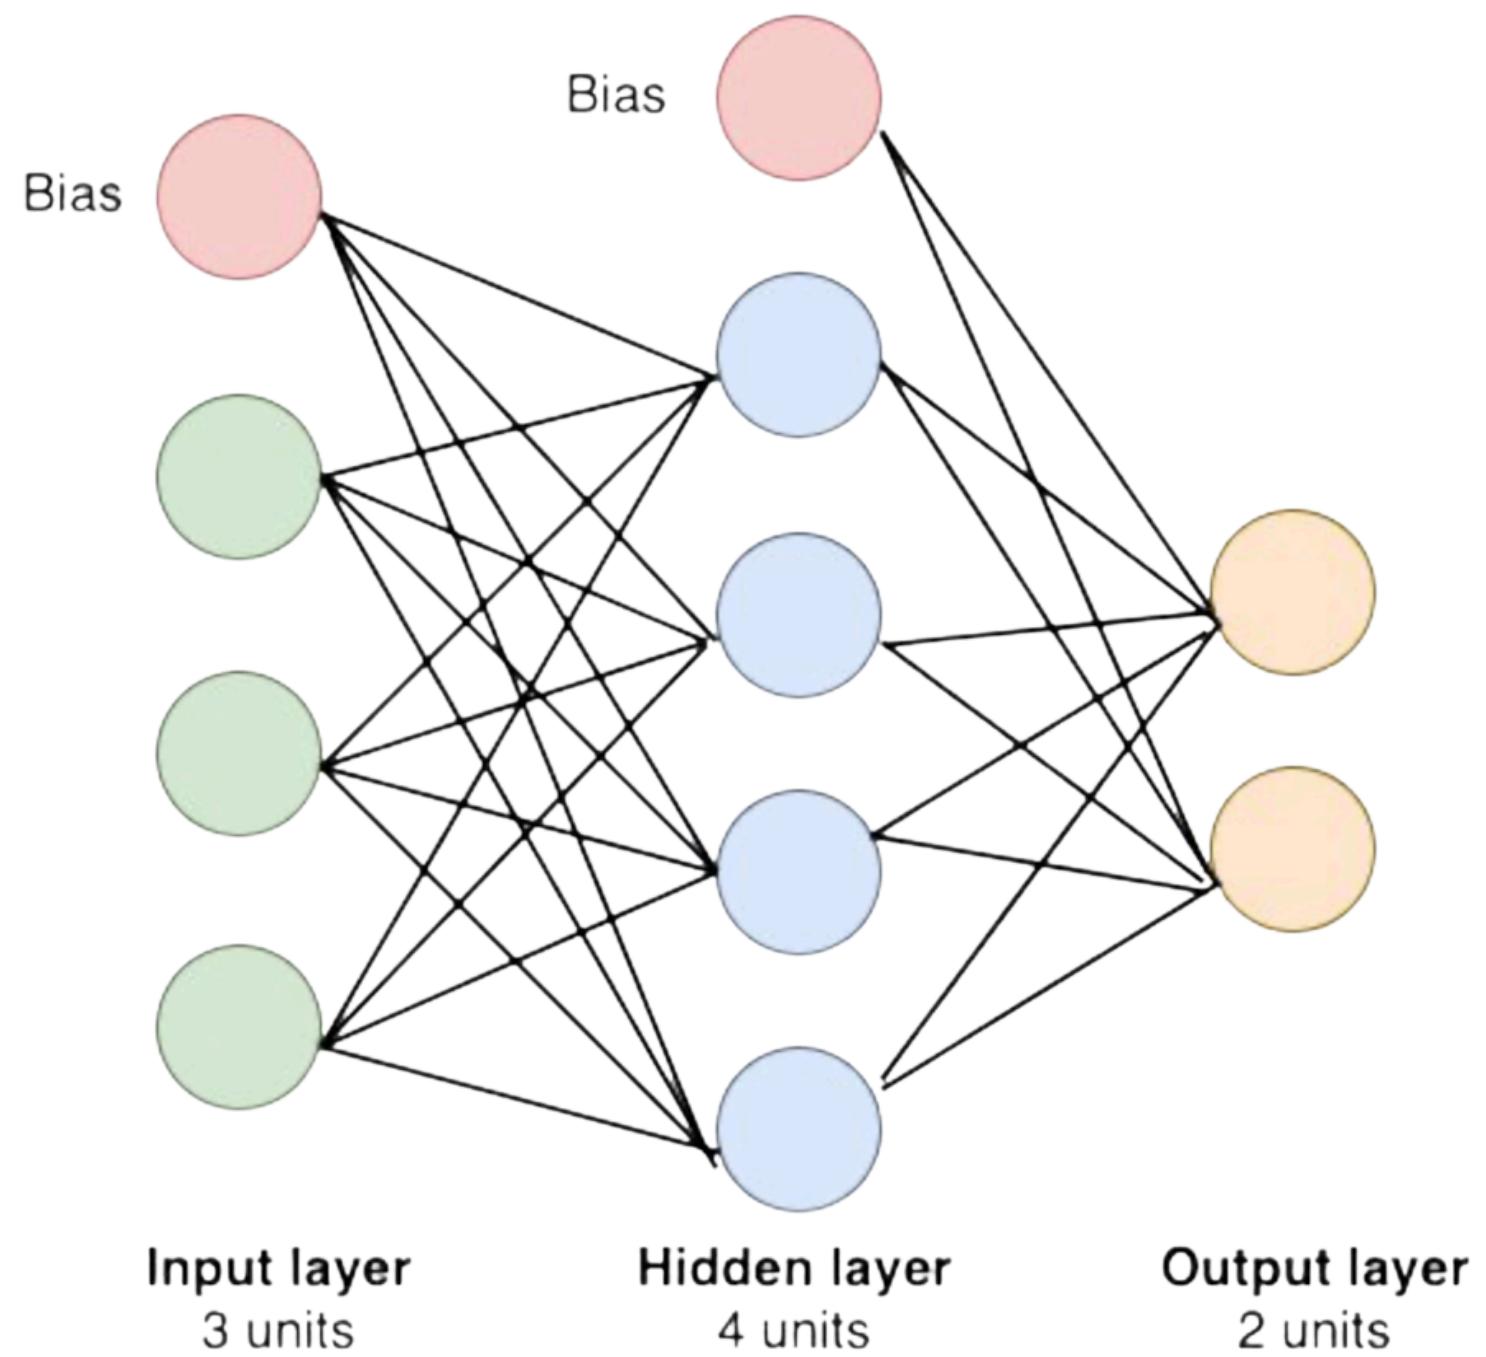
# Dense Layer

- It combines the features learned by earlier convolutional and pooling layers to perform high-level reasoning or classification.



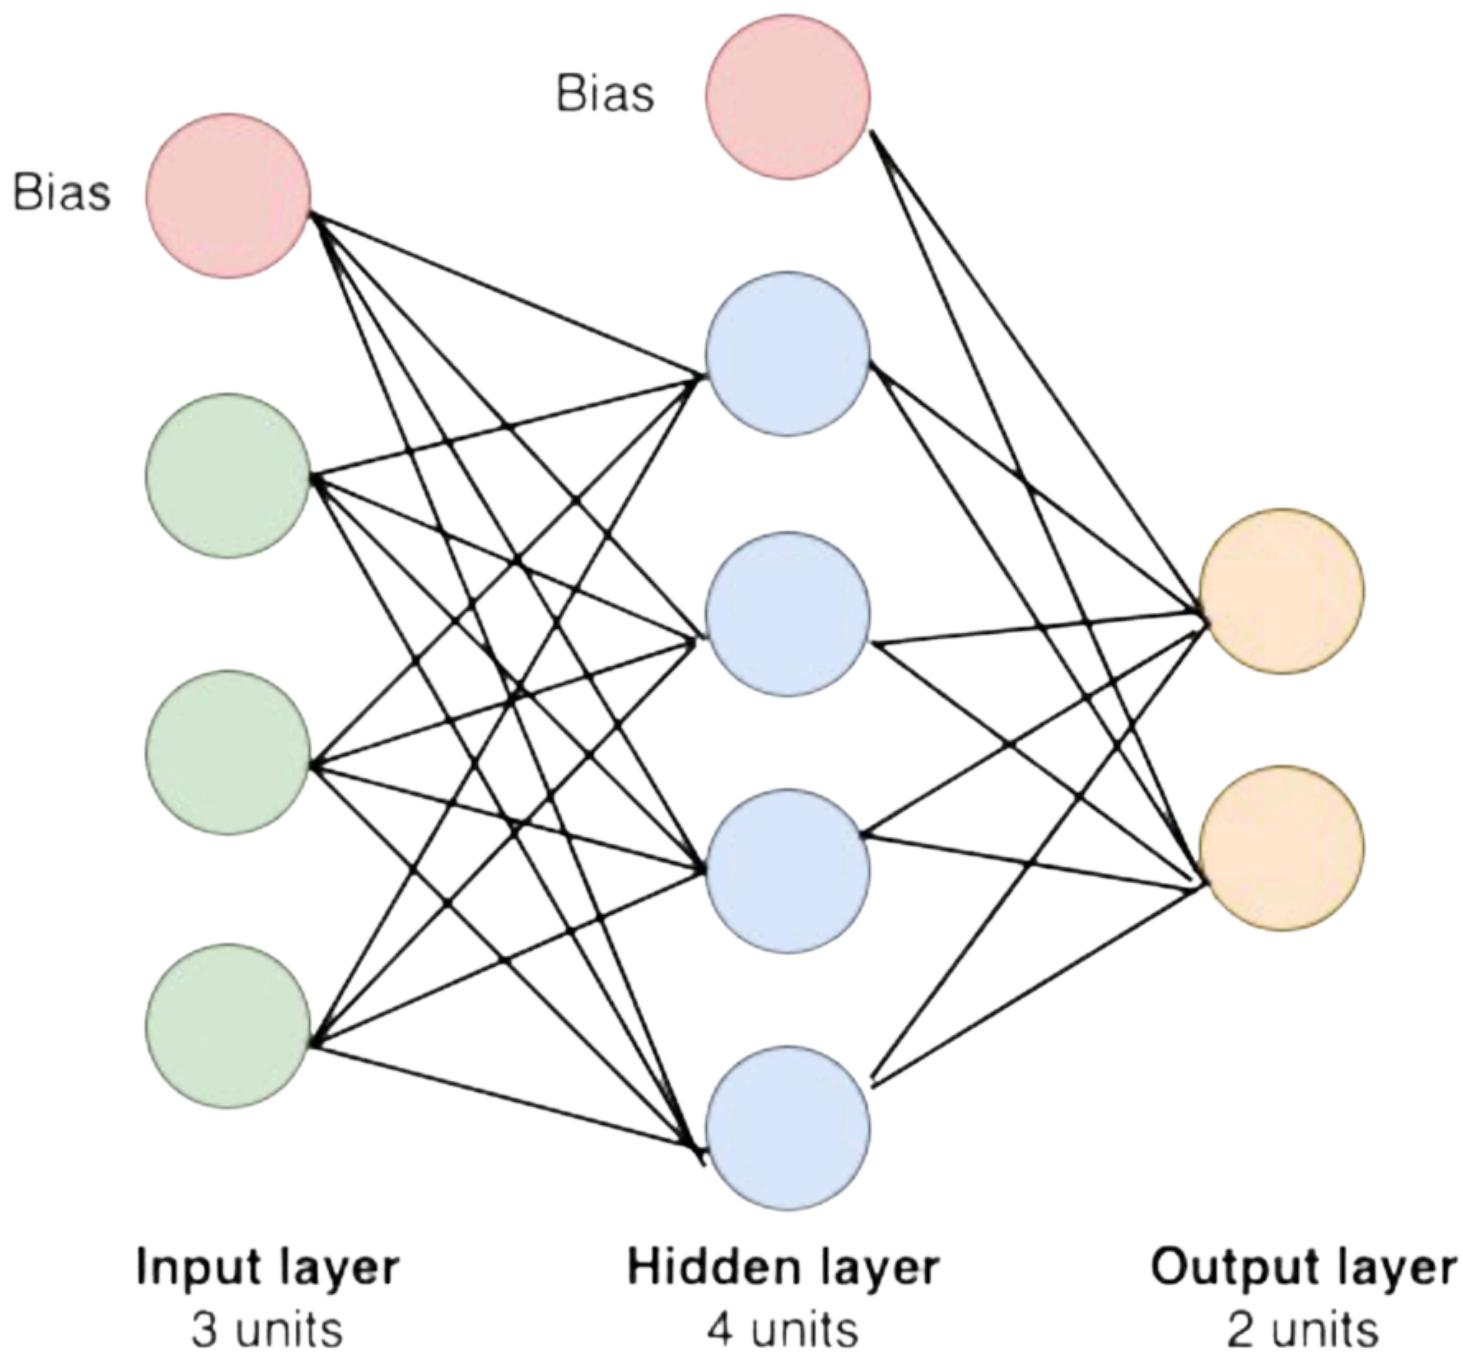
## Learnable Parameters in NN

Learnable Parameters in a Neural Network:  $= (n_{\text{in}} \times n_{\text{out}}) + n_{\text{out}}$



## Learnable Parameters in NN

Learnable Parameters in a Neural Network:  $= (n_{\text{in}} \times n_{\text{out}}) + n_{\text{out}}$



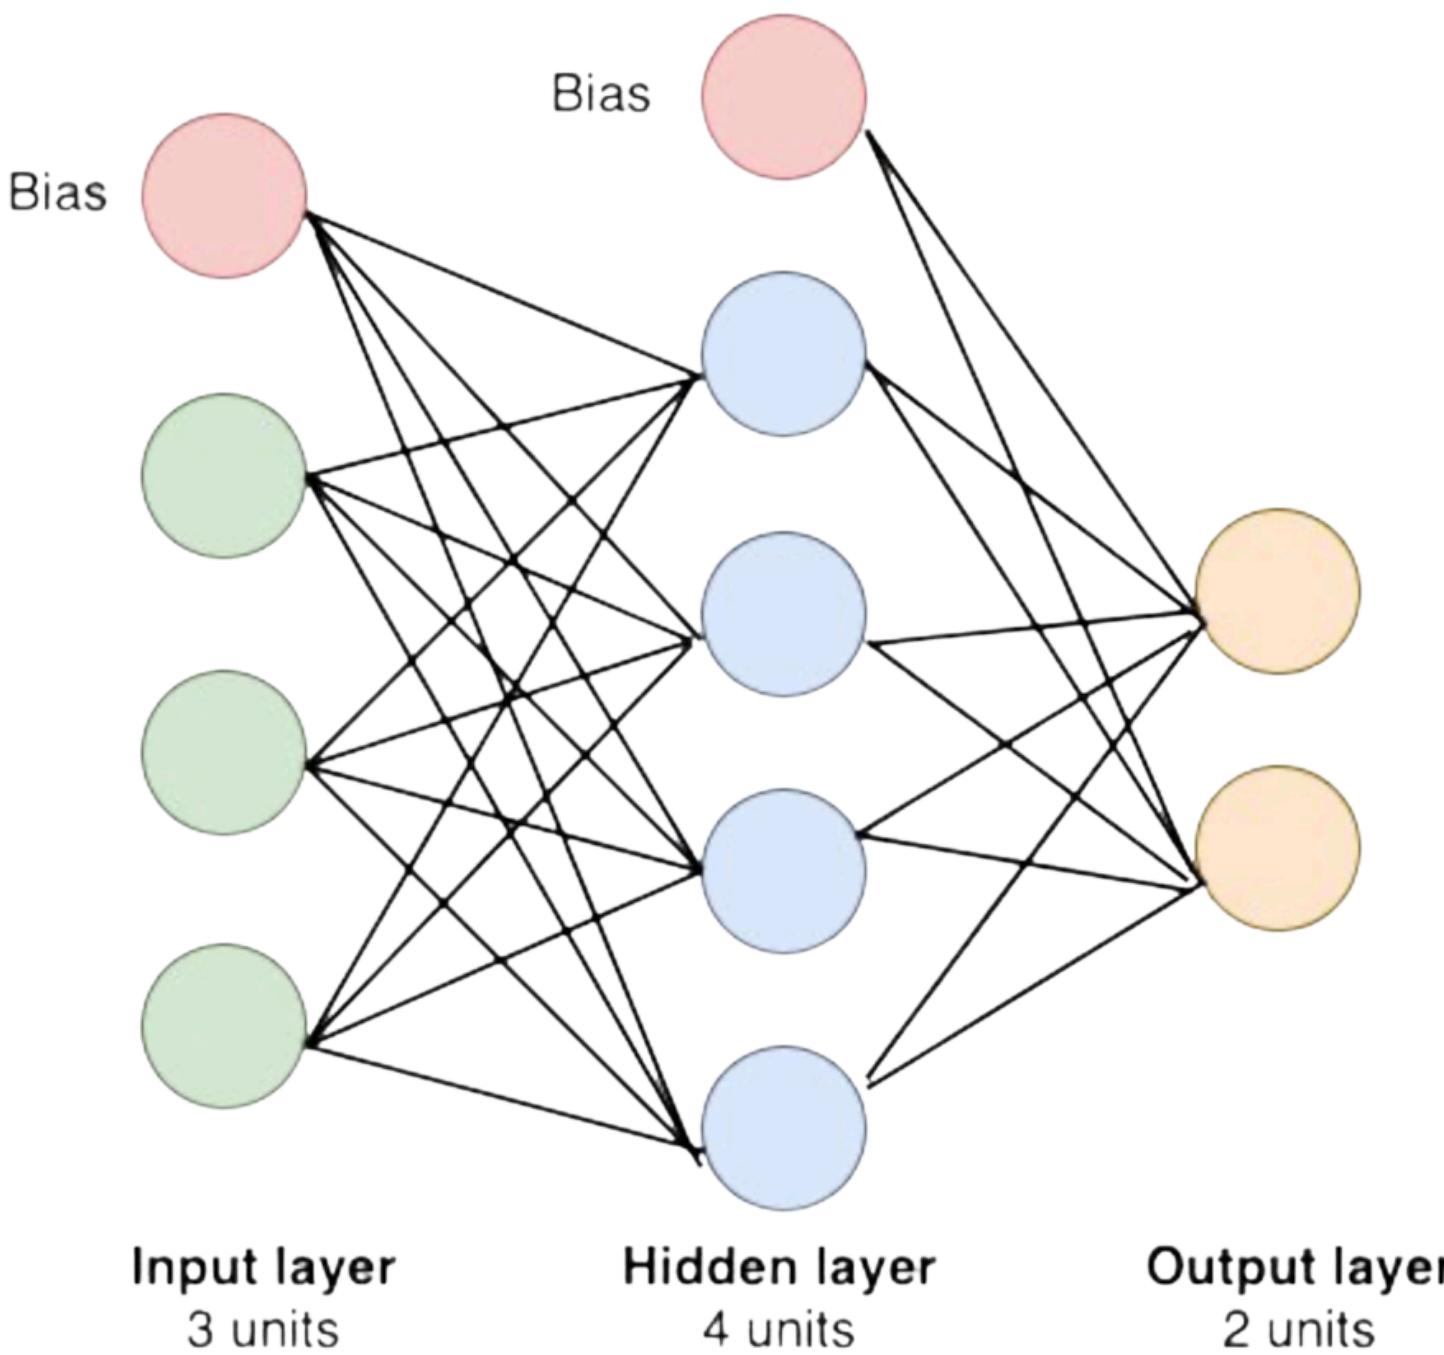
$$\text{Weights} = 3 \times 4 = 12$$

$$\text{Biases} = 4$$

$$\text{Total}_{\text{IH}} = 12 + 4 = 16$$

## Learnable Parameters in NN

Learnable Parameters in a Neural Network:  $= (n_{\text{in}} \times n_{\text{out}}) + n_{\text{out}}$



$$\text{Weights} = 3 \times 4 = 12$$

$$\text{Biases} = 4$$

$$\text{Total}_{IH} = 12 + 4 = 16$$

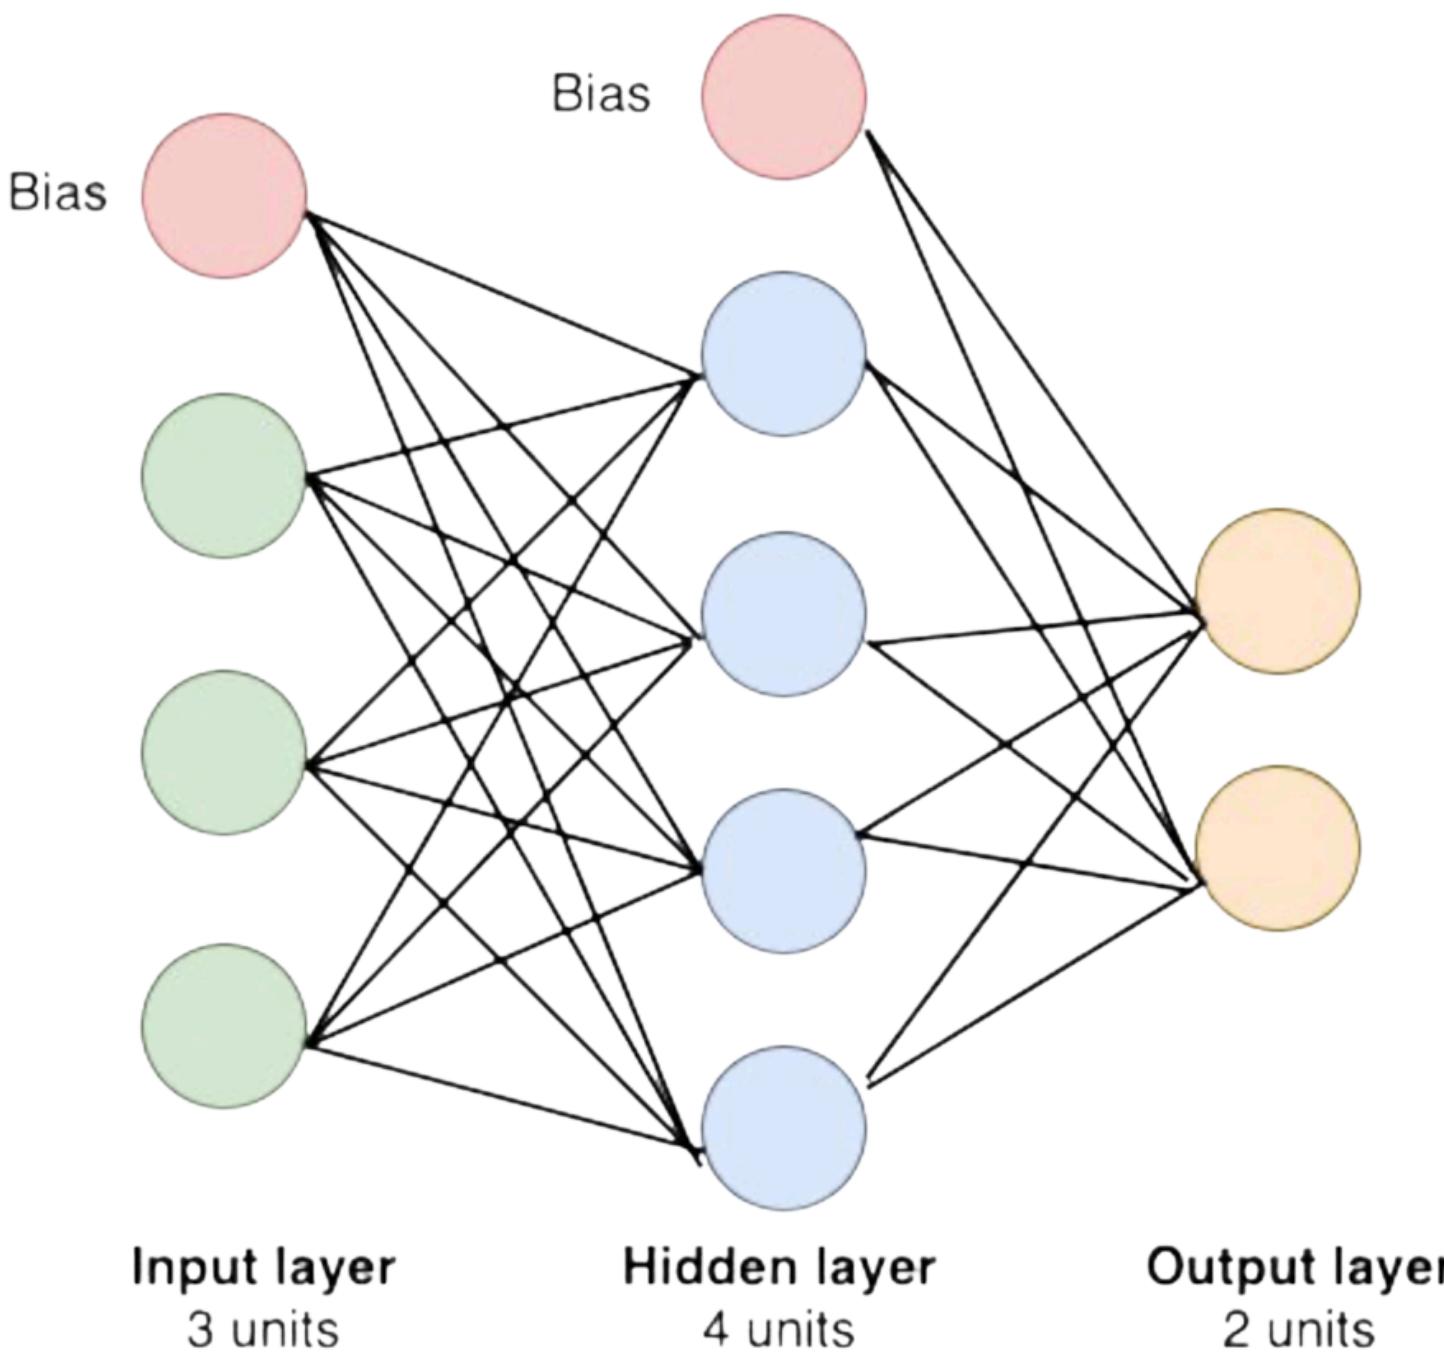
$$\text{Weights} = 4 \times 2 = 8$$

$$\text{Biases} = 2$$

$$\text{Total}_{HO} = 8 + 2 = 10$$

## Learnable Parameters in NN

Learnable Parameters in a Neural Network:  $= (n_{\text{in}} \times n_{\text{out}}) + n_{\text{out}}$



$$\text{Weights} = 3 \times 4 = 12$$

$$\text{Biases} = 4$$

$$\text{Total}_{IH} = 12 + 4 = 16$$

$$\text{Weights} = 4 \times 2 = 8$$

$$\text{Biases} = 2$$

$$\text{Total}_{HO} = 8 + 2 = 10$$

$$\text{Total} = 16 + 10 = 26$$

# Learnable Parameters in CNN

Learnable Parameters in a CNN:

$$\textbf{Params} = (K_h \times K_w \times C_{in} + 1) \times C_{out}$$

A convolution layer with:

- Kernel size:  $3 \times 3$
- Input channels: 3 (RGB)
- Output channels: 16 filters

# Learnable Parameters in CNN

Learnable Parameters in a CNN:

$$\textbf{Params} = (K_h \times K_w \times C_{in} + 1) \times C_{out}$$

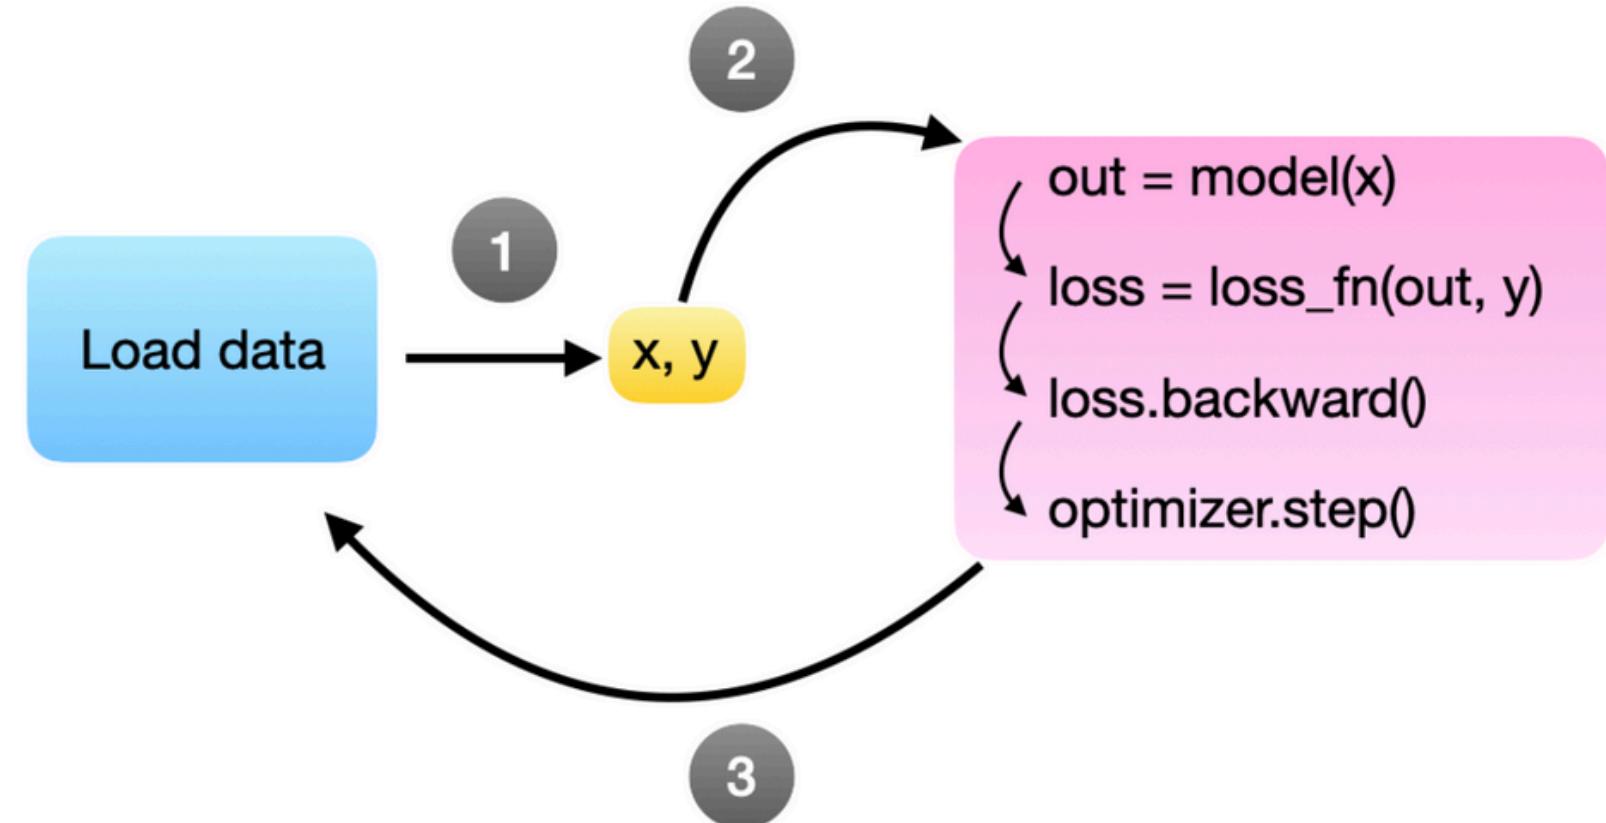
A convolution layer with:

- Kernel size:  $3 \times 3$   $(3 \times 3 \times 3 + 1) \times 16$
- Input channels: 3 (RGB)  $28 \times 16 = 448$
- Output channels: 16 filters

# Efficient Data Loading

- Training speed depends not only on the model
- GPU should never wait for data

```
for epoch in range(num_epochs):  
    inputs = []  
    for f in filename_minibatch:  
        img = Image.open(os.path.join(img_dir, f))  
        img_tensor = o_tensor(img).to('cuda')  
        inputs.append(img_tensor)  
    inputs = torch.cat(inputs)  
    logits = model(inputs)  
    loss.backward()  
    optimizer.step()
```

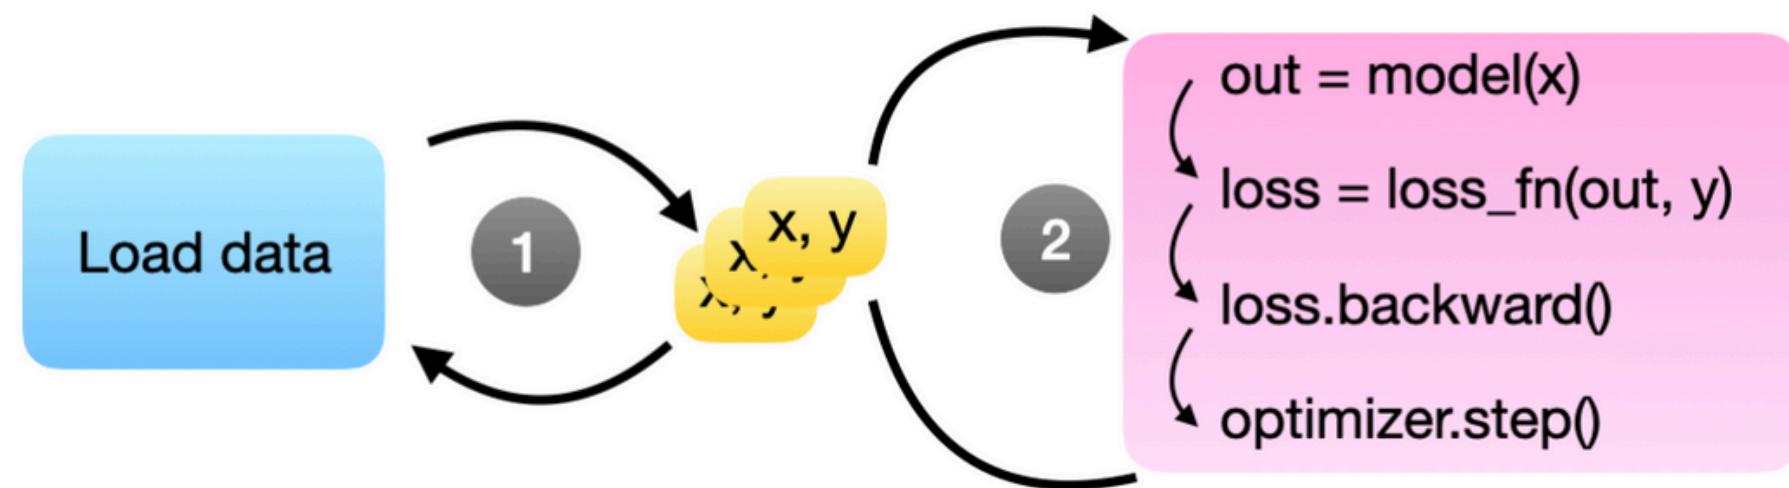


# Efficient Data Loading

- Load next mini-batch in background
- Supports multiple worker processes
- Automatic shuffling, batching, and collation

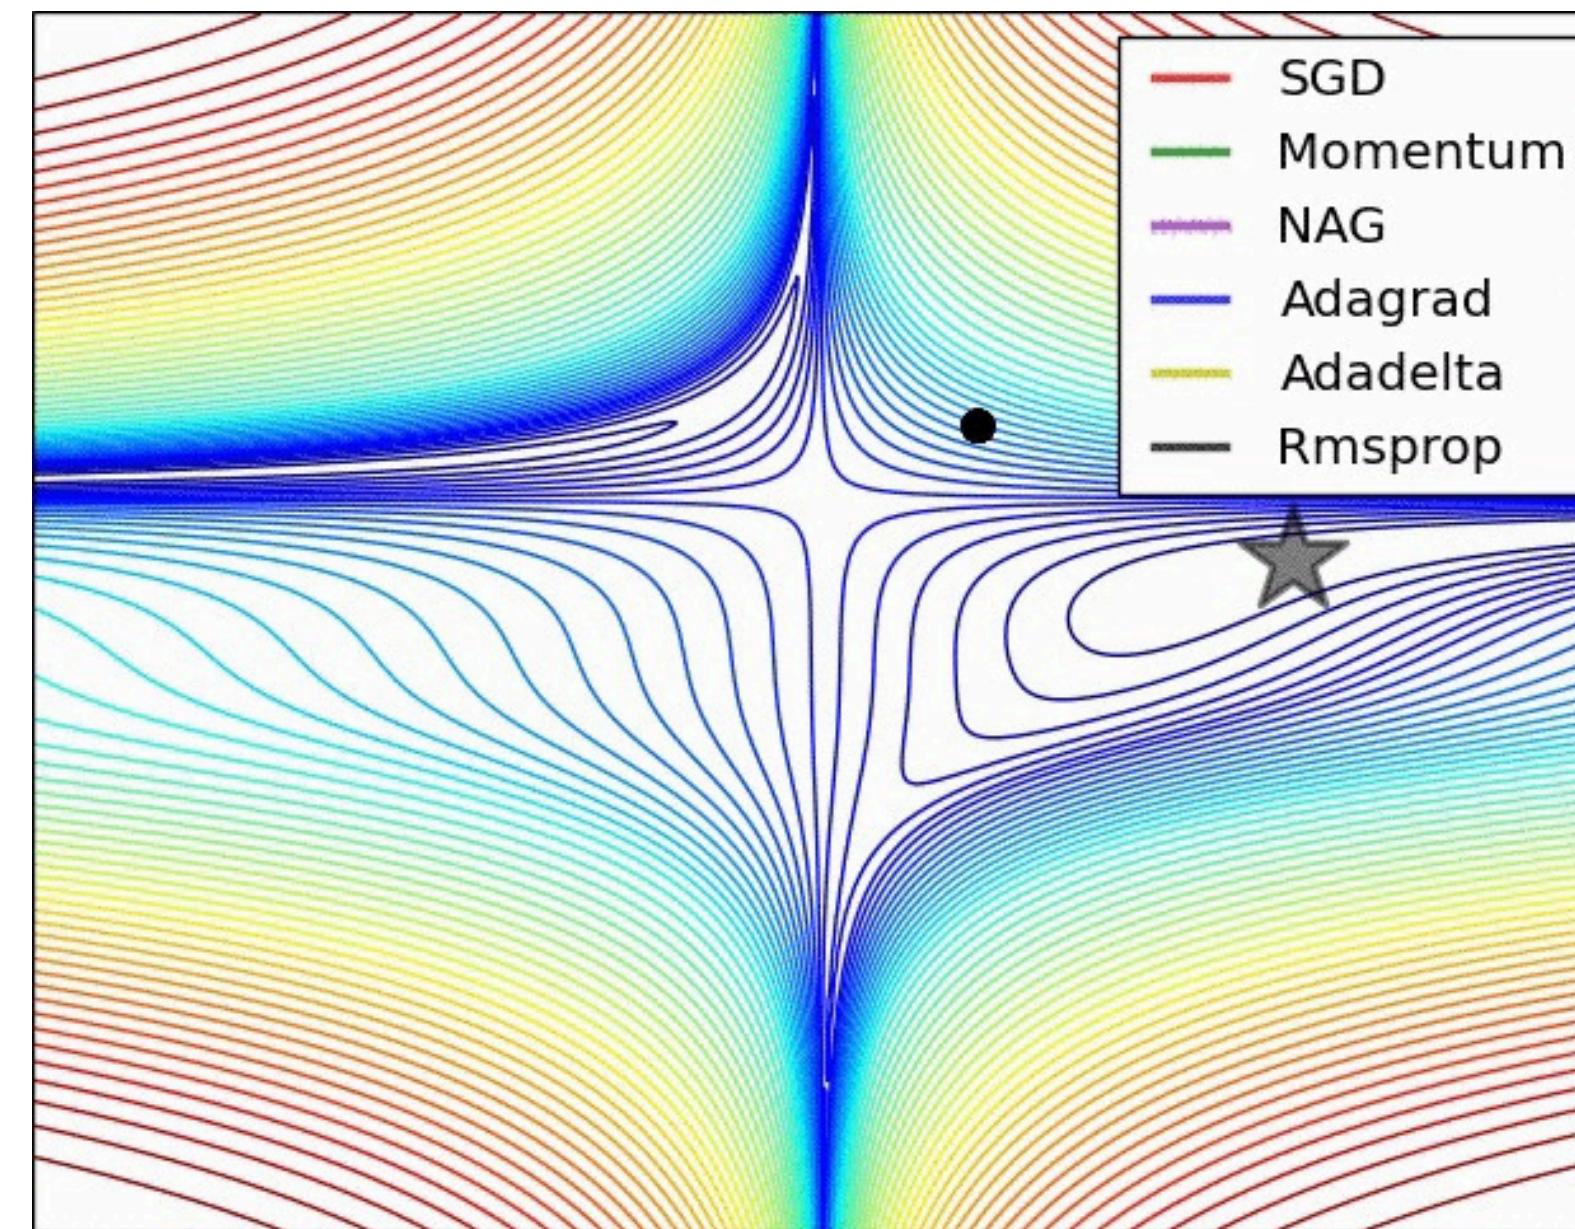
```
from torch.utils.data import Dataset, DataLoader

train_loader = DataLoader(dataset, batch_size=64,
                           shuffle=True, num_workers=4)
```



# Optimization

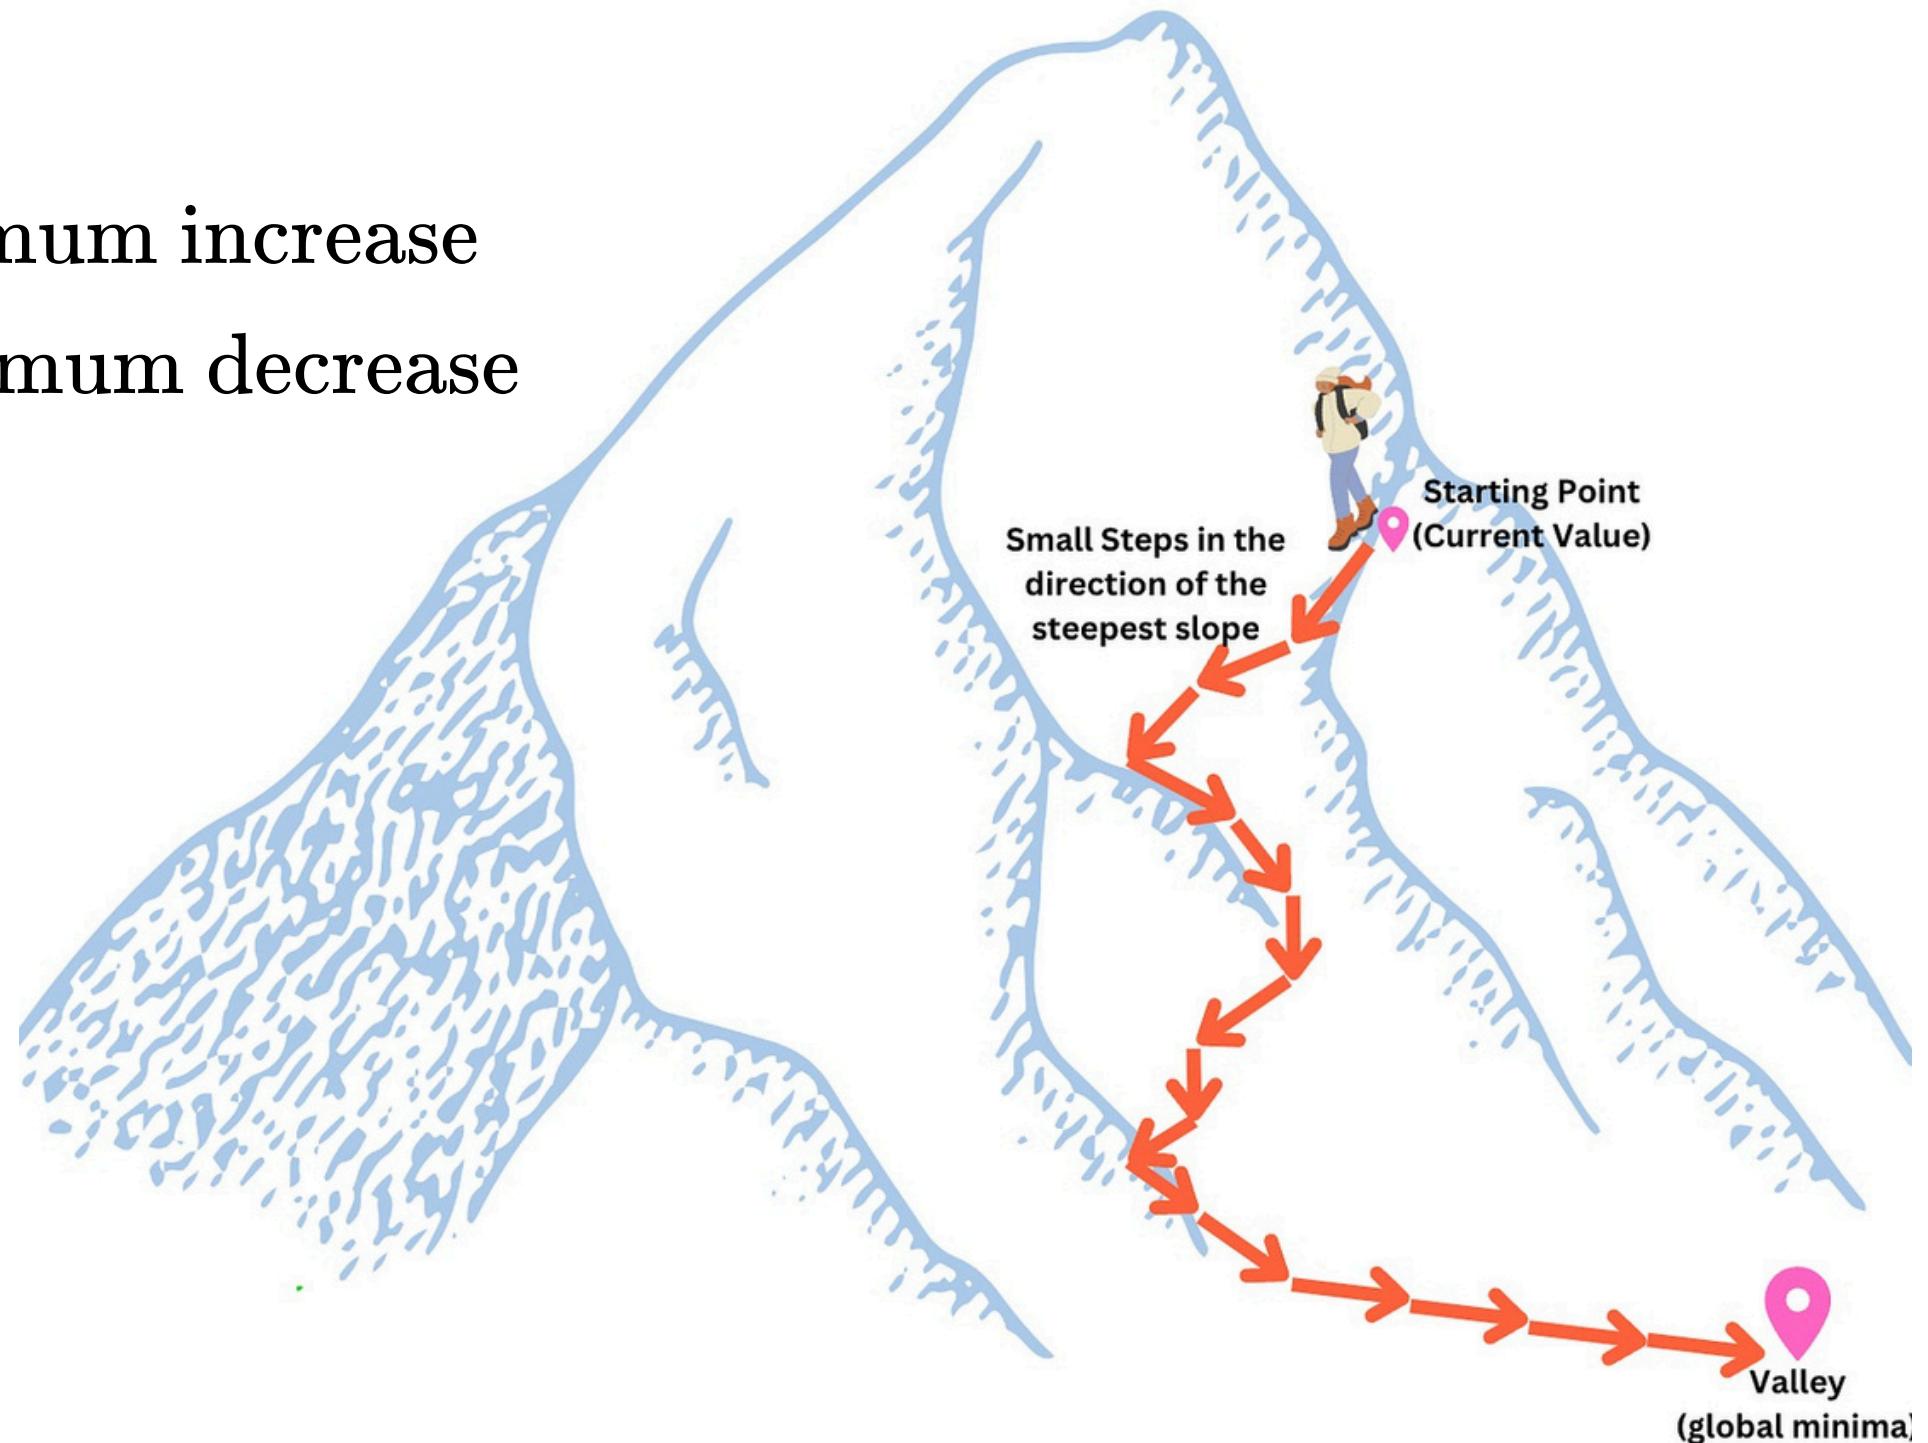
- Optimizers are algorithms that update a model's parameters in order to minimize the loss function during training.



# Batch Gradient Descent

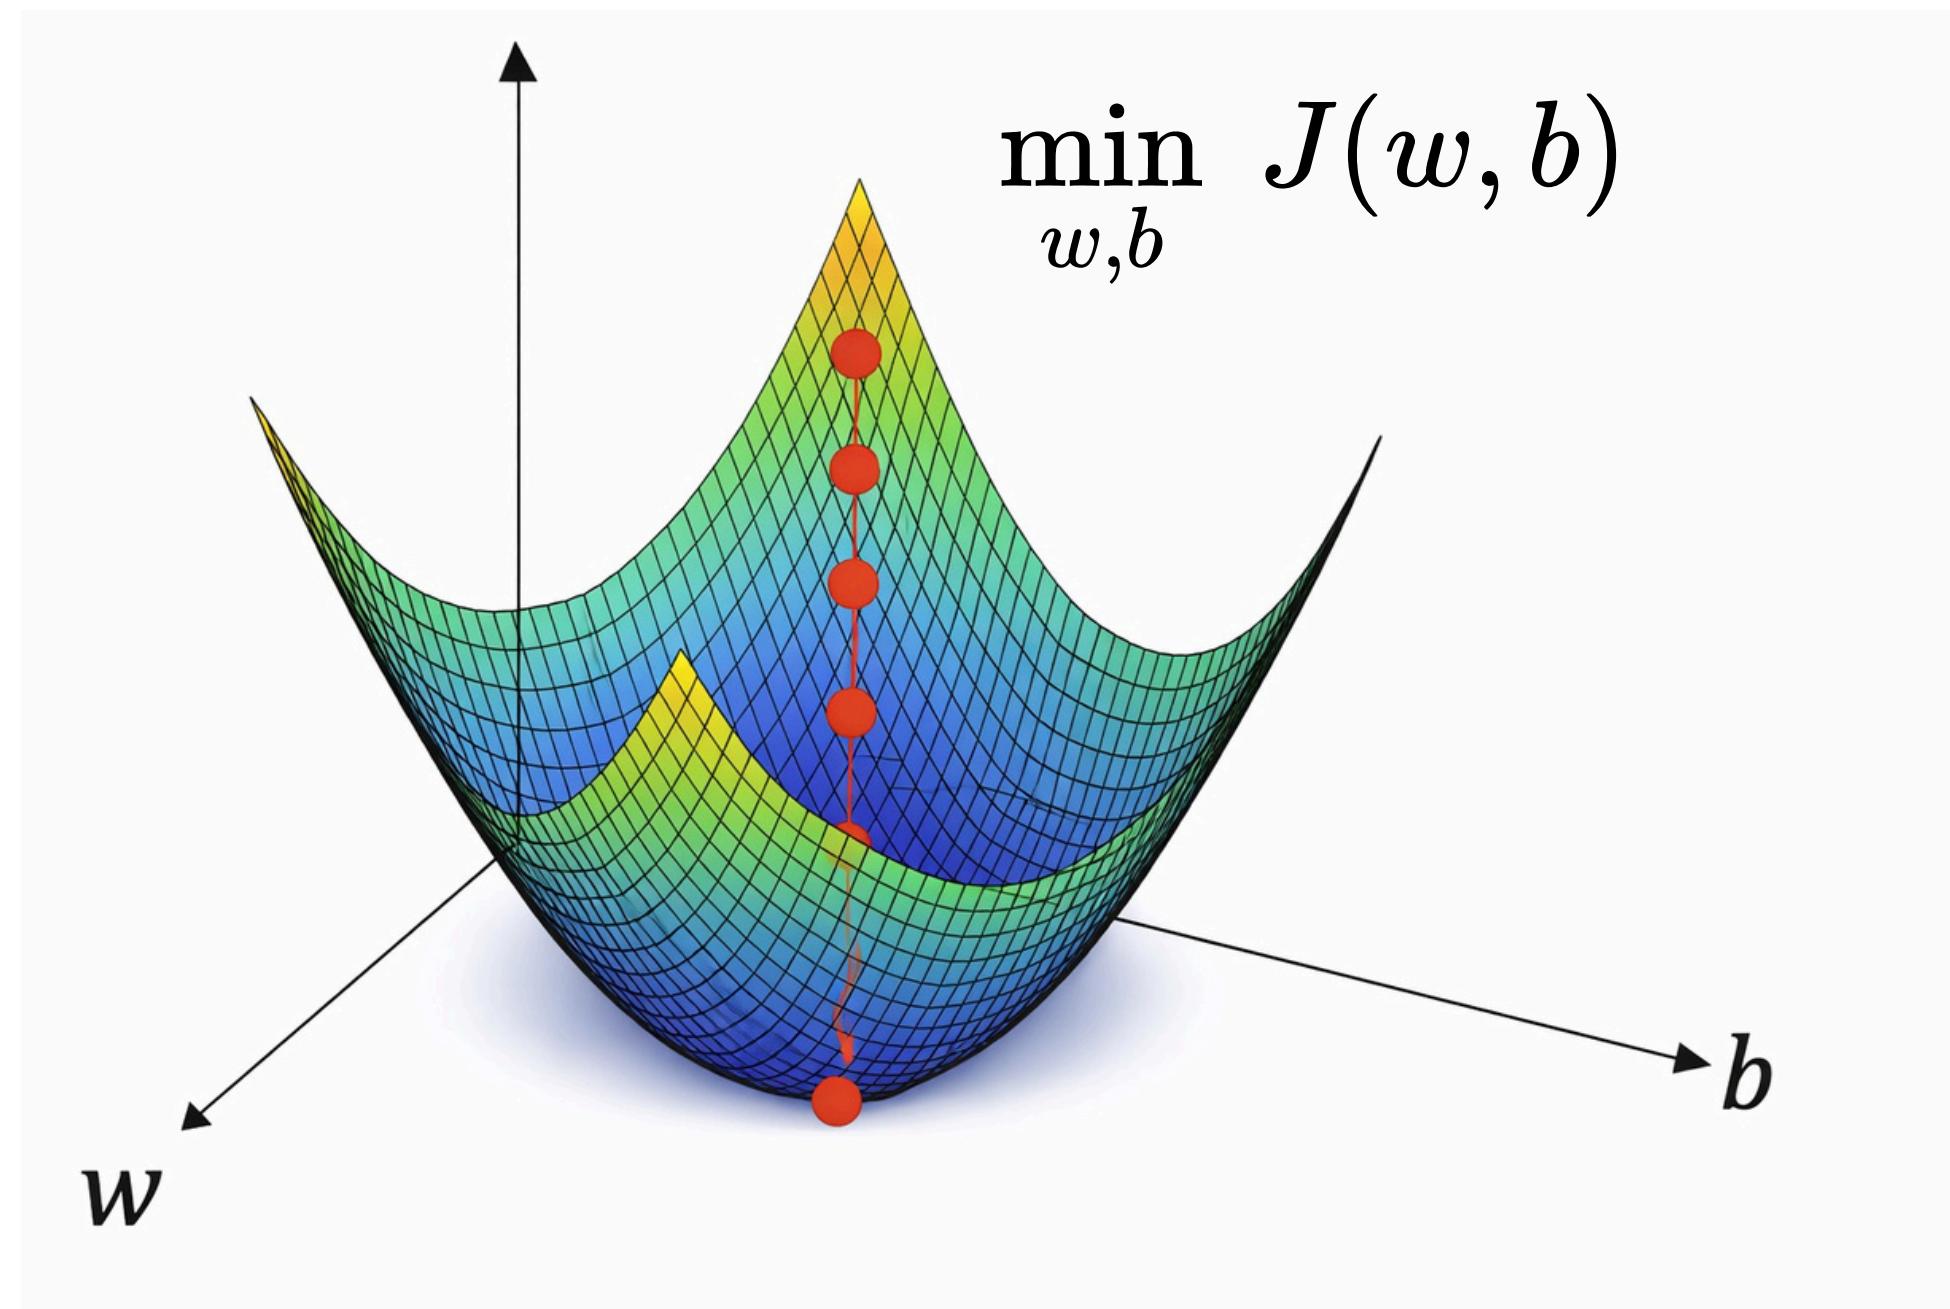
- Direction of Maximum Increase and Decrease

$\nabla f(x)$  : direction of maximum increase  
 $-\nabla f(x)$  : direction of maximum decrease



## Gradient Descent

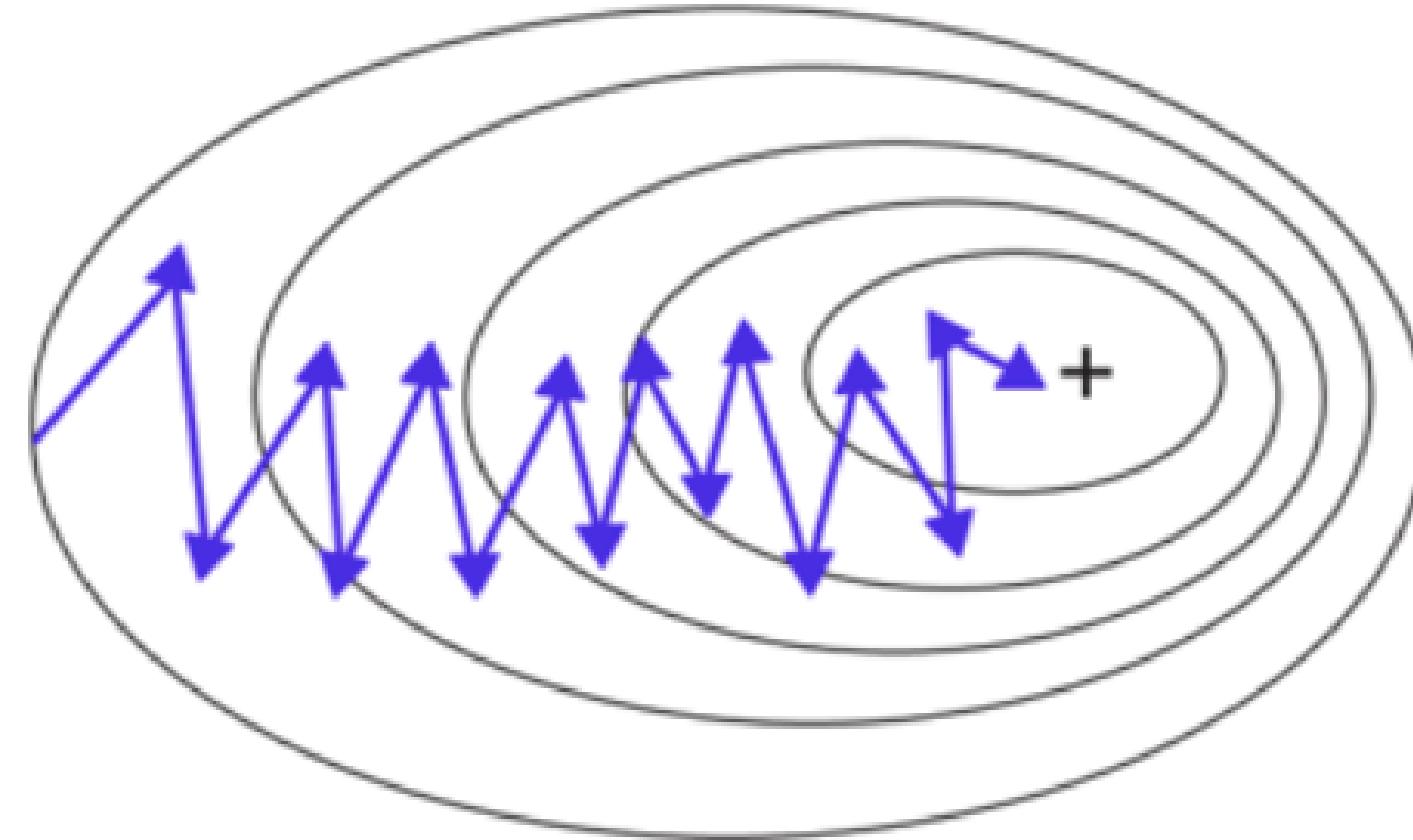
- Moving in the direction of steepest loss decrease



## Stochastic Gradient Descent

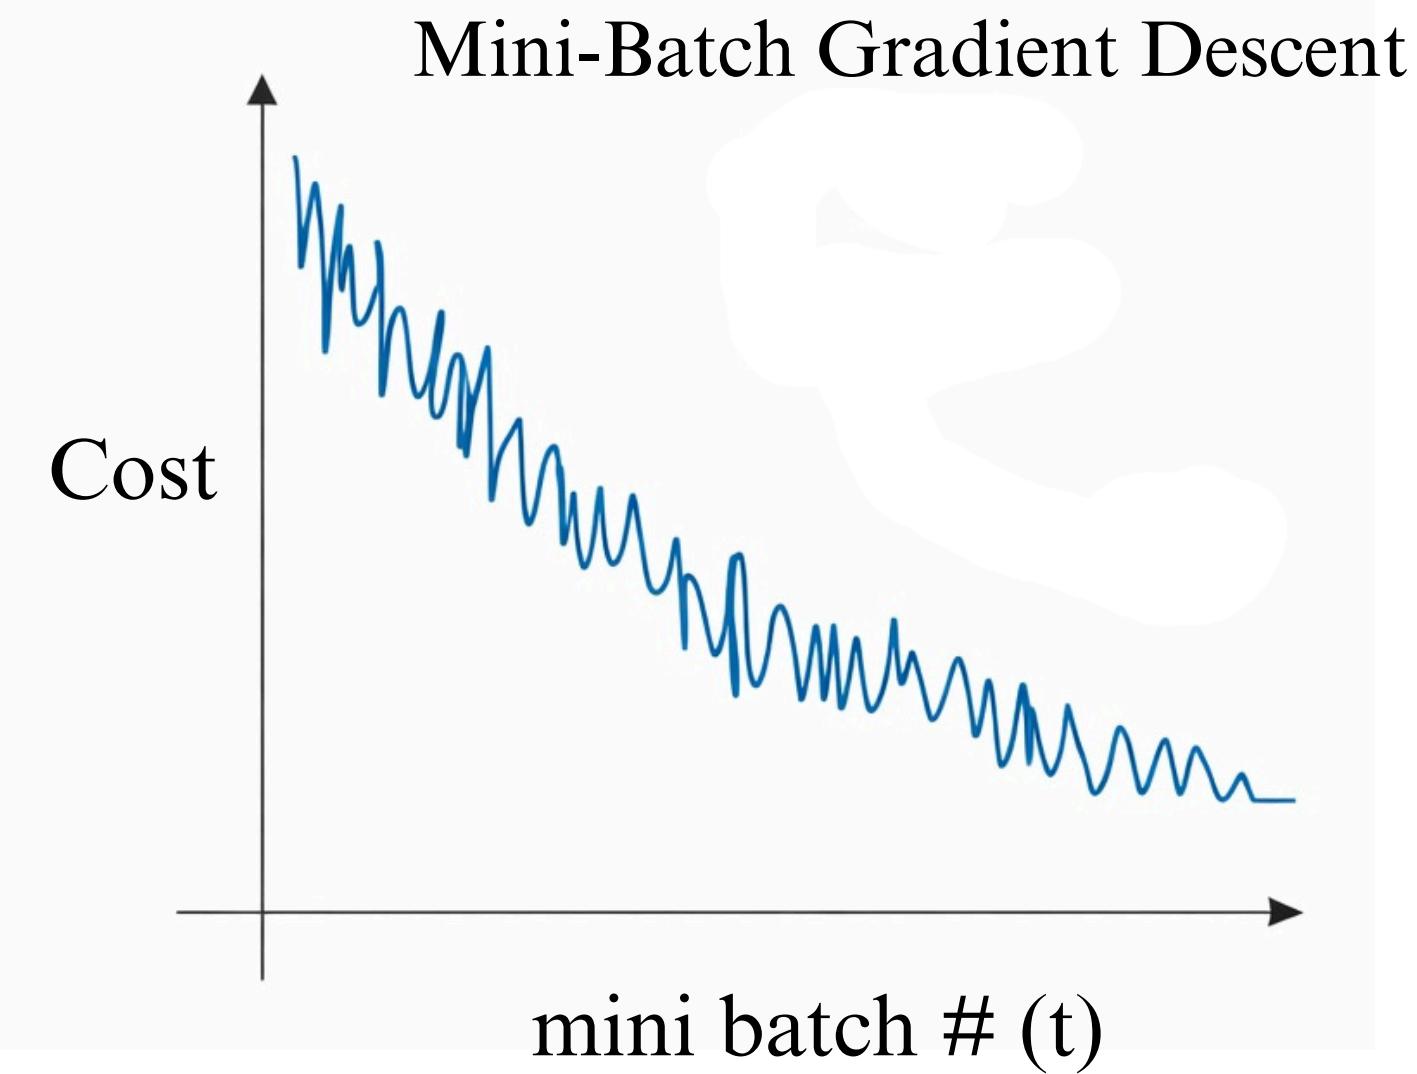
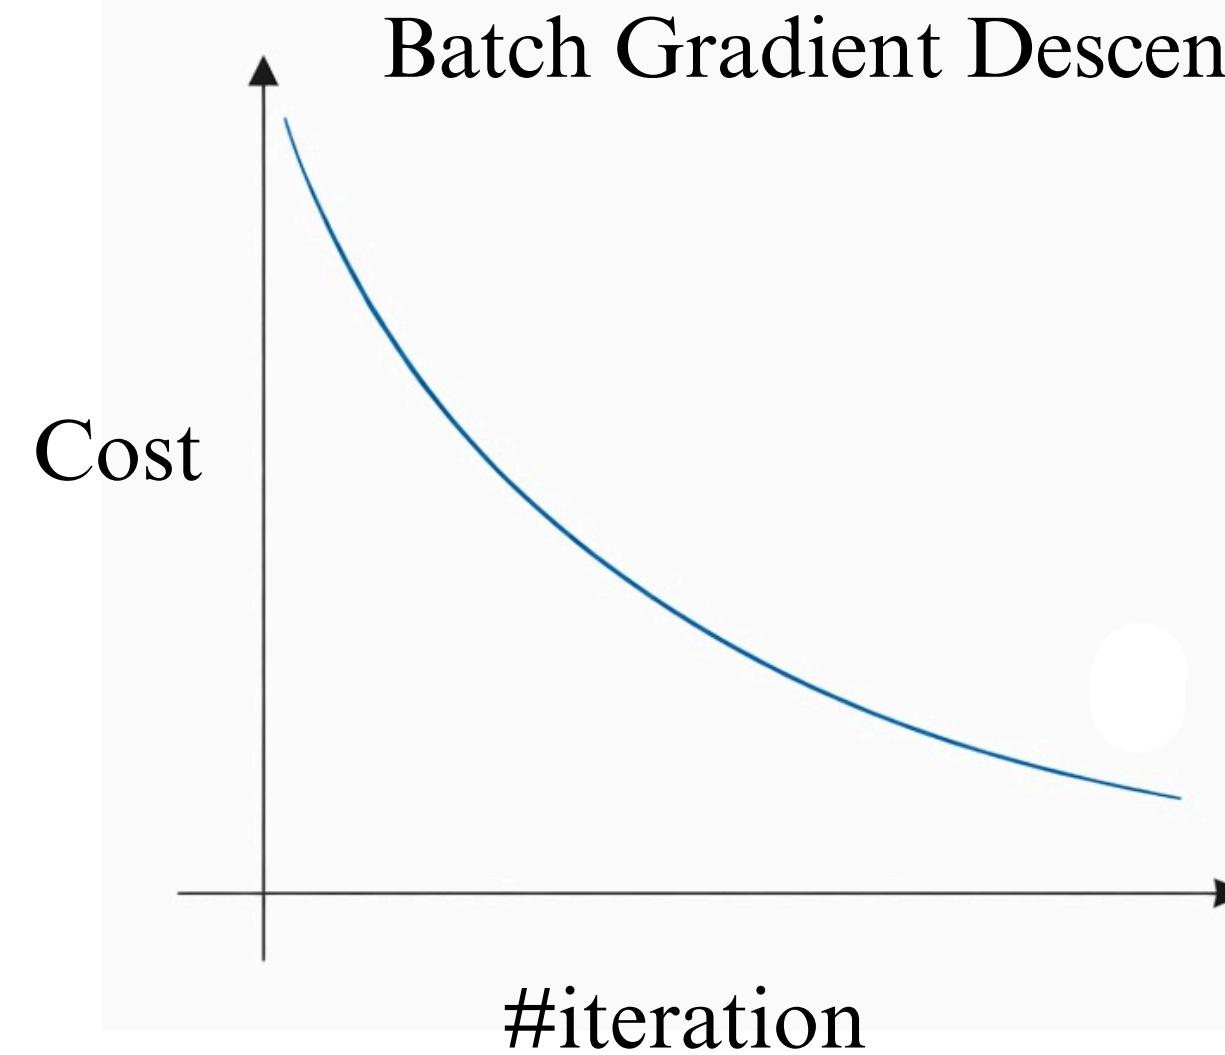
- Updates model parameters using one training sample at a time

$$\theta = \theta - \eta \cdot \nabla L(x_i, y_i)$$



## Mini-Batch Gradient Descent

- Instead of using all training examples or just one, we update parameters using a small batch of samples.



*Common mini-batch sizes are 32/64/128 examples*

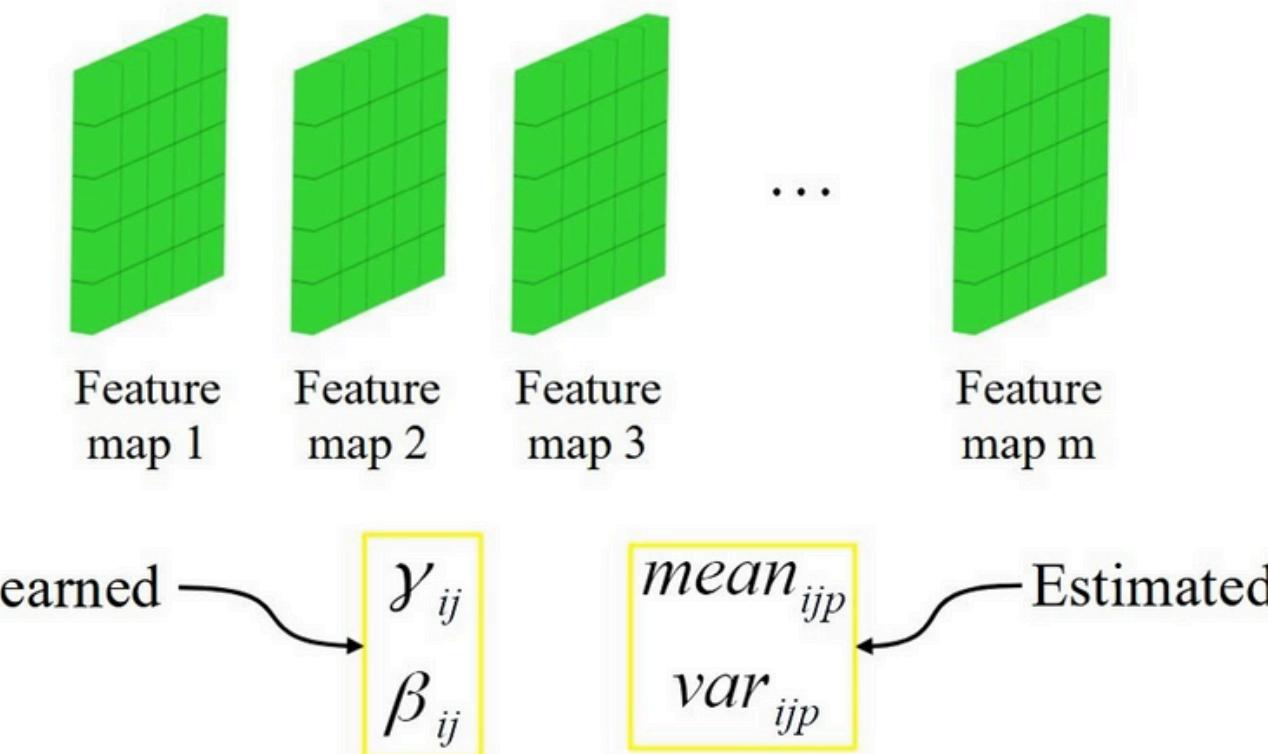


# Batch Normalization

- Uses batch statistics: mean ( $\mu$ ) and variance ( $\sigma^2$ ) of the current mini-batch
- has 2 learnable parameters:
  - $\gamma$  (gamma) controls the scale (how stretched the distribution is)
  - $\beta$  (beta) controls the shift (moves the center left or right)

$$\hat{x} = \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

$$y = \gamma \hat{x} + \beta$$



# Batch Normalization

$$\hat{x} = \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

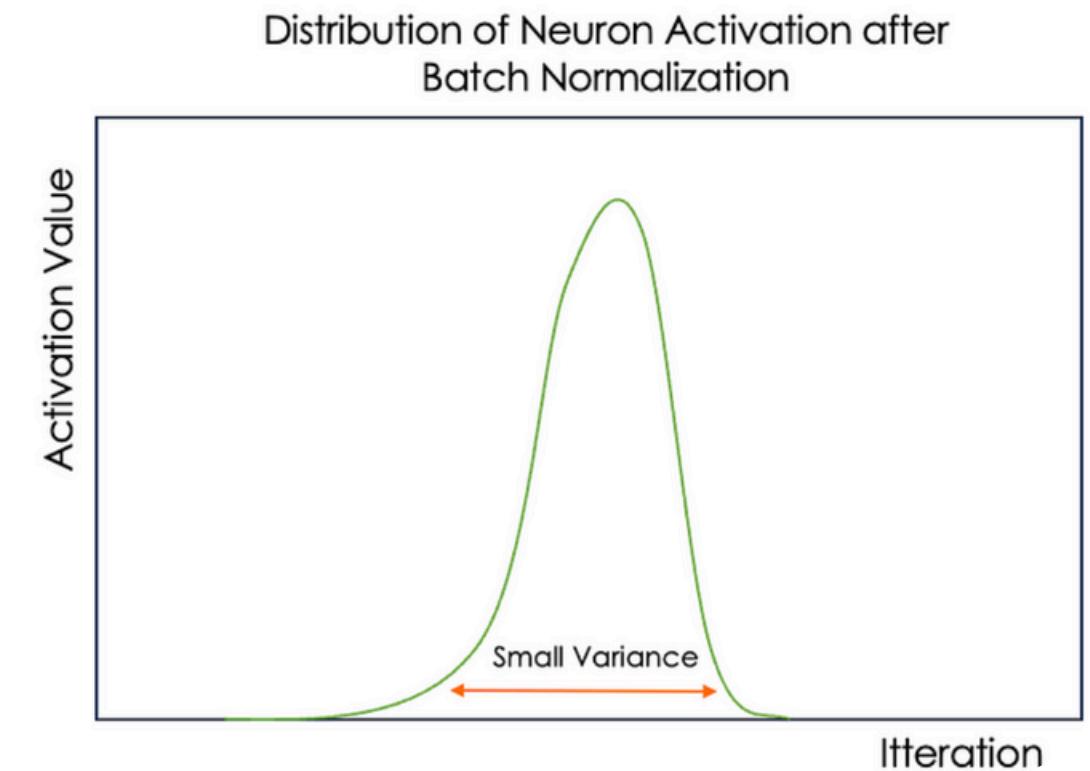
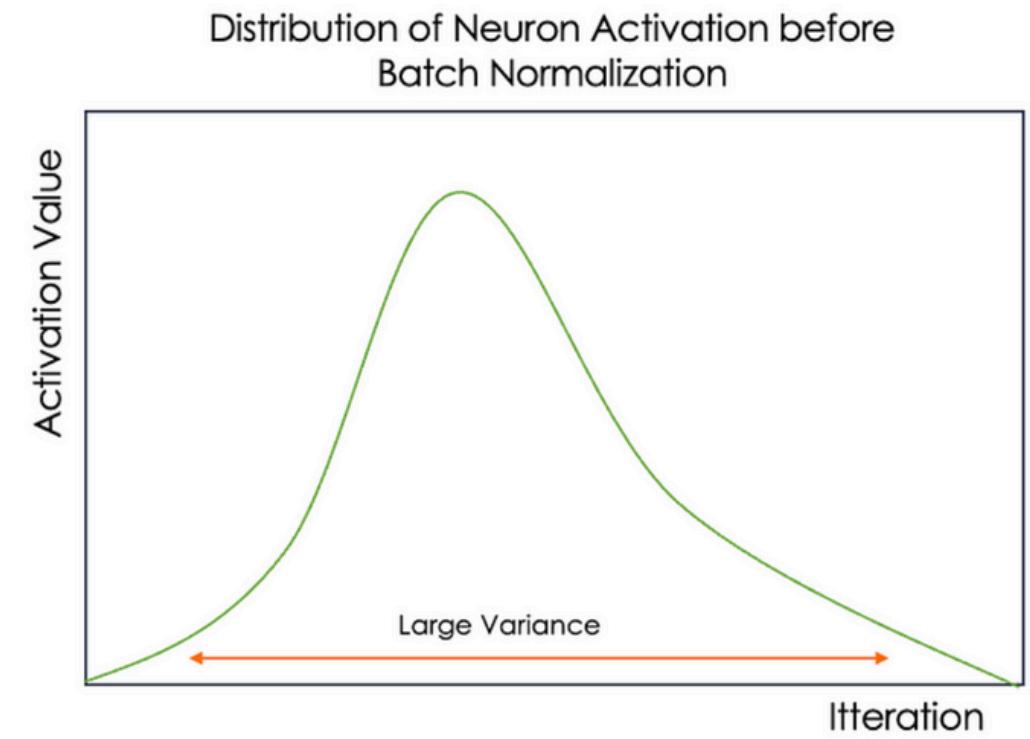
$$x = [2, 4, 6, 8]$$

$$\mu = \frac{2 + 4 + 6 + 8}{4} = 5$$

$$\sigma^2 = \frac{(2 - 5)^2 + (4 - 5)^2 + (6 - 5)^2 + (8 - 5)^2}{4} = \frac{9 + 1 + 1 + 9}{4} = 5$$

$$\hat{x} = [-1.342, -0.447, 0.447, 1.342]$$

$$\mu_{\hat{x}} = 0 \quad \sigma_{\hat{x}}^2 = 1$$



## Batch Normalization

$$y = \gamma \hat{x} + \beta$$

$$\gamma = 2, \beta = 1$$

$$y = 2 \cdot \hat{x} + 1$$

$$y = [-1.684, 0.106, 1.894, 3.684]$$

$$Var(y) = \gamma^2 = 4$$

- Applying  $\gamma$  (scale) stretches the distribution.
- Applying  $\beta$  (shift) moves the entire distribution left/right.
- The distribution is now  $4\times$  wider than before scaling
- $\beta$  does not affect variance only shifts the mean

\*BatchNorm learns statistics during training and reuses them during inference.

# Dropout

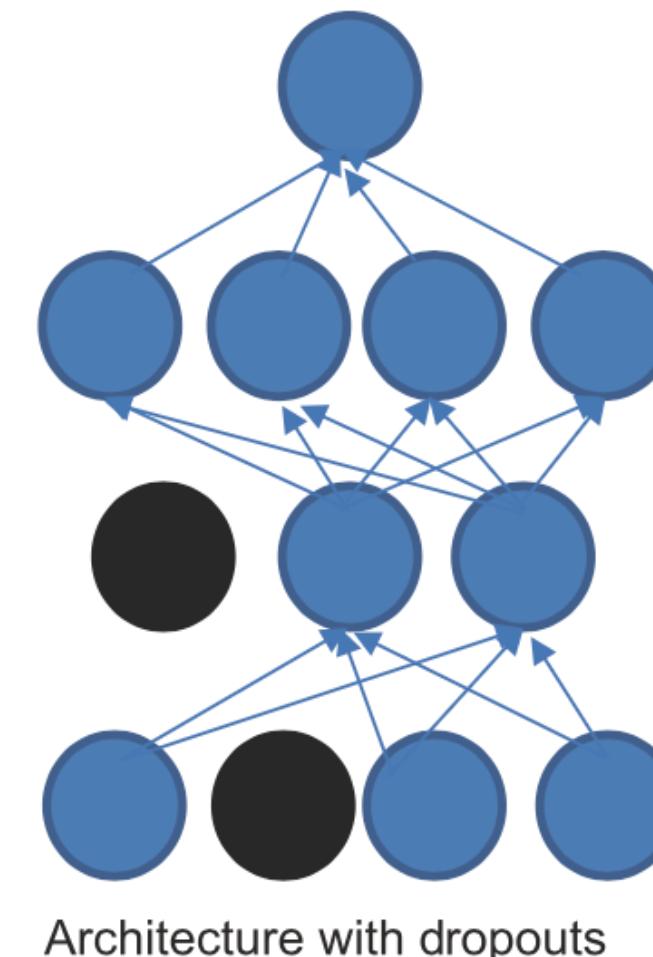
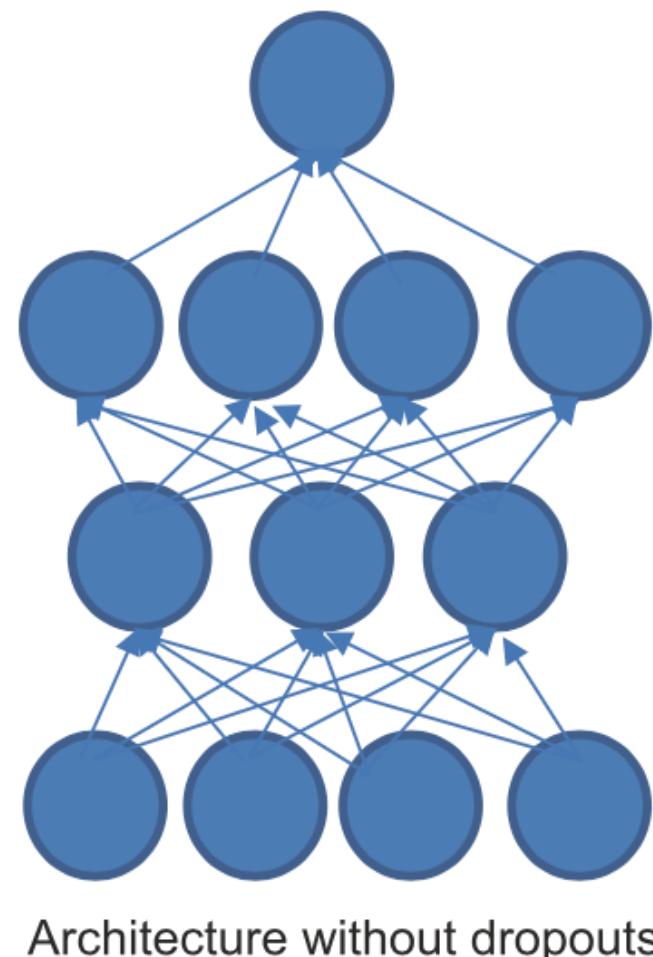
**Input  $x = \{1, 2, 3, 4, 5\}$ ,**

**Dropout  $p = 0.2$ , keep probability = 0.8.**

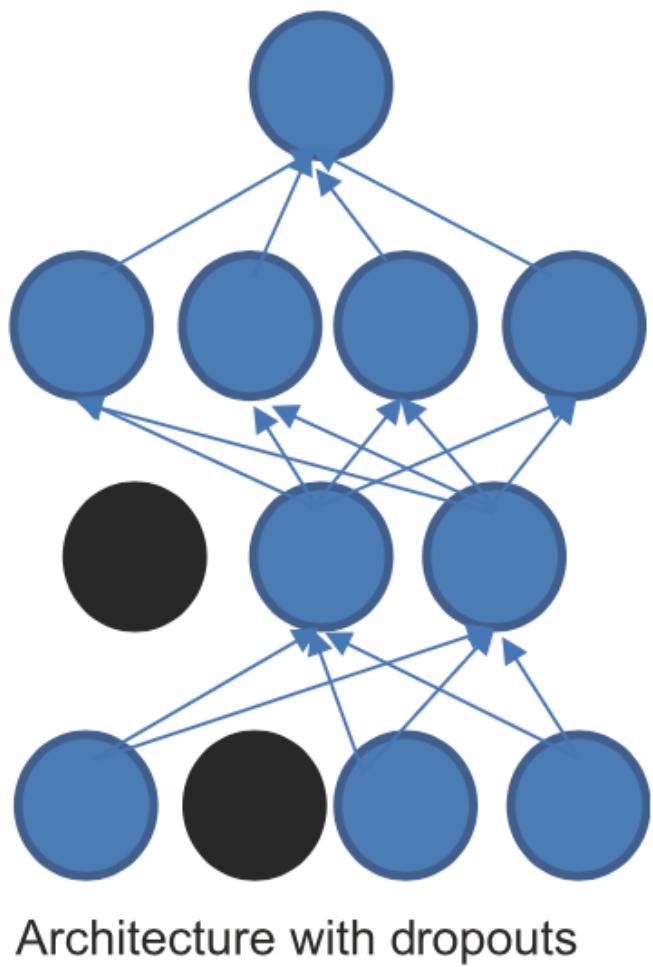
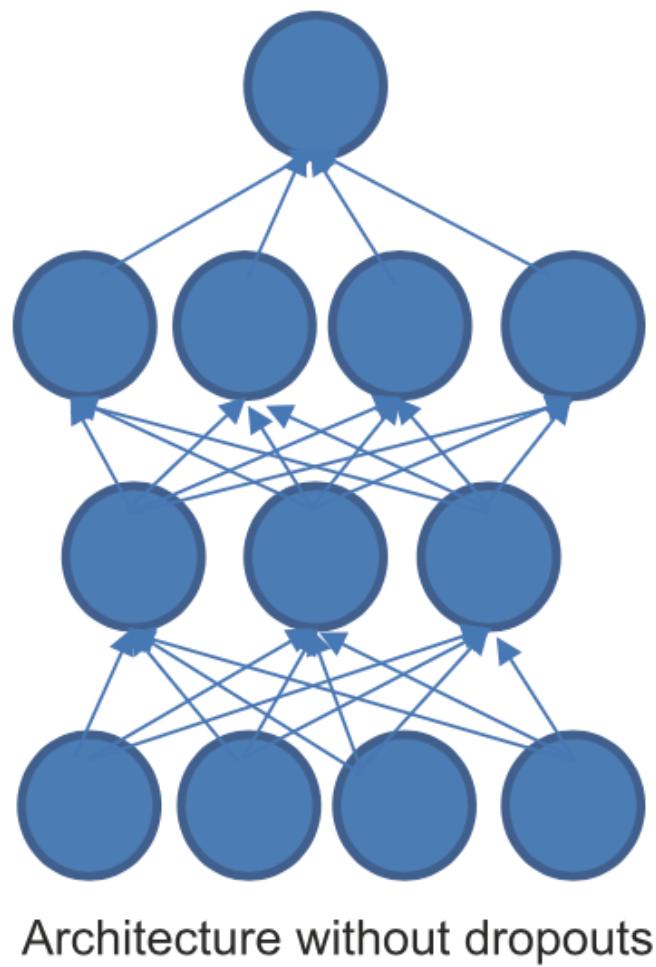
Randomly dropped version may be

$\{1, 0, 3, 4, 5\}$  or  $\{1, 2, 0, 4, 5\}$ .

**Hidden layer example:** 1000 neurons,  $p = 0.5 \rightarrow 500$  neurons removed each iteration.



# Dropout



$$x = [2, 4, 6, 8]$$

$$w = [1, 2, 3, 4]$$

$$y = x \cdot w^T = (2)(1) + (4)(2) + (6)(3) + (8)(4)$$

$$y = 2 + 8 + 18 + 32 = 60$$

$$p = 0.5$$

# Dropout

## Training

$$r = [1, 0, 1, 0]$$

$$x_{\text{drop}} = x \odot r = [2, 0, 6, 0]$$

$$y_{\text{train}} = (2)(1) + (0)(2) + (6)(3) + (0)(4)$$

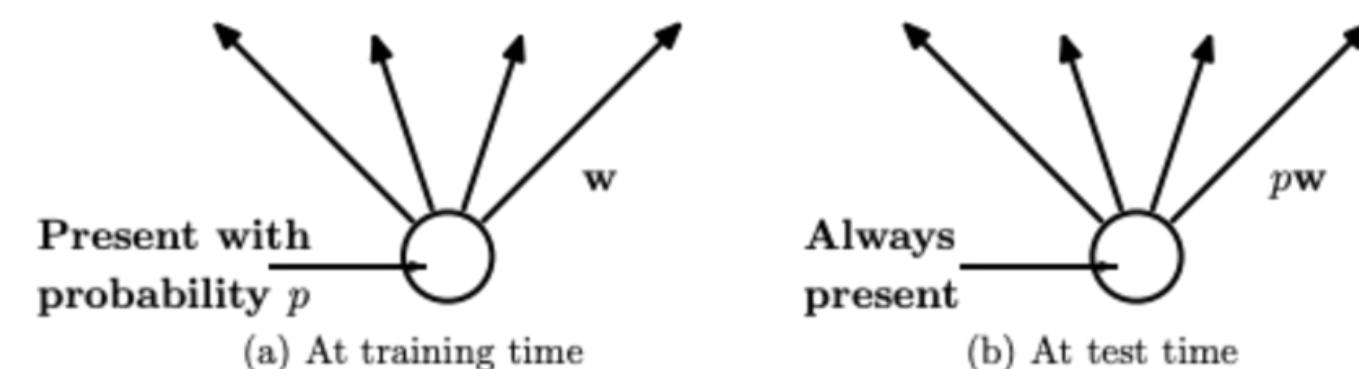
$$= 2 + 0 + 18 + 0 = 20$$

## Inference

$$w' = p \cdot w = [0.5, 1, 1.5, 2]$$

$$y_{\text{test}} = x \cdot w'^T = (2)(0.5) + (4)(1) + (6)(1.5) + (8)(2)$$

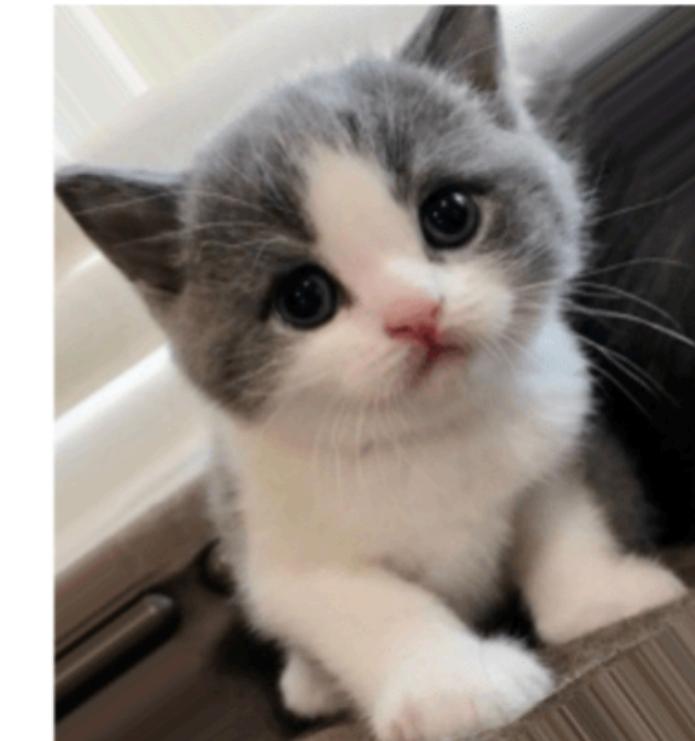
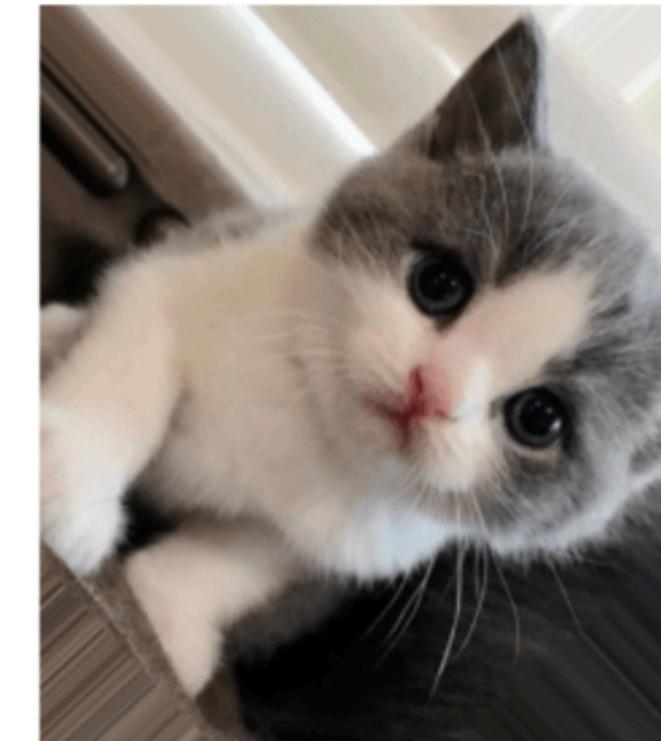
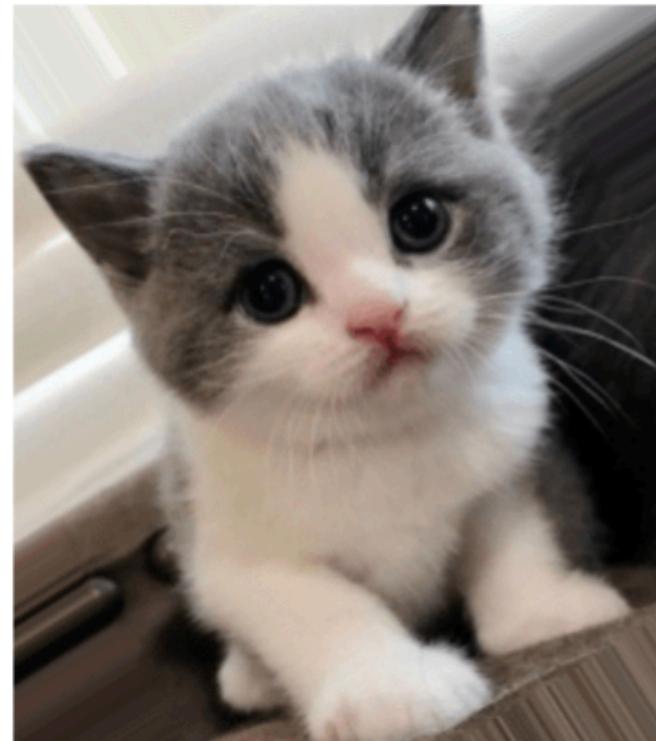
$$y_{\text{test}} = 1 + 4 + 9 + 16 = 30$$



- Training avg = 30, Testing = 30

# Data Augmentation

## Random Rotation



Rotation technique

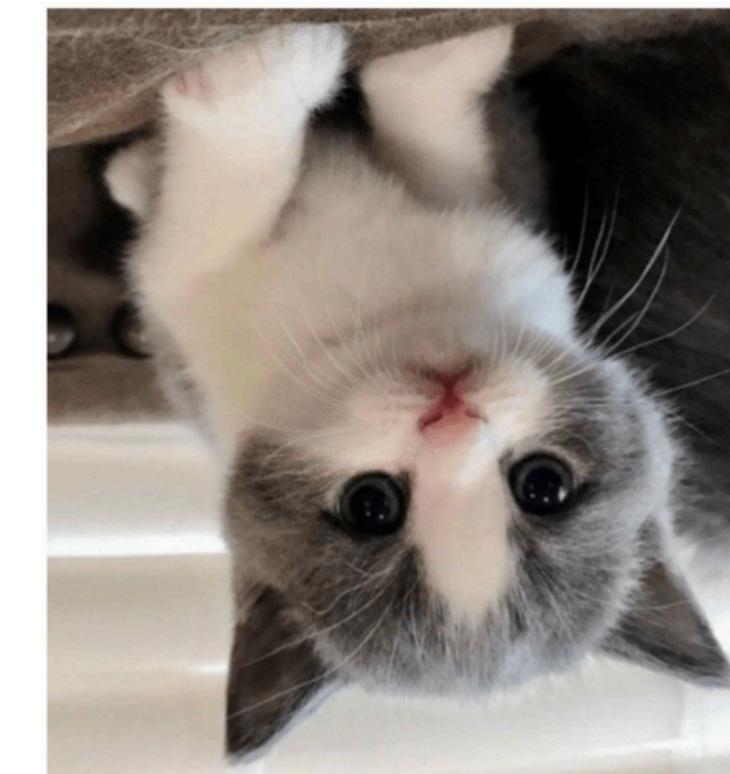
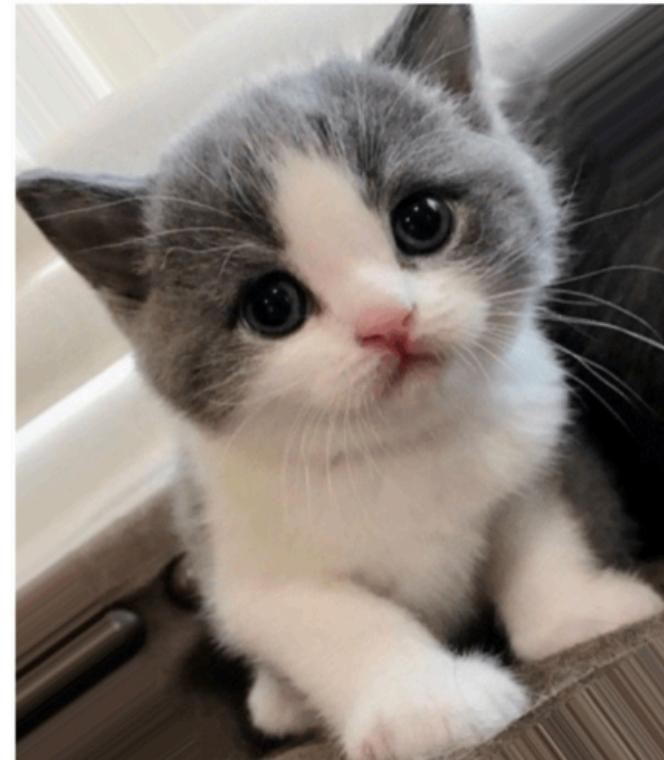
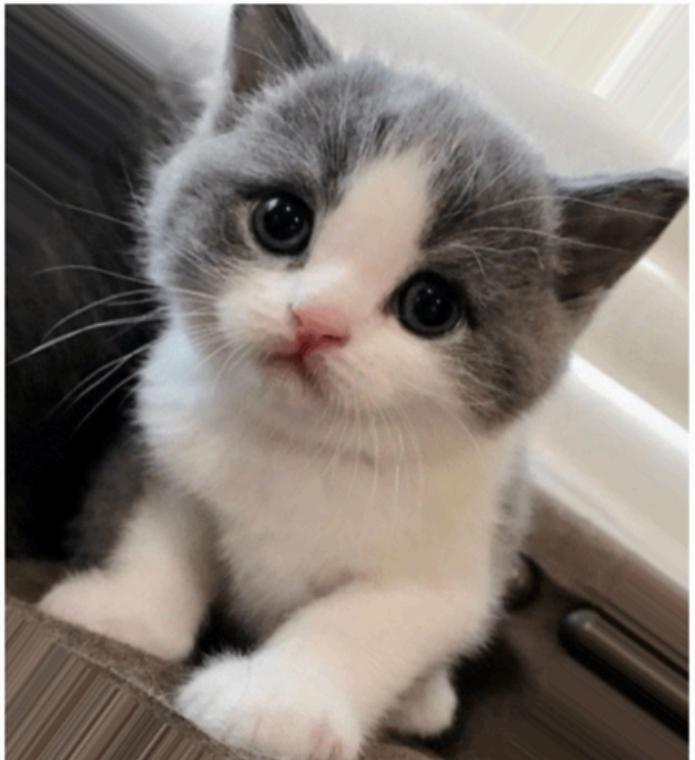
# Data Augmentation

Brightness



# Data Augmentation

Flip

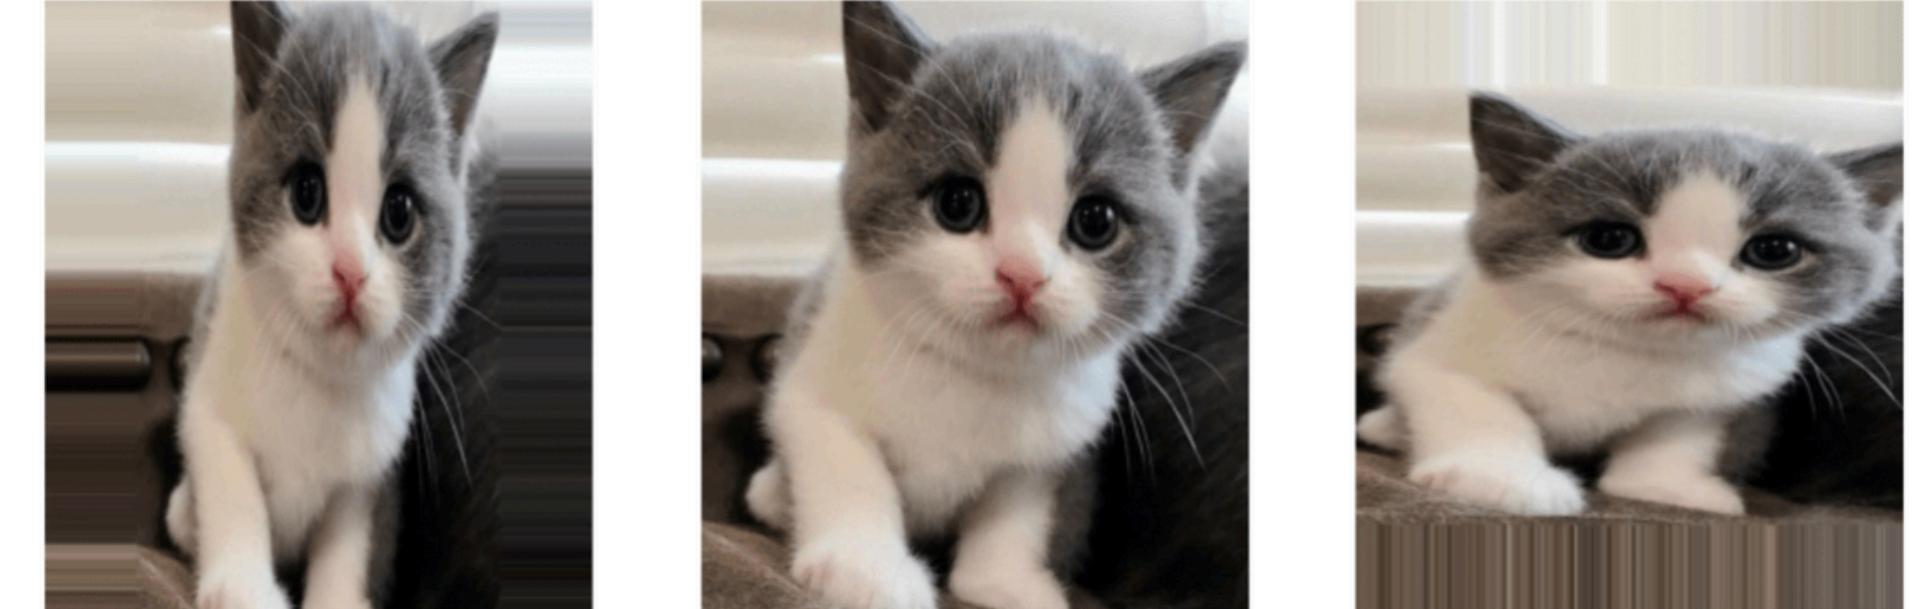


Horizontal flip technique

Vertical flip technique

# Data Augmentation

Zoom



Zoom technique

Random Shift



Width shift technique



Q A

*Thank  
You*