

Talent Bootcamp Series -1

Day-2

Computer Vision Lab

Lab slides prepared by: Hasan Shahid Kamal



Deep Learning Lab Workflow



Load & Preprocess

Load images and labels, then transform raw data for neural networks



Split Data

Divide into training, validation, and test sets



Train & Optimize

Feed batches, compute loss, update parameters over epochs



Evaluate

Measure performance and visualize results

MNIST Dataset Explained

What is MNIST?

A collection of handwritten digits (0-9). Each image is 28×28 grayscale pixels.

The Task

Classification: Given an image, predict which digit it represents across 10 possible classes.

Why MNIST?

Simple, balanced, and clean—perfect for learning deep learning concepts without noisy data.



PyTorch Ecosystem



torch

Tensor operations & automatic differentiation



torch.nn

Neural network layers: Linear, ReLU, Conv2d



torch.optim

Optimization algorithms for parameter updates



torchvision

Vision datasets, transforms, and pretrained models

torchvision simplifies image data handling—automatically downloads, verifies, and loads datasets like MNIST so you can focus on model design.

Torchvision

torchvision is the **computer vision companion library for PyTorch**. It provides everything you typically need to work with **images and videos**—from datasets and preprocessing to models and utilities.

torchvision.datasets — Ready-to-use Vision Datasets

This module provides **pre-built dataset classes** so you don't have to manually download, unzip, and parse data.

What it includes

- Popular **image classification datasets**
- Automatic **download & verification**
- Returns **(image, label)** pairs

Common datasets

- MNIST
- Fashion-MNIST
- CIFAR-10 / CIFAR-100
- ImageNet
- SVHN

Image Preprocessing & Augmentation-torchvision.transforms

Neural networks **cannot work with raw images**. Transforms convert and prepare images before training.

What it includes

- Image → Tensor conversion
- Normalization
- Data augmentation

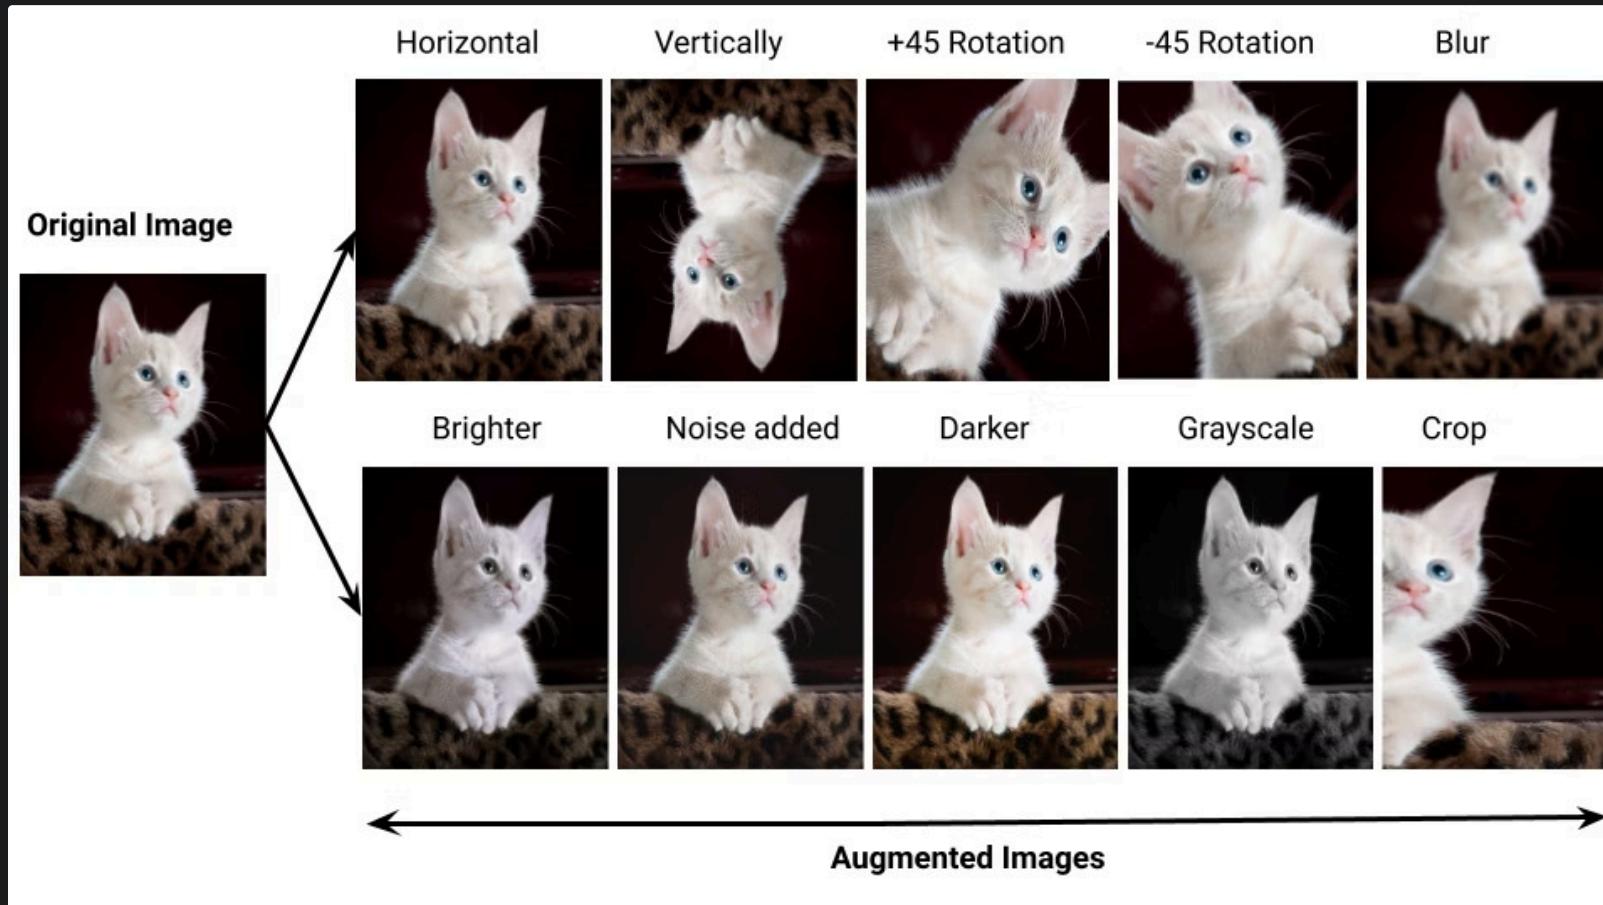
Common transforms

- `ToTensor()` – converts image to PyTorch tensor
- `Normalize()` – scales pixel values
- `Resize()`, `CenterCrop()`
- `RandomHorizontalFlip()`, `RandomRotation()`

Why this matters

- Improves **training stability**
- Augmentation helps **prevent overfitting**
- Standard preprocessing ensures fair comparisons

Example for different transforms



torchvision.models — Pretrained Vision Models

This module provides **state-of-the-art neural network architectures** already implemented and optionally pretrained.

What it includes

- CNN architectures
- Pretrained weights (trained on ImageNet)

Popular models

- ResNet
- VGG
- AlexNet
- MobileNet
- EfficientNet

Why this matters

- No need to implement complex models from scratch
- Enables **transfer learning**
- Faster convergence and better accuracy

Torch.nn module

`torch.nn` **is the neural-network module of PyTorch.**

It provides **everything needed to define, organize, and run neural networks**, but **not** to train them fully (training also needs optimizers, data, etc.).

`torch.nn` = **building blocks + structure for neural networks**

Why do we need `torch.nn`?

Neural networks consist of:

- Layers (Linear, Conv, ReLU, etc.)
- Parameters (weights & biases)
- Loss functions
- Model structure

`torch.nn` provides **ready-made, optimized implementations** of all these.

Without `torch.nn`, you would have to:

- Define weights manually
- Write forward computations yourself
- Track parameters by hand

What does torch.nn contain?

We can divide `torch.nn` into **4 major categories**:

Layers (Core Building Blocks)

Model Base Class: `nn.Module`

Loss Functions (Error Measurement)

Container Modules (Organizing Layers)

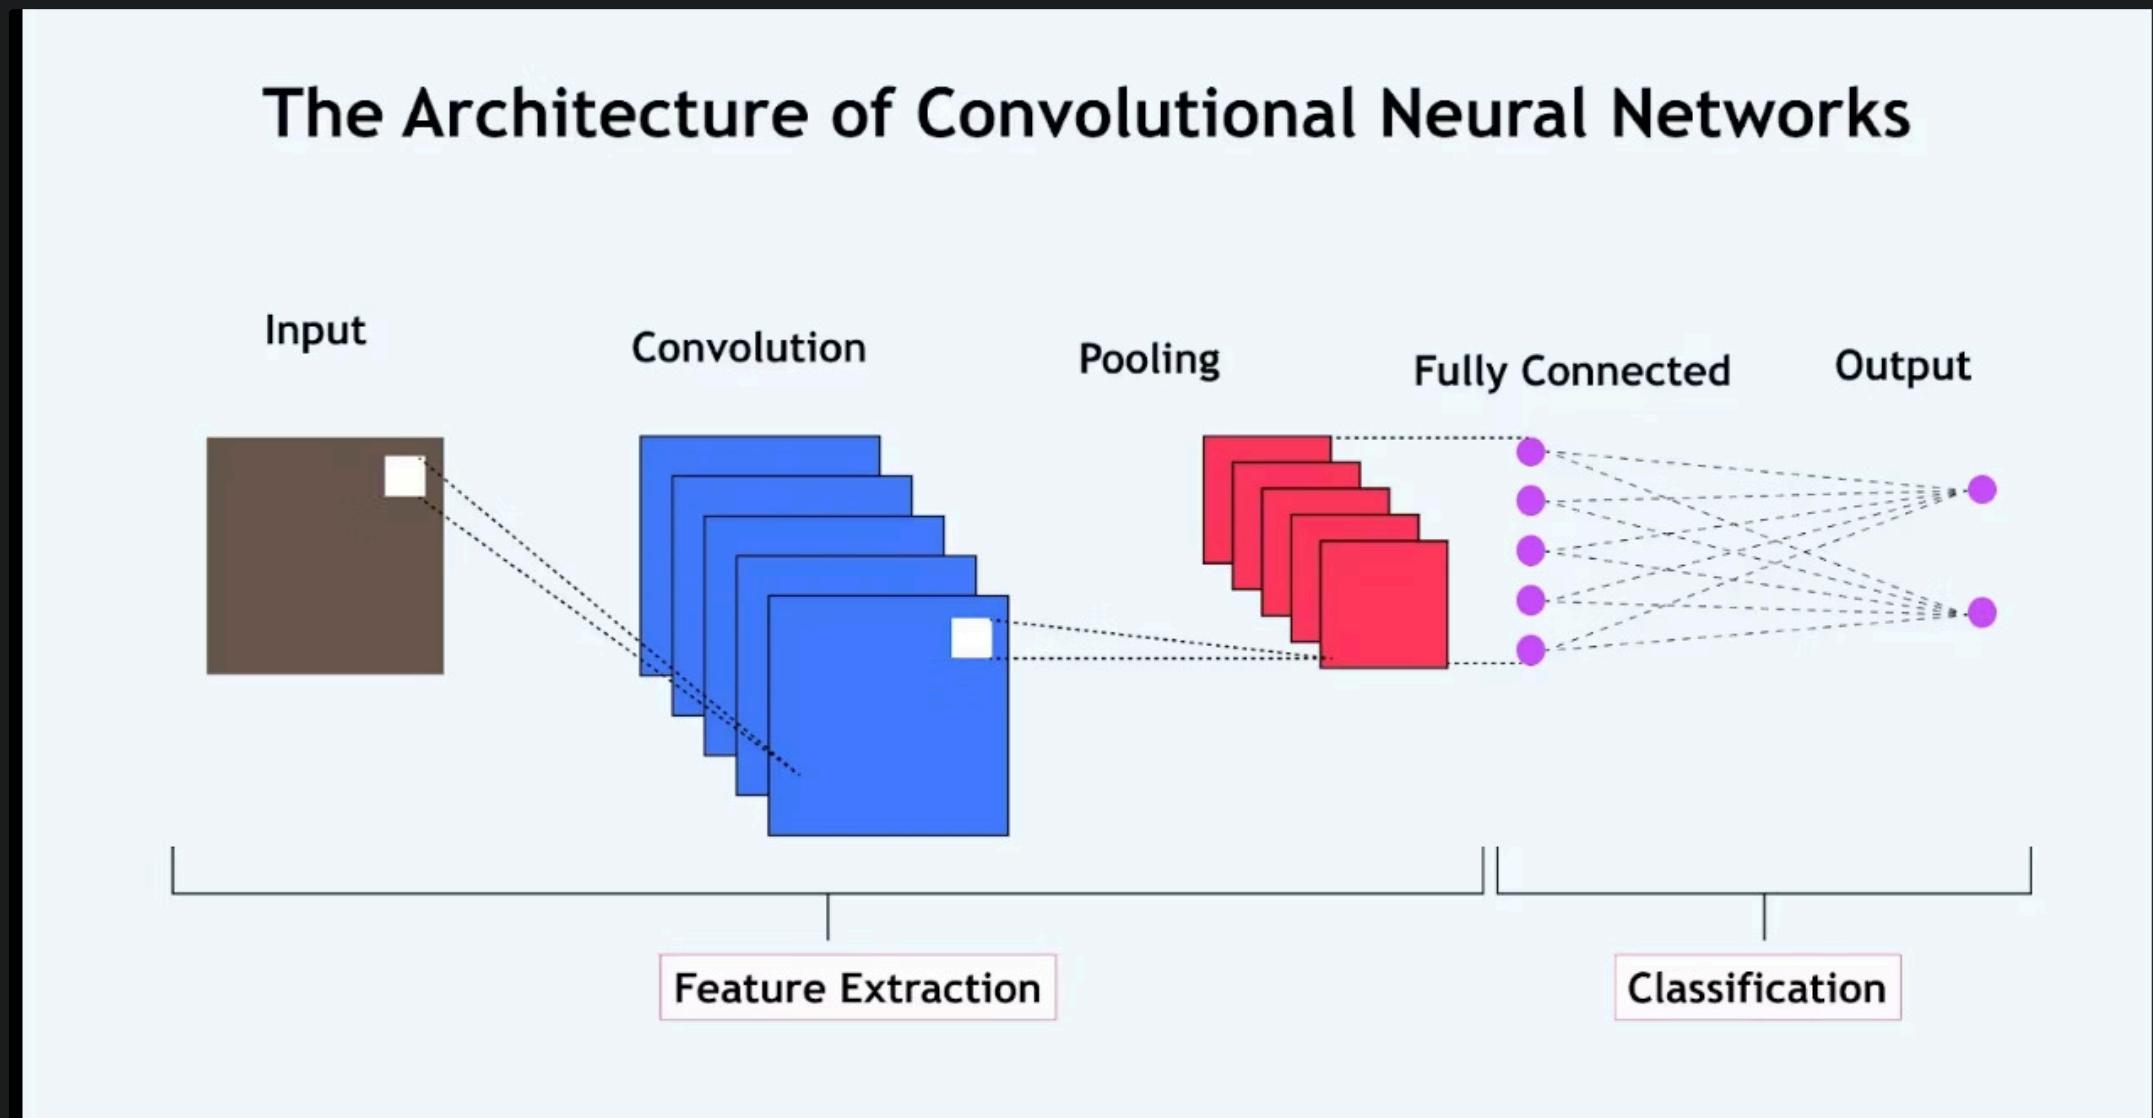
Let's go through each one in detail.

Layers (Core Building Blocks)

These are **learnable transformations**.

Common layers

- `nn.Linear` – fully connected layer
- `nn.Conv2d` – convolutional layer
- `nn.ReLU` – activation function
- `nn.MaxPool2d` – pooling
- `nn.Dropout` – regularization



Loss Functions (Error Measurement)

Loss functions **quantify how wrong the model's prediction is.**

Model Base Class: nn.Module

This is the **most important concept** in torch.nn.

What is nn.Module?

- Base class for **all models and layers**
- Handles parameter registration
- Enables training & evaluation modes

Why models inherit from it

- Automatically tracks parameters
- Works with optimizers
- Supports .train() and .eval()

Example:

python

```
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc = nn.Linear(784, 128)

    def forward(self, x):
        return self.fc(x)
```

Copy code

Container Modules (Organizing Layers)

These help **organize layers cleanly**.

Common containers

- nn.Sequential – layers in order
- nn.ModuleList
- nn.ModuleDict

Problem with nn.Sequential

- Does **not naturally handle reshaping**
- Assumes every step is a module

Example:

```
python
model = nn.Sequential(
    nn.Linear(784, 128),
    nn.ReLU(),
    nn.Linear(128, 10)
)
```

 Copy code

What `torch.nn` does NOT do

`torch.nn` does **not**:

- Update weights
- Perform optimization
- Handle datasets or batching

Those are handled by:

- `torch.optim`
- `torch.utils.data`

DataLoaders, Epochs & Data Splitting

1

DataLoader

Groups data into batches, shuffles to prevent order bias, enables parallel loading

2

Epochs & Batches

Epoch = full dataset pass. $60K$ samples \div 100 batch size = 600 batches per epoch

3

Avoid Data Leakage

Always evaluate on unseen test data
—never on training data

Dataset Class

Used to create a custom Dataset class when the data isn't directly available to use.

Defines how data is accessed:

- `__len__()`: Returns total samples
- `__getitem__()`: Retrieves single data point + label



Neural Network Architecture

01

`nn.Module`

Foundation of PyTorch models—tracks parameters and manages devices

02

`_init_0`

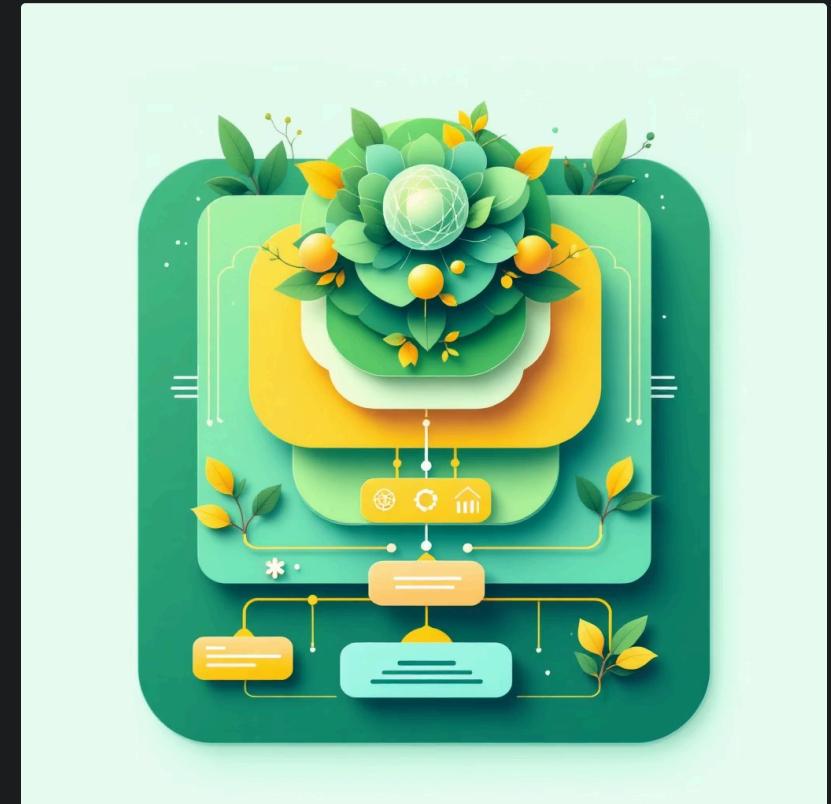
Defines layers and initializes parameters

03

`forward()`

Defines how input flows through layers

The model learns patterns like curves and edges that define each digit.



The enumerate() function

enumerate() returns both an index and a value when looping

```
for i, (images, labels) in enumerate(train_loader):
```

- i → **batch index**
- (images, labels) → **data from the DataLoader**

Why is enumerate() useful in training loops?

- 1 Progress tracking
- 2 Debugging
- 3 Logging & visualization

Loss Functions & Optimizers

Loss Functions

Quantify prediction error. Lower loss = better performance.

- **CrossEntropyLoss:** Classification
- **MSELoss:** Regression

Optimizers

Update parameters based on gradients from backpropagation.

- **SGD:** Fixed learning rate
- **Adam:** Adaptive per-parameter rates

When to use what Loss function?

Loss Function	Problem Type	Model Output	Target Format	When to Use	Example Use Case
<code>nn.CrossEntropyLoss</code>	Multi-class classification	Raw scores (logits), shape (N, C)	Class labels $(N,)$	When each input belongs to exactly one class	MNIST (digits 0-9), CIFAR-10
<code>nn.NLLLoss</code>	Multi-class classification	Log-probabilities	Class labels	When model already applies LogSoftmax	Advanced/custom architectures
<code>nn.BCEWithLogitLoss</code>	Binary / multi-label classification	Raw scores (logits)	0/1 labels	Binary or multiple independent labels	Spam detection, multi-tag images
<code>nn.BCELoss</code>	Binary classification	Probabilities (after sigmoid)	0/1 labels	When sigmoid is applied manually	Rarely used in practice
<code>nn.MSELoss</code>	Regression	Continuous values	Continuous values	Predicting real-valued outputs	House price prediction
<code>nn.L1Loss</code>	Regression	Continuous values	Continuous values	When robustness to outliers is needed	Noisy sensor data
<code>nn.SmoothL1Loss(Huber)</code>	Regression	Continuous values	Continuous values	Balance between MSE and L1	Bounding box regression
<code>nn.KLDivLoss</code>	Distribution matching	Log-probabilities	Probabilities	Comparing probability distributions	Knowledge distillation
<code>nn.CosineEmbeddingLoss</code>	Similarity learning	Embeddings	Similar / dissimilar labels	Measuring similarity	Face verification
<code>nn.TripletMarginLoss</code>	Metric learning	Embeddings	Anchor/positive /negative	Learning distances	Face recognition
<code>nn.CTCLoss</code>	Sequence alignment	Probabilities	Sequences	Unaligned sequence tasks	Speech recognition, OCR

What are logits?

Logits are the raw, unnormalized outputs of a neural network

They are the **numbers produced by the last layer of a model *before* applying any activation like Softmax or Sigmoid.**

Logits = raw class scores

Logits vs Probabilities

Property	Logits	Probabilities
Range	$(-\infty, +\infty)$	$[0, 1]$
Sum to 1?	<input checked="" type="checkbox"/> No	<input checked="" type="checkbox"/> Yes
Direct model output	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> No
Used by loss function	<input checked="" type="checkbox"/> Yes (CrossEntropy)	<input checked="" type="checkbox"/> No
Human-interpretable	<input checked="" type="checkbox"/> No	<input checked="" type="checkbox"/> Yes

QUICK EXAMPLE FOR LOGITS

Example:

text

 Copy code

Logits: [2.3, 0.4, 5.6]

Probabilities: [0.03, 0.01, 0.96]

Now:

- Values are between 0 and 1
- Sum to 1
- Can be interpreted as probabilities

Optimizers

Optimizers in PyTorch (with Functions)

Optimizer	PyTorch Function	Learning Rate Type	Speed	Stability	When to Use
SGD	<code>torch.optim.SGD()</code>	Fixed	Slow	Medium	When you want full control; theory-focused training
SGD + Momentum	<code>torch.optim.SGD(momentum=...)</code>	Fixed	Medium	Good	CNNs and vision tasks; smoother convergence
Adam	<code>torch.optim.Adam()</code>	Adaptive	Fast	Very good	Default choice for most deep learning tasks
RMSprop	<code>torch.optim.RMSprop()</code>	Adaptive	Fast	Good	RNNs, noisy or non-stationary problems
AdamW	<code>torch.optim.AdamW()</code>	Adaptive	Fast	Very good	Modern best practice; better weight decay
Adagrad	<code>torch.optim.Adagrad()</code>	Adaptive	Slow	Medium	Sparse data (e.g., NLP embeddings)

What is `model.parameters()`?

`model.parameters()` returns all the trainable parameters of the model

These parameters are:

- **Weights**
- **Biases**

stored inside layers such as:

- `nn.Linear`
- `nn.Conv2d`
- etc.

These are the values the optimizer updates during training.

Training & Evaluation Pipeline



`model.train()`

Enables dropout & training behaviors

`model.eval()`

Disables for consistent evaluation

Visualize

Loss & accuracy curves show progress

Model.train() vs Model.eval()

model.train() vs model.eval()

Aspect	model.train()	model.eval()
Purpose	Puts the model in training mode	Puts the model in evaluation (inference) mode
When used	During training phase	During validation / testing / inference
Affects weights?	✗ No (weights updated by optimizer, not this)	✗ No
Dropout layers	✓ Enabled (random neurons dropped)	✗ Disabled (all neurons active)
Batch Normalization	✓ Uses batch statistics	✓ Uses running statistics
Parameter updates	✗ Does not update parameters by itself	✗ Does not update parameters
Gradient computation	✓ Gradients are usually computed	✗ Gradients usually disabled (<code>torch.no_grad()</code>)
Random behavior	✓ Yes (dropout, batch stats)	✗ No (deterministic)
<code>Loss.backward()</code>	✓ Yes	✗ No
<code>Optimizer.step()</code>	✓ Yes	✗ No
Common mistake	Forgetting to switch back to train mode	Forgetting to switch to eval mode



FINAL CHAPTER

Key Takeaways



Device Management

GPUs accelerate matrix operations. Model and data must be on the same device.



Reproducibility

Set random seeds to ensure consistent results and fair comparisons.



Understanding > Memorizing

MNIST provides a foundation for advanced computer vision tasks.